



POLITECNICO DI TORINO

Corso di laurea in
Ingegneria Matematica

Numerical optimization for large scale problems and Stochastic
Optimization - Assignment on Unconstrained Optimization

Elisabetta Roviera s328422

21th February 2025

Contents

1	Introduction	2
2	Description of the Methods	2
2.1	Modified Newton Method	3
2.2	Truncated Newton Method	5
3	Results and comparison between methods	7
3.1	Rosenbrook Function	7
3.2	Problem 16. Banded trigonometric problem.	12
3.2.1	Dimension $n = 10^3$	13
3.2.2	Dimension $n = 10^4$	17
3.2.3	Dimension $n = 10^5$	20
3.3	Problem 27. Penalty Function	21
3.3.1	Dimension $n = 10^3$	22
3.3.2	Dimension $n = 10^4$	26
3.3.3	Dimension $n = 10^5$	27
3.4	Problem 76.	29
3.4.1	Dimension $n = 10^3$	30
3.4.2	Dimension $n = 10^4$	34
3.4.3	Dimension $n = 10^5$	38
4	Final considerations	39
A	Appendix: Problem implementation and Main Code	40
B	Appendix: Modified Newton Method	45
C	Appendix: Truncated Newton Method	48
D	Appendix: Finite Differences	61

1 Introduction

This document presents the implementation and analysis of two numerical methods for unconstrained optimisation: the Modified Newton Method and the Truncated Newton Method. These methods, which are variations of the classical Newton Method, aim to minimise functions of the form

$$\min_x f(x),$$

where x is a real-valued variable and f is a real-valued function. The study adheres to the guidelines stipulated in the project assignment, which necessitates the implementation of two optimisation techniques integrated with a backtracking line search strategy to ensure sufficient decrease conditions. The implementation has been conducted in MATLAB, and the performance of these methods has been evaluated on four distinct optimisation problems, encompassing the well-known Rosenbrock function and three additional problems derived from reference datasets [1].

A pivotal element of this study is the assessment of the methods' efficiency and accuracy under varied conditions. The Rosenbrock function is utilised as an initial test case, with specified starting points, to observe the convergence behaviour. The remaining test problems are studied in higher dimensions with random initial conditions to further assess the robustness of the system. Additionally, experiments have been conducted to compare the use of exact derivatives versus numerical approximations obtained via finite differences. The accuracy of finite difference approximations is explored by varying the step size parameters across multiple values to understand their impact on convergence.

Following the execution of the optimisation methods in accordance with the prescribed experimental setup, a comprehensive comparison has been conducted. The evaluation criteria encompass the number of successful runs, the iterations required to meet the stopping criterion, the experimental rate of convergence, the execution time, and other pertinent performance indicators. The results are presented in a structured manner using tabular and graphical representations, followed by an analytical discussion that contextualises the findings in relation to theoretical expectations. The entire MATLAB code used to generate the results is available in the appendix.

2 Description of the Methods

Newton's method, also referred to as the Newton-Raphson method, is a fundamental iterative technique used to find successively better approximations to the roots (or zeroes) of a real-valued function. In the context of optimisation, it is employed to locate stationary points where the gradient of the objective function vanishes, indicating potential minima, maxima, or saddle points. The method utilises the first and second derivatives of the function to inform its iterations, providing a more direct route to the solution compared to first-order methods like gradient descent.

The primary objective of Newton's method in optimization is to solve for x in $\nabla f(x) = 0$, where f is a twice-differentiable function mapping \mathbb{R}^n to \mathbb{R} , and $\nabla f(x)$ denotes the gradient vector of f at x . The iterative update rule is given by:

$$x_{k+1} = x_k - [\nabla^2 f(x_k)]^{-1} \nabla f(x_k)$$

Here, $\nabla^2 f(x_k)$ represents the Hessian matrix of second-order partial derivatives at x_k . This update formula adjusts the current estimate x_k by the product of the inverse Hessian and the gradient, effectively utilizing curvature information to inform the step direction and magnitude.

A notable strength of Newton's method is its convergence rate, which, under appropriate conditions, exhibits quadratic convergence. That is to say, when the initial guess is sufficiently close to a solution and the Hessian is positive definite, the method demonstrates a decrease in error between the iterates and the actual solution that is proportional to the square of the previous error. This results in rapid convergence in the vicinity of the solution. However, it should be noted that this rapid convergence is local; the method may not perform well if the initial guess is far from the optimal point or if the Hessian is not positive definite. Despite its advantages, Newton's method has limitations. Computing the inverse of the Hessian matrix can be computationally expensive, especially for high-dimensional problems, as it

requires operations of order $O(n^3)$. Additionally, if the Hessian is singular or nearly singular, the method can encounter difficulties. To address these challenges, alternative methods such as the Damped Newton Method have been proposed. This method incorporates a step size parameter, ensuring global convergence and enhancing stability. Furthermore, quasi-Newton methods have been developed to approximate the Hessian, thereby reducing the computational burden while maintaining superlinear convergence rates. The methods implemented are variations of the Newton method, which is based on the approximation of the function f at each iteration k using a second-order Taylor expansion around the point $x^{(k+1)} = x^{(k)} + p$:

$$f(x^{(k)} + p) \approx f(x^{(k)}) + \nabla f(x^{(k)})^T p + \frac{1}{2} p^T \nabla^2 f(x^{(k)}) p.$$

The descent direction $p^{(k)}$ is then obtained by solving the linear system:

$$\nabla^2 f(x^{(k)}) p = -\nabla f(x^{(k)}).$$

It is evident that the direction p_k is a descent direction at iteration k only if the Hessian $\text{Hess}(f(x_k))$ is positive definite. In the event that this is not the case, the obtained direction may not be the correct one, requiring corrections. In such cases, the Modified Newton Method or the Truncated Newton Method may be employed. The convergence rate of the Truncated Newton and Modified Newton methods varies depending on the problem and specific implementation. Under certain assumptions, the pure Newton method exhibits quadratic convergence to the solution, and in the best-case scenario, both the modified and truncated Newton methods can also achieve quadratic convergence. However, due to the use of inexact line search techniques, convergence often results in only linear progress. Moreover, since the minima of the functions are not known except for the first iteration, an approximation of the true error at step k is given by:

$$e_k \approx \|x_k - x_{k-1}\|_2.$$

To compute the convergence rate, the following formula is used:

$$p \approx \frac{\log\left(\frac{\|e_{k+1}\|}{\|e_k\|}\right)}{\log\left(\frac{\|e_k\|}{\|e_{k-1}\|}\right)}.$$

In the ensuing sections, a concise discourse on the Modified and Truncated Newton methods will be explained.

2.1 Modified Newton Method

The Modified Newton Method is a variant of the classical Newton method that addresses one of its primary weaknesses: the possible indefiniteness of the Hessian matrix $\nabla^2 f(x)$. In cases where $\nabla^2 f(x)$ is not positive definite, a correction is applied to ensure that the modified Hessian is sufficiently positive definite. This is achieved by adding a corrective matrix E_k , resulting in a new matrix

$$B_k = \nabla^2 f(x_k) + E_k,$$

which is designed to have all its eigenvalues bounded away from zero by a positive constant. This modification ensures that the linear system used to compute the Newton direction is well-posed and that the search direction is indeed a descent direction.

Hessian Modification. At each iteration k , the method examines the Hessian $\nabla^2 f(x_k)$ and determines whether it is sufficiently positive definite. The goal is to enforce

$$\lambda_{\min}(B_k) \geq \delta > 0,$$

where $\lambda_{\min}(B_k)$ denotes the minimum eigenvalue of B_k , and δ is a predefined tolerance. A common and effective choice for the correction is to set

$$E_k = \tau_k I,$$

with I being the identity matrix and τ_k a nonnegative scalar that is adjusted adaptively. If $\nabla^2 f(x_k)$ is already positive definite, one may choose $\tau_k = 0$. Otherwise, τ_k is increased until B_k meets the condition for positive definiteness. Since computing the exact eigenvalues of B_k can be expensive, a bound such as

$$\lambda_i(B_k) \leq \|B_k\|_F,$$

where $\|B_k\|_F$ denotes the Frobenius norm, is sometimes used to guide the selection of τ_k .

Algorithm Description. Given an initial point x_0 , the Modified Newton Method proceeds iteratively as follows:

Iteration 1: Hessian Modification and Factorization. For the current iterate x_k , form the modified Hessian

$$B_k = \nabla^2 f(x_k) + \tau_k I.$$

The parameter τ_k is chosen based on a preliminary check. For instance, one may set an initial value τ_0 as follows:

- Compute a measure β , e.g., the Frobenius norm of $\nabla^2 f(x_k)$.
- If the minimum of the diagonal entries of $\nabla^2 f(x_k)$ is positive, set $\tau_0 = 0$.
- Otherwise, set $\tau_0 = \beta/2$.

An attempt is then made to perform an (incomplete) Cholesky factorization on B_k . If the factorization is successful, the resulting factor (typically a lower triangular matrix L) is used to solve the linear system. If the factorization fails (indicating that B_k is not yet sufficiently positive definite), τ_k is updated, for example by

$$\tau_{k+1} = \max(2\tau_k, \beta/2),$$

and the process is repeated until a successful factorization is achieved.

Iteration 2: Computing the Search Direction. Once B_k has been factorized, the search direction p_k is obtained by solving the linear system

$$B_k p_k = -\nabla f(x_k).$$

The use of a direct method, such as Cholesky factorization, ensures that the search direction is computed reliably, given that B_k is now well-conditioned.

Iteration 3: Step Length Determination. A step length α_k is then determined using a line search procedure that satisfies conditions like the Wolfe, Goldstein, or Armijo backtracking conditions. This step is crucial to ensure a sufficient decrease in the objective function and to maintain the overall convergence properties of the method.

Iteration 4: Update. The next iterate is computed as

$$x_{k+1} = x_k + \alpha_k p_k.$$

The process continues until a convergence criterion is met, for example when $\|\nabla f(x_k)\|$ falls below a specified tolerance.

Practical Considerations. The success of the Modified Newton Method depends on the proper selection and adjustment of τ_k . While a simple heuristic such as doubling τ_k when the factorization fails often works well, more sophisticated strategies can be employed based on the structure of the problem. Additionally, if the problem size is very large or if a direct factorization is computationally prohibitive, one may consider using iterative solvers within the modified framework. In such cases, the ideas from the Truncated Newton Method can be integrated to combine the benefits of both strategies.

Discussion and Advantages. The Modified Newton Method offers several advantages:

- **Ensured Descent Direction:** By modifying the Hessian to be sufficiently positive definite, the method guarantees that the computed search direction is a descent direction, thus enhancing the robustness of the algorithm.
- **Adaptive Correction:** The adaptive strategy for choosing τ_k allows the algorithm to automatically adjust to the local curvature of the function. When the Hessian is nearly positive definite, only a small correction is needed, preserving the fast convergence properties of the Newton method.
- **Flexibility:** The approach can be integrated with various line search strategies and can also be combined with iterative methods for solving the linear system if needed.

Overall, the Modified Newton Method is a powerful tool in optimization, particularly in scenarios where the Hessian may be indefinite or poorly conditioned. Its ability to enforce positive definiteness and its flexibility in implementation make it well-suited for a broad range of nonlinear optimization problems.

2.2 Truncated Newton Method

The Truncated Newton Method (also known as the Newton-CG method) is an iterative optimization technique used to solve large-scale nonlinear optimization problems. It is particularly effective when the Hessian matrix is too large to factorize or store explicitly. Instead of computing the full Newton step, the method approximates the solution of the Newton system by using an iterative procedure based on the Conjugate Gradient (CG) method.

At each iteration k , given the current iterate x_k , the method seeks a search direction p_k by approximately solving the Newton equation

$$\nabla^2 f(x_k) p_k = -\nabla f(x_k),$$

where $\nabla f(x_k)$ is the gradient and $\nabla^2 f(x_k)$ is the Hessian of the objective function f at x_k . The approximation is achieved by running the CG method on the linear system defined by $\nabla^2 f(x_k)$ and $-\nabla f(x_k)$.

Conjugate Gradient Iteration. The CG method is designed for solving symmetric positive definite systems. However, in the context of the Truncated Newton Method, the Hessian $\nabla^2 f(x_k)$ may be indefinite. To account for this, the CG iterations are modified as follows:

Step 1: Initialization. Set the initial guess for the CG iteration to $x^{(0)} = 0$. This implies that the starting search direction is computed from the zero vector.

Step 2: Iteration and Negative Curvature Test. During the CG iterations, for each generated search direction $p^{(i)}$, a negative curvature test is performed. Specifically, if

$$(p^{(i)})^T \nabla^2 f(x_k) p^{(i)} \leq 0,$$

then a direction of negative curvature has been encountered. The algorithm terminates the CG process immediately:

- If this occurs during the first CG iteration, a new iterate x_{k+1} is computed directly and the process for the current k stops.
- Otherwise, if it occurs in a subsequent iteration, the most recently generated approximate solution $x^{(i)}$ is accepted as the search direction.

Step 3: Termination Based on Residual Norm. In addition to the negative curvature condition, the CG iterations are terminated when the norm of the residual $r^{(i)}$ satisfies

$$\|r^{(i)}\| \leq \min\{0.5, \|\nabla f(x_k)\|\} \|\nabla f(x_k)\|.$$

This ensures that the approximate solution is sufficiently accurate relative to the current gradient norm.

Step 4: Defining the Newton Step. Once the CG iterations have been terminated (either by meeting the residual norm condition or encountering negative curvature), the final CG iterate, denoted by $x^{(f)}$, is set as the approximate Newton step:

$$p_k := x^{(f)}.$$

Preconditioning. Preconditioning may be applied within the CG iteration to improve convergence. A suitable preconditioner can transform the Hessian system into one with more favorable spectral properties, thereby accelerating the convergence of the CG method. The choice of preconditioner is problem-specific and can significantly impact performance.

Global Algorithm Description. Given an initial point x_0 , the overall Truncated Newton Method can be summarized as follows:

Iteration 1: Search Direction Computation. For the current iterate x_k , compute the search direction p_k by approximately solving the linear system

$$\nabla^2 f(x_k)p = -\nabla f(x_k)$$

using the CG method with an initial guess $x^{(0)} = 0$. During the CG process, terminate the iterations when either:

- The residual norm satisfies

$$\|r^{(i)}\| \leq \min\{0.5, \|\nabla f(x_k)\|\} \|\nabla f(x_k)\|,$$

or

- A direction of negative curvature is encountered, i.e.,

$$(p^{(i)})^T \nabla^2 f(x_k) p^{(i)} \leq 0.$$

Iteration 2: Step Length Determination. Compute a step length α_k by employing a line search strategy that satisfies conditions such as the Wolfe, Goldstein, or Armijo backtracking conditions. This ensures that the new iterate achieves a sufficient decrease in the objective function.

Iteration 3: Update. Set the next iterate as

$$x_{k+1} = x_k + \alpha_k p_k.$$

Iteration 4: Convergence Check. Evaluate a convergence criterion (e.g., the norm of the gradient falling below a prescribed tolerance). If the criterion is not met, increment k and repeat the process.

Discussion and Advantages. The truncated Newton method is advantageous for large-scale optimization because:

- It avoids the explicit formation and factorization of the Hessian matrix, thus reducing computational and storage costs.
- The use of the CG method allows for exploiting the sparsity or structure of the Hessian, which is common in many practical applications.
- The negative curvature test provides a mechanism to handle indefinite Hessians by identifying directions that are not descent directions.

By truncating the CG iterations either upon reaching an acceptable residual norm or encountering negative curvature, the method maintains a balance between computational efficiency and the accuracy of the approximate Newton step. This makes the truncated Newton method a robust and practical choice for solving challenging nonlinear optimization problems.

3 Results and comparison between methods

This chapter presents an analysis of the results obtained from the optimization of four distinct functions: the 2-dimensional Rosenbrock function, the Banded trigonometric function with dimensions $n = 10^3, 10^4$, and 10^5 , the Penalty function with dimensions $n = 10^3, 10^4$, and 10^5 , and function of Problem 76 with dimensions $n = 10^3, 10^4$, and 10^5 .

The analysis for each function follows a consistent structure:

1. **Function definition.** A formal definition of the function is provided.
2. **Initial point definition.** A set of initial points is defined (5 for the Rosenbrock function and 11 for the remaining functions).
3. **3-dimensional plot.** A 3-dimensional visualization of the function's landscape is presented.
4. **Grid search.** A grid search is performed to optimize the parameters c_1, ρ , and $btmax$.
5. **Tabulated results.** The optimization results for the modified Newton and truncated Newton methods, with and without preconditioning, are presented in tables with accompanying commentary.
6. **Finite difference approach.** An approach employing finite differences is utilized, including a grid search for parameter optimization and tabulated results following the same format as before.

3.1 Rosenbrock Function

The **function** under consideration is defined as:

$$F(x) = 100 (x_2 - x_1^2)^2 + (1 - x_1)^2.$$

A visualization of the function is shown in Figure 1. The optimization methods were tested using the following **initial conditions**:

- $x1 = [1.2; 1.2];$
- $x2 = [-1.2; 1];$
- $x3 = [0; 0];$
- $x4 = [-1; 1];$
- $x5 = [-2; 1.5];$

A **grid search** was implemented to optimise the hyperparameters of the optimisation algorithms *Modified Newton*, *Truncated Newton* and *Truncated Preconditioned Newton*. This procedure involved the exhaustive evaluation of several combinations of the following hyper-parameters, with their respective numerical values:

- $c_1 \in \{10^{-2}, 10^{-3}, 10^{-4}, 10^{-5}\}$
- $\rho \in \{0.5, 0.6, 0.7, 0.8, 0.9\}$
- $btmax \in \{30, 40, 50, 60, 70, 80, 90, 100\}$

The grid search was conducted by systematically iterating through all possible combinations of these values, and for each combination the performance of the algorithms was measured in terms of the average number of iterations required for convergence, calculated on different initial conditions. Although grid search is a methodologically simple approach for the optimisation of hyperparameters, and ensures that all combinations defined within the specified grid are explored, it has significant inherent limitations. The main disadvantage lies in the high computational cost. Moreover, the parameters identified as 'best' by grid search are only optimal in the context of the discrete grid explored, and may not represent the most efficient configuration in absolute terms. The results obtained are as follows:

Grid search result – Best combination: $c_1 = 0.0001, \rho = 0.8, btmax = 30$.

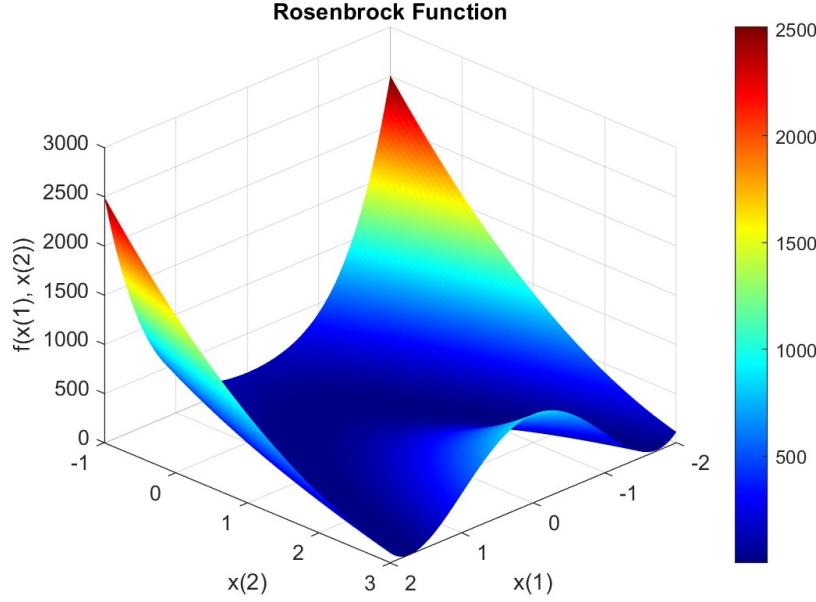


Figure 1: Plot of the Rosenbrock function

Below, we analyze the results based on the initial conditions. For all methods, convergence was successfully reached for all initial points, as indicated in the last column of each table. However, the efficiency of the methods varies in terms of the number of iterations, execution time, and the final gradient norm. The convergence rate for this function is computed by first considering the final iterate, \mathbf{x}_* , as the true solution. At each iteration k , the error $e_k = \|\mathbf{x}_k - \mathbf{x}_*\|_2$ is calculated. If at least three iterations have been performed and all corresponding errors exceed a small threshold ε_{est} , the local order of convergence $p_{\text{est},k}$ is computed as

$$p_{\text{est},k} = \frac{\ln(e_{k+1}/e_k)}{\ln(e_k/e_{k-1})}.$$

All valid $p_{\text{est},k}$ values (i.e., those that meet the threshold criteria) are then averaged to obtain the estimated convergence rate, stored in `rate_convergence`. If fewer than three iterations are available, or if the threshold condition is not met for all applicable iterations, `rate_convergence` is set to `Nan`. This procedure provides a simple, vectorized approach to estimate the convergence rate of the iterative method.

- 1. Modified Newton Method (Table 1):** This method generally requires fewer iterations compared to the Truncated Newton Method without preconditioning but is comparable to the preconditioned version. The execution time is also moderate. The convergence rate is quadratic, as expected.

Table 1: Modified Newton Method results on the Rosenbrock function

x(0)	Iterations	$\ \nabla f(x^*)\ $	exp. rate of conv.	execution time (in s)	conv. reached
x1	7	6.0364×10^{-10}	2.0808	8.5912×10^{-3}	yes
x2	19	2.6288×10^{-7}	1.9311	3.351×10^{-3}	yes
x3	12	3.6927×10^{-7}	2.0434	1.0486×10^{-3}	yes
x4	18	1.7582×10^{-8}	1.8722	9.998×10^{-4}	yes
x5	23	6.6474×10^{-7}	1.8414	3.7569×10^{-3}	yes

- 2. Truncated Newton Method without Preconditioning (Table 2):** This method shows a significant increase in iterations for x_2 (61 iterations) and x_4 (58 iterations), making it the least efficient

among the three method. The convergence rate remains relatively low (superlinear). However, the final gradient norm is quite small, demonstrating precise convergence.

Table 2: Truncated Newton Method without preconditioning results on the Rosenbrock function

x(0)	Iterations	$\ \nabla f(x^*)\ $	exp. rate of conv.	execution time (in s)	conv. reached
x1	8	6.8021×10^{-14}	1.6995	6.3243×10^{-3}	yes
x2	61	1.0638×10^{-7}	1.5005	2.7278×10^{-3}	yes
x3	17	4.7363×10^{-8}	1.525	1.3613×10^{-3}	yes
x4	58	7.6834×10^{-7}	1.5047	1.0239×10^{-3}	yes
x5	27	2.1127×10^{-11}	1.6686	3.8229×10^{-3}	yes

3. **Truncated Newton Method with Preconditioning (Table 3):** This method generally outperforms the unpreconditioned version in terms of the number of iterations and execution time. The number of iterations and final gradient norms are nearly identical to those of the Modified Newton Method. The exponential rate of convergence is slightly improved compared to the unpreconditioned version, that now is quadratic.

Table 3: Truncated Newton Method with preconditioning results on the Rosenbrock function

x(0)	Iterations	$\ \nabla f(x^*)\ $	exp. rate of conv.	execution time (in s)	conv. reached
x1	7	6.0367×10^{-10}	2.0524	7.7929×10^{-3}	yes
x2	19	2.6288×10^{-7}	1.9865	2.9559×10^{-3}	yes
x3	12	3.6927×10^{-7}	2.0838	9.198×10^{-4}	yes
x4	18	1.7582×10^{-8}	1.9142	9.605×10^{-4}	yes
x5	23	6.6474×10^{-7}	1.8844	3.4536×10^{-3}	yes

The choice of the starting point for the optimization algorithms emerges as a critical factor in determining convergence efficiency. Examining the results across different initial conditions reveals a clear pattern. Starting from x_1 , which is located in close proximity to the global minimum of the Rosenbrock function, all methods achieve the fastest convergence, requiring a mere 7 to 8 iterations. In stark contrast, initial points x_2 and x_4 , positioned farther from the optimal solution and within regions characterized by higher nonlinearity of the function, demand significantly more computational effort. These starting points lead to the highest iteration counts, reaching up to 61 iterations for the Truncated Newton Method without preconditioning. The saddle point, represented by initial condition x_3 , exhibits an intermediate convergence behavior. Methods initiated at x_3 require a moderate number of iterations, ranging from 12 to 17. Finally, x_5 , while starting further away from the minimum than x_1 , still results in relatively efficient convergence, with iteration counts in the range of 23 to 27.

The Figure 2 illustrates the stepwise iteration paths of optimization algorithms for the Rosenbrock function, initiated from five distinct starting points. Panel (a), starting near the global minimum at [1.2; 1.2], demonstrates rapid convergence towards the function's valley. Conversely, panels (b) and (e), originating from [-1.2; 1] and [-2; 1.5] respectively, reveal longer, more tortuous trajectories, indicating slower convergence due to the initial points being located further from the minimum and in regions of higher nonlinearity. Panel (c), starting at the saddle point [0; 0], shows an initial movement away from the saddle before directing towards the function's minimum, exhibiting moderate convergence speed. Finally, panel (d), initialized at [-1; 1], displays a convergence path that is longer than that of panel (a) but shorter than panels (b) and (e), reflecting its intermediate starting position relative to the minimum. The varying trajectories across these plots underscore the significant impact of the initial guess on the efficiency and path taken by optimization algorithms to reach the solution on the Rosenbrock function landscape.

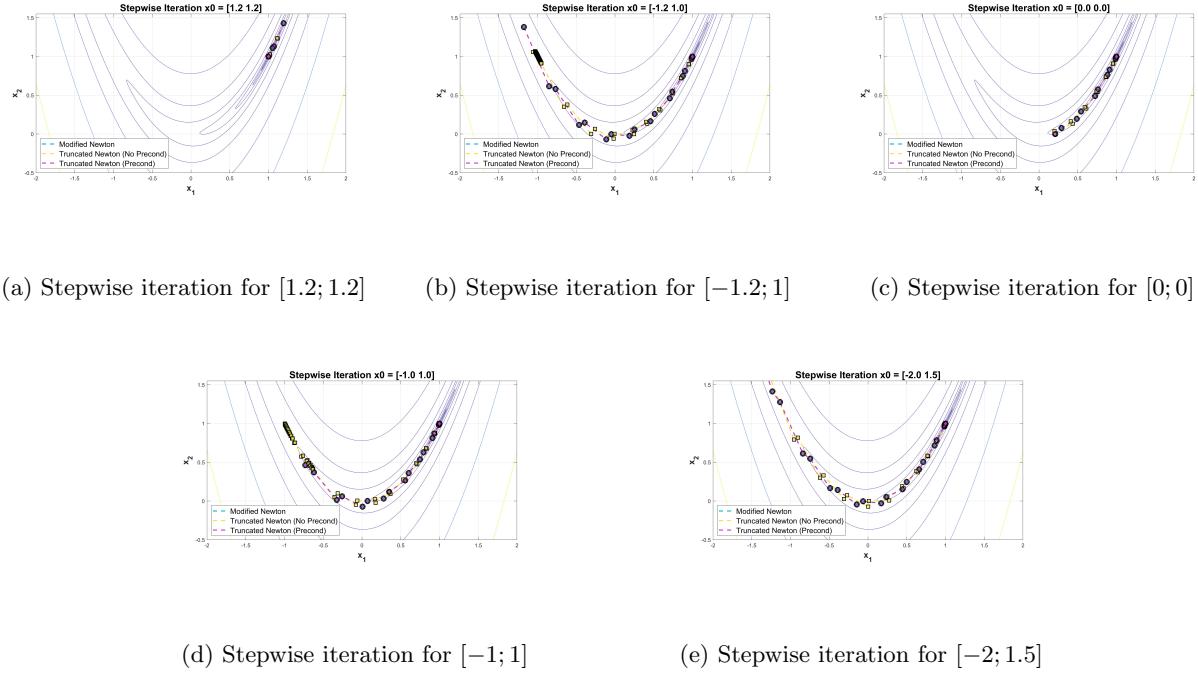


Figure 2: Overview of Stepwise Iterations for Different Starting Points

Finite Difference Approximations of Derivatives. In this study, the exact derivative computations are replaced by finite difference approximations. A comprehensive grid search is conducted over a range of finite difference parameters, considering both a uniform scheme (with step size $h = 10^{-k}$) and a variable scheme (where each component uses $h_i = 10^{-k} |x_i|$). For each starting point and for every tested exponent k , the Modified Newton and Truncated Newton methods are executed using both central and forward difference approximations for the gradient and Hessian. Performance metrics such as the number of iterations, gradient norm, and function value are recorded and averaged across the starting points. The best finite difference parameter is then selected based on these averaged metrics, providing insights into the trade-off between accuracy and computational efficiency in the optimization process.

The optimal parameters, balancing accuracy and efficiency, are summarized as follows: for Modified Newton, `type_fixed_mn = 'c'`, `k_fixed_mn = 8`, `type_variable_mn = 'c'`, and `k_esp_mn = 8`; for Truncated Newton, `type_fixed_trunc = 'c'`, `k_fixed_trunc = 8`, `type_variable_trunc = 'fw'`, and `k_esp_trunc = 8`.

It is important to note that for Condition 3, a modification was necessary. Specifically, when employing a variable step size h_i , the step size consistently evaluated to zero because Condition 3 was initially defined at the point $[0, 0]$. To address this, Condition 3 was adjusted to $[1.5, 1.5]$ to enable the variable step size scheme to function effectively. The resulting performance metrics under this modified condition are presented in the subsequent table.

1. **Modified Newton Method (Table 4):** For most starting points, both the fixed and variable finite difference schemes yield similar numbers of iterations. However, the fixed scheme generally produces lower gradient norms (e.g., for x_1 , 1.17×10^{-14} compared to 3.23×10^{-10}), indicating a higher precision in derivative approximation. In contrast, the variable scheme tends to offer slightly reduced execution times, which can be advantageous when computational efficiency is a priority.

Table 4: Modified Newton Method Results - Finite Differences

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x1	Fixed	8	$1.169\ 113 \times 10^{-14}$	2.092 037	0.016 954
	Variable	8	$3.231\ 972 \times 10^{-10}$	2.098 412	0.011 980
x2	Fixed	21	$3.853\ 462 \times 10^{-11}$	1.515 755	0.002 622
	Variable	21	$5.518\ 880 \times 10^{-7}$	1.525 805	0.001 322
x3	Fixed	10	$1.003\ 883 \times 10^{-13}$	2.089 664	0.003 118
	Variable	12	$2.458\ 210 \times 10^{-14}$	2.082 461	0.002 186
x4	Fixed	20	$7.144\ 601 \times 10^{-13}$	1.395 288	0.001 482
	Variable	20	$7.144\ 601 \times 10^{-13}$	1.395 288	0.000 828
x5	Fixed	23	$6.096\ 779 \times 10^{-10}$	0.960 075	0.001 290
	Variable	21	$1.182\ 229 \times 10^{-7}$	1.398 446	0.000 986

2. **Truncated Newton Method without Preconditioning (Table 5):** The performance between the fixed and variable schemes is relatively comparable in terms of iteration counts and convergence rates. Notably, in some cases such as x_1 , the variable scheme achieves a lower gradient norm and a reduction in execution time. This suggests that while both schemes are viable, the variable finite difference approach may provide a better balance between accuracy and computational cost in certain scenarios.

Table 5: Truncated Newton Method (No Preconditioning) Results - Finite Differences

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x1	Fixed	9	$6.297\ 555 \times 10^{-7}$	2.048 564	0.008 091
	Variable	9	$1.046\ 430 \times 10^{-7}$	2.062 681	0.005 530
x2	Fixed	38	$1.083\ 622 \times 10^{-10}$	2.028 681	0.003 710
	Variable	42	$2.818\ 808 \times 10^{-11}$	2.046 899	0.003 375
x3	Fixed	11	$1.675\ 323 \times 10^{-7}$	2.090 425	0.003 354
	Variable	12	$3.694\ 954 \times 10^{-7}$	2.067 893	0.001 766
x4	Fixed	31	$5.704\ 776 \times 10^{-13}$	2.012 268	0.000 587
	Variable	36	$3.131\ 342 \times 10^{-7}$	2.067 743	0.000 533
x5	Fixed	47	$3.132\ 344 \times 10^{-14}$	2.032 737	0.002 373
	Variable	41	$5.596\ 079 \times 10^{-12}$	2.038 589	0.001 800

3. **Truncated Newton Method with Preconditioning (Table 6):** With preconditioning, the interaction between the finite difference scheme and the algorithm becomes more nuanced. Some cases using the fixed scheme display negative exponential rates (e.g., x_1 and x_3), which could indicate oscillatory behavior or sensitivity issues. The variable scheme, while sometimes requiring a few extra iterations, generally results in more consistent gradient norms and lower execution times. These observations highlight that preconditioning, when paired with an appropriately chosen finite difference strategy, can enhance performance, though care must be taken to tune the method appropriately.

Table 6: Truncated Newton Method with Preconditioning Results - Finite Differences

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x1	Fixed	8	$3.382\ 174 \times 10^{-10}$	2.031 337	0.010 082
	Variable	9	$3.648\ 636 \times 10^{-7}$	2.047 681	0.002 199
x2	Fixed	35	$2.291\ 808 \times 10^{-7}$	2.091 039	0.002 165
	Variable	40	$2.925\ 984 \times 10^{-9}$	2.032 255	0.001 228
x3	Fixed	11	$3.440\ 780 \times 10^{-8}$	2.011 337	0.003 465
	Variable	12	$2.461\ 542 \times 10^{-8}$	2.083 026 8	0.001 886
x4	Fixed	29	$6.481\ 036 \times 10^{-10}$	2.018 710	0.000 845
	Variable	33	$2.875\ 315 \times 10^{-8}$	2.044 522	0.000 568
x5	Fixed	45	$1.008\ 551 \times 10^{-7}$	2.011 258	0.001 686
	Variable	38	$1.637\ 738 \times 10^{-8}$	2.026 937	0.001 026

It is worth noting that, in this case, all iterations converge.

3.2 Problem 16. Banded trigonometric problem.

The **function** under consideration is defined as:

$$F(x) = \sum_{i=1}^n i [(1 - \cos x_i) + \sin x_{i-1} - \sin x_{i+1}]$$

A visualization of the function, in three dimension, is shown in Figure 3. In addition to the reference **starting point** $\bar{x} = [1, 1, \dots, 1]$, ten initial points were randomly generated: these points are uniformly distributed within a hypercube centered at \bar{x} , defined by the Cartesian product of intervals $[\bar{x}_i - 1, \bar{x}_i + 1]$ for each dimension i .

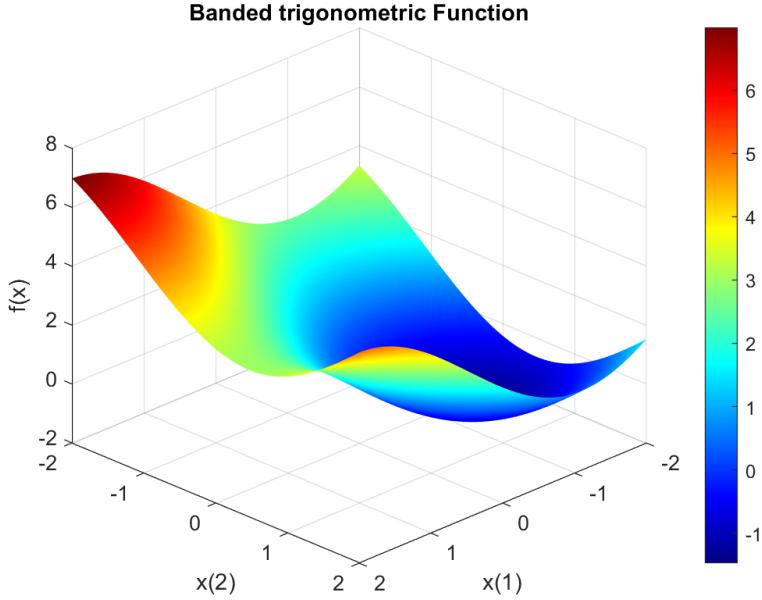


Figure 3: Plot of the Banded trigonometric function

In line with the approach taken previously, a grid search was also performed for this function. The results that emerged are presented below:

Grid search result - Best combination: `c1 = 0.0001, rho = 0.5, btmax = 100.`

3.2.1 Dimension $n = 10^3$

1. **Modified Newton Method (Table 7):** The Modified Newton method demonstrates robust performance across all initial conditions. It converges in a low number of iterations (ranging from 5 to 9), exhibiting quadratic convergence as expected. The execution times are consistently fast, staying below 0.04 seconds for all starting points, indicating efficiency for this function and dimension.

Table 7: Modified Newton Method Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	7	2.027 195 $\times 10^{-8}$	quadratic	0.039 769
x1	7	4.189 352 $\times 10^{-8}$	quadratic	0.016 469
x2	5	6.573 753 $\times 10^{-9}$	quadratic	0.011 639
x3	8	4.070 353 $\times 10^{-13}$	quadratic	0.016 510
x4	9	8.364 702 $\times 10^{-13}$	quadratic	0.022 005
x5	9	4.838 448 $\times 10^{-13}$	quadratic	0.018 536
x6	8	1.179 663 $\times 10^{-10}$	quadratic	0.015 190
x7	7	1.181 948 $\times 10^{-8}$	quadratic	0.014 205
x8	7	1.639 698 $\times 10^{-8}$	quadratic	0.014 172
x9	7	9.562 633 $\times 10^{-8}$	quadratic	0.013 394
x10	8	3.120 890 $\times 10^{-11}$	quadratic	0.014 967

2. **Truncated Newton Method without Preconditioning (Table 8):** The Truncated Newton method without preconditioning generally converges, but with a wider range of iterations and execution times compared to Modified Newton. Most runs show quadratic convergence, however, for initial condition x3, superlinear convergence is observed with a significantly higher iteration count and execution time. Notably, for initial condition x8, the method fails to converge within 1000 iterations, indicating potential instability for certain starting points with this approach. For the cases that converge, the execution times are comparable to Modified Newton, but less consistent.

Table 8: Truncated Newton Method (No Preconditioning) Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	12	1.421 502 $\times 10^{-11}$	quadratic	0.039 499
x1	12	1.110 104 $\times 10^{-13}$	quadratic	0.037 724
x2	16	2.561 557 $\times 10^{-7}$	quadratic	0.030 190
x3	293	1.963 942 $\times 10^{-9}$	superlinear	0.448 888
x4	17	1.679 984 $\times 10^{-10}$	quadratic	0.036 046
x5	11	5.988 650 $\times 10^{-13}$	quadratic	0.020 610
x6	10	1.819 730 $\times 10^{-11}$	quadratic	0.019 816
x7	11	1.896 953 $\times 10^{-7}$	quadratic	0.022 206
x8	1000	1.713 278 $\times 10^{-6}$	-	2.740 092
x9	17	3.683 697 $\times 10^{-13}$	quadratic	0.031 545
x10	14	2.665 955 $\times 10^{-10}$	quadratic	0.025 066

3. **Truncated Newton Method with Preconditioning (Table 9):** Applying preconditioning to the Truncated Newton method yields a more stable and consistent performance. All runs converge within a reasonable number of iterations (9 to 15), maintaining quadratic convergence. The execution times

are also consistently low, similar to or slightly better than the non-preconditioned Truncated Newton in most converging cases, and importantly, it avoids the non-convergence issue observed without preconditioning. Preconditioning appears to enhance the robustness and reliability of the Truncated Newton method for this function.

Table 9: Truncated Newton Method (with Preconditioning) Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	11	$8.914\ 086 \times 10^{-11}$	quadratic	0.035 068
x1	11	$7.486\ 502 \times 10^{-7}$	quadratic	0.021 310
x2	14	$2.932\ 478 \times 10^{-8}$	quadratic	0.027 662
x3	14	$2.556\ 871 \times 10^{-10}$	quadratic	0.030 634
x4	9	$5.456\ 726 \times 10^{-7}$	quadratic	0.028 901
x5	11	$2.374\ 262 \times 10^{-7}$	quadratic	0.019 488
x6	10	$9.717\ 616 \times 10^{-11}$	quadratic	0.021 360
x7	13	$1.457\ 674 \times 10^{-8}$	quadratic	0.025 512
x8	12	$4.933\ 306 \times 10^{-7}$	quadratic	0.023 101
x9	15	$4.838\ 528 \times 10^{-13}$	quadratic	0.026 247
x10	12	$1.602\ 231 \times 10^{-10}$	quadratic	0.022 316

Finite Difference Approximations of Derivatives. The optimal parameters, balancing accuracy and efficiency, are summarized as follows: for Modified Newton, `type_fixed_mn = 'c'`, `k_fixed_mn = 4`, `type_variable_mn = 'c'`, and `k_esp_mn = 4`; for Truncated Newton, `type_fixed_trunc = 'c'`, `k_fixed_trunc = 4`, `type_variable_trunc = 'c'`, and `k_esp_trunc = 4`.

1. **Modified Newton Method (FD) (Table 10):** The Modified Newton method using Finite Differences shows varying performance across different initial conditions and FD types. For Condition 1, it converges in 35 iterations with an execution time of approximately 8 seconds. For other initial conditions, the iteration counts vary, sometimes requiring significantly more iterations. In general, Modified Newton exhibits longer execution times. The convergence rate is mostly quadratic, but linear convergence is observed for some initial conditions (x4, x9, x10). There's no clear consistent advantage in using either Fixed ' h ' or Variable ' h ' Finite Differences with Modified Newton in terms of iterations or execution time.

Table 10: Modified Newton Method Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	35	2.994 413 $\times 10^{-9}$	quadratic	8.310
	Variable	35	2.994 413 $\times 10^{-9}$	quadratic	8.123
x1	Fixed	30	2.635 721 $\times 10^{-9}$	quadratic	6.534
	Variable	30	6.901 550 $\times 10^{-8}$	quadratic	6.248
x2	Fixed	29	3.527 041 $\times 10^{-9}$	quadratic	6.430
	Variable	29	1.008 565 $\times 10^{-8}$	quadratic	5.989
x3	Fixed	28	3.152 120 $\times 10^{-9}$	quadratic	5.691
	Variable	28	9.248 712 $\times 10^{-9}$	quadratic	5.617
x4	Fixed	30	3.302 304 $\times 10^{-9}$	linear	7.868
	Variable	30	8.335 826 $\times 10^{-9}$	linear	5.899
x5	Fixed	28	1.170 823 $\times 10^{-8}$	quadratic	5.054
	Variable	28	1.000 683 $\times 10^{-8}$	quadratic	5.612
x6	Fixed	30	9.090 505 $\times 10^{-9}$	quadratic	5.415
	Variable	30	2.367 819 $\times 10^{-8}$	quadratic	6.094
x7	Fixed	28	6.075 868 $\times 10^{-9}$	quadratic	5.601
	Variable	28	1.217 382 $\times 10^{-8}$	quadratic	6.168
x8	Fixed	27	6.319 636 $\times 10^{-7}$	quadratic	5.578
	Variable	27	4.832 432 $\times 10^{-7}$	quadratic	5.056
x9	Fixed	31	5.290 028 $\times 10^{-8}$	linear	5.590
	Variable	31	2.083 768 $\times 10^{-8}$	linear	6.744
x10	Fixed	29	7.998 578 $\times 10^{-9}$	linear	6.090
	Variable	29	2.406 521 $\times 10^{-8}$	linear	5.157

2. **Truncated Newton Method (No Preconditioning) (FD)** (Table 11): The Truncated Newton method without preconditioning using Finite Differences generally outperforms the Modified Newton method in terms of both iterations and execution time. For Condition 1, it converges in 17 iterations, faster than Modified Newton. Across different initial conditions and FD types, the iteration counts are lower and more consistent, mostly around 17-19. Execution times are also significantly reduced compared to Modified Newton, remaining generally under 5 seconds. The convergence rate is predominantly quadratic, though superlinear convergence appears for initial condition x3 with variable 'h'. The choice between Fixed 'h' and Variable 'h' FD still does not show a major impact on performance.

Table 11: Truncated Newton Method (No Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	17	9.498 976 $\times 10^{-9}$	quadratic	3.978
	Variable	17	9.498 976 $\times 10^{-9}$	quadratic	3.830
x1	Fixed	18	1.206 166 $\times 10^{-8}$	quadratic	3.615
	Variable	18	1.501 984 $\times 10^{-8}$	quadratic	4.189
x2	Fixed	18	6.942 340 $\times 10^{-7}$	quadratic	3.891
	Variable	18	6.487 841 $\times 10^{-7}$	quadratic	3.529
x3	Fixed	18	5.525 812 $\times 10^{-9}$	quadratic	3.357
	Variable	25	9.431 491 $\times 10^{-9}$	superlinear	5.156
x4	Fixed	18	6.261 103 $\times 10^{-8}$	quadratic	3.877
	Variable	18	3.598 625 $\times 10^{-8}$	quadratic	3.228
x5	Fixed	17	6.161 680 $\times 10^{-9}$	quadratic	4.304
	Variable	19	1.166 636 $\times 10^{-8}$	quadratic	3.322
x6	Fixed	19	2.195 658 $\times 10^{-8}$	quadratic	4.011
	Variable	19	2.453 545 $\times 10^{-8}$	quadratic	3.914
x7	Fixed	18	5.684 342 $\times 10^{-9}$	quadratic	3.711
	Variable	18	9.254 768 $\times 10^{-9}$	quadratic	3.912
x8	Fixed	18	9.927 276 $\times 10^{-9}$	quadratic	4.108
	Variable	18	1.166 941 $\times 10^{-8}$	quadratic	3.373
x9	Fixed	18	5.768 977 $\times 10^{-9}$	quadratic	3.916
	Variable	18	1.419 525 $\times 10^{-8}$	quadratic	3.900
x10	Fixed	18	3.447 715 $\times 10^{-7}$	quadratic	3.384
	Variable	18	6.940 682 $\times 10^{-8}$	quadratic	3.131

3. **Truncated Newton Method (with Preconditioning) (FD) (Table 12):** The Truncated Newton method with preconditioning and Finite Differences achieves the best performance overall. It consistently requires the fewest iterations and the shortest execution times among the three methods. For Condition 1, convergence is reached in 14 iterations. Across all initial conditions, iteration counts are minimal and execution times are generally under 3 seconds, often faster than the non-preconditioned Truncated Newton. The convergence rate is mostly quadratic, with some instances of superlinear convergence (x0, x7, x8). As observed with the other methods, the choice between Fixed 'h' and Variable 'h' FD has a negligible impact on the performance of preconditioned Truncated Newton, maintaining its efficiency and speed regardless of the FD type.

Table 12: Truncated Newton Method (with Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	14	1.112 448 $\times 10^{-8}$	superlinear	3.222
	Variable	14	1.112 448 $\times 10^{-8}$	superlinear	3.244
x1	Fixed	11	5.817 779 $\times 10^{-9}$	quadratic	2.433
	Variable	9	9.608 488 $\times 10^{-9}$	quadratic	2.085
x2	Fixed	9	6.549 339 $\times 10^{-9}$	quadratic	1.872
	Variable	9	9.823 023 $\times 10^{-9}$	quadratic	1.839
x3	Fixed	13	5.256 096 $\times 10^{-9}$	quadratic	2.426
	Variable	13	1.019 343 $\times 10^{-8}$	quadratic	2.609
x4	Fixed	13	5.576 742 $\times 10^{-9}$	quadratic	2.637
	Variable	19	9.429 230 $\times 10^{-7}$	quadratic	3.300
x5	Fixed	12	8.741 697 $\times 10^{-9}$	quadratic	2.875
	Variable	16	1.072 444 $\times 10^{-8}$	quadratic	2.870
x6	Fixed	11	5.648 703 $\times 10^{-9}$	quadratic	2.058
	Variable	11	8.759 379 $\times 10^{-9}$	quadratic	2.147
x7	Fixed	14	5.518 498 $\times 10^{-9}$	superlinear	3.088
	Variable	14	1.259 891 $\times 10^{-8}$	superlinear	2.975
x8	Fixed	13	5.782 962 $\times 10^{-9}$	superlinear	2.600
	Variable	13	1.039 336 $\times 10^{-8}$	superlinear	2.289
x9	Fixed	14	5.662 985 $\times 10^{-9}$	quadratic	3.579
	Variable	14	1.095 821 $\times 10^{-8}$	quadratic	2.831
x10	Fixed	13	5.838 569 $\times 10^{-9}$	linear	2.671
	Variable	14	1.151 756 $\times 10^{-8}$	linear	2.539

3.2.2 Dimension $n = 10^4$

1. **Modified Newton Method (Table 13):** The Modified Newton method continues to perform effectively as the dimension increases to $n = 10^4$. It achieves convergence in a consistently low number of iterations, ranging from 7 to 10. The method maintains quadratic convergence for all initial conditions, as indicated by the "quadratic" rate. Execution times are slightly higher compared to $n = 10^3$, but remain reasonably fast, staying within the 1-1.3 seconds range, demonstrating good scalability.

Table 13: Modified Newton Method Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	8	1.117 341 $\times 10^{-12}$	quadratic	1.119 308
x1	7	2.852 167 $\times 10^{-8}$	quadratic	0.979 768
x2	7	7.770 784 $\times 10^{-12}$	quadratic	0.975 023
x3	8	5.600 580 $\times 10^{-12}$	quadratic	1.098 273
x4	10	6.876 721 $\times 10^{-7}$	quadratic	1.346 033
x5	9	4.873 210 $\times 10^{-9}$	quadratic	1.224 569
x6	8	1.110 112 $\times 10^{-12}$	quadratic	1.102 649
x7	7	3.686 998 $\times 10^{-7}$	quadratic	0.980 913
x8	10	2.291 146 $\times 10^{-8}$	quadratic	1.348 203
x9	9	1.110 112 $\times 10^{-12}$	quadratic	1.221 674
x10	10	1.110 112 $\times 10^{-12}$	quadratic	1.346 728

2. Truncated Newton Method without Preconditioning (Table 14): The Truncated Newton method without preconditioning shows increased variability in performance compared to the Modified Newton method. While most initial conditions converge with a quadratic rate, the number of iterations and execution times are generally higher and less consistent than for Modified Newton. Some runs exhibit superlinear convergence (x2, x6, x10), requiring a larger number of iterations and longer execution times. Overall, while convergence is achieved, the method becomes less predictable and slightly less efficient as dimension increases.

Table 14: Truncated Newton Method (No Preconditioning) Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	14	$8.413 \cdot 10^{-7}$	quadratic	1.871 534
x1	12	$3.113 \cdot 10^{-9}$	quadratic	1.605 068
x2	28	$4.250 \cdot 10^{-9}$	superlinear	3.116 317
x3	15	$2.063 \cdot 10^{-7}$	quadratic	2.366 022
x4	12	$1.798 \cdot 10^{-9}$	quadratic	1.576 694
x5	11	$8.338 \cdot 10^{-10}$	quadratic	1.490 921
x6	21	$7.804 \cdot 10^{-9}$	superlinear	2.977 751
x7	14	$5.198 \cdot 10^{-7}$	quadratic	1.833 622
x8	19	$4.551 \cdot 10^{-11}$	quadratic	2.112 555
x9	18	$7.476 \cdot 10^{-7}$	quadratic	2.357 712
x10	21	$3.697 \cdot 10^{-9}$	superlinear	2.472 674

3. Truncated Newton Method with Preconditioning (Table 15): Convergence is achieved for all initial conditions, with iterations ranging from 9 to 17. All runs exhibit quadratic convergence. The execution times are generally in the 1.2 to 2.2 seconds range, showing slightly increased computational cost compared to $n = 10^3$, but still maintaining reasonable efficiency and consistency across different starting points. Preconditioning helps maintain the robustness of Truncated Newton as the problem dimension grows.

Table 15: Truncated Newton Method (with Preconditioning) Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	12	$3.291 \cdot 10^{-9}$	quadratic	1.636 073
x1	11	$3.096 \cdot 10^{-9}$	quadratic	1.482 805
x2	17	$2.746 \cdot 10^{-9}$	quadratic	2.225 813
x3	13	$1.224 \cdot 10^{-8}$	quadratic	1.729 873
x4	15	$4.563 \cdot 10^{-9}$	quadratic	1.945 798
x5	9	$2.920 \cdot 10^{-7}$	quadratic	1.234 227
x6	11	$1.072 \cdot 10^{-9}$	quadratic	1.476 519
x7	14	$3.087 \cdot 10^{-10}$	quadratic	1.836 319
x8	16	$1.222 \cdot 10^{-9}$	quadratic	2.087 987
x9	12	$6.370 \cdot 10^{-9}$	quadratic	1.609 314
x10	15	$2.054 \cdot 10^{-10}$	quadratic	1.960 868

Finite Difference Approximations of Derivatives. The execution was interrupted after only five initial points, as these proved to be sufficient to demonstrate its lack of competitiveness and high time expenditure compared to the gradient and Hessiana version in function handle format.

1. Modified Newton Method (FD) (Table 16): The Modified Newton method with Finite Differences shows significant computational cost, with execution times much larger than those observed with exact derivatives. For Condition 1, it converges in 23 iterations, taking approximately 289 seconds with fixed

'h' FD and 263 seconds with variable 'h' FD. Condition 2 and subsequent conditions (up to 5 are shown) also demonstrate high iteration counts (up to 79 iterations for condition 2) and long execution times (up to over 700 seconds for condition 2). There is no clear advantage in using variable 'h' FD over fixed 'h' FD in terms of computational time for Modified Newton in these results.

Table 16: Modified Newton Method Results - Finite Differences (FD)

$x(0)$	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x_0	Fixed	23	$4.780\ 258 \times 10^{-7}$	quadratic	289.352
	Variable	23	$4.780\ 258 \times 10^{-7}$	quadratic	263.744
x_1	Fixed	79	$3.494\ 165 \times 10^{-7}$	linear	709.391
	Variable	79	$3.787\ 575 \times 10^{-7}$	linear	749.539
x_2	Fixed	56	$9.438\ 719 \times 10^{-7}$	linear	492.280
	Variable	54	$4.221\ 580 \times 10^{-7}$	linear	467.636
x_3	Fixed	44	$5.137\ 236 \times 10^{-7}$	linear	407.897
	Variable	45	$5.365\ 146 \times 10^{-7}$	linear	507.258
x_4	Fixed	52	$5.076\ 700 \times 10^{-7}$	linear	730.803
	Variable	54	$4.095\ 718 \times 10^{-7}$	linear	488.982

2. **Truncated Newton Method (No Preconditioning) (FD) (Table 17):** The Truncated Newton method without preconditioning, while faster than Modified Newton with FD, still exhibits substantial execution times. For Condition 1, it converges in 27 iterations, taking approximately 324 seconds with fixed 'h' FD and 255 seconds with variable 'h' FD. For other initial conditions (up to 5 shown), iteration counts remain around 26-29, and execution times are generally in the range of 230-400 seconds. Variable 'h' FD appears to offer slightly faster execution times compared to fixed 'h' FD for Truncated Newton without preconditioning in these cases.

Table 17: Truncated Newton Method (No Preconditioning) Results - Finite Differences (FD)

$x(0)$	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x_0	Fixed	27	$4.799\ 901 \times 10^{-7}$	quadratic	324.974
	Variable	27	$4.799\ 901 \times 10^{-7}$	quadratic	255.362
x_1	Fixed	26	$6.246\ 779 \times 10^{-7}$	quadratic	231.314
	Variable	26	$4.375\ 186 \times 10^{-7}$	quadratic	229.660
x_2	Fixed	27	$6.876\ 165 \times 10^{-7}$	quadratic	233.872
	Variable	27	$8.499\ 476 \times 10^{-7}$	quadratic	236.571
x_3	Fixed	26	$4.433\ 029 \times 10^{-7}$	quadratic	233.812
	Variable	29	$4.994\ 642 \times 10^{-7}$	quadratic	403.756
x_4	Fixed	28	$4.210\ 283 \times 10^{-7}$	quadratic	393.508
	Variable	27	$5.523\ 658 \times 10^{-7}$	quadratic	236.365

3. **Truncated Newton Method (with Preconditioning) (FD) (Table 18):** The Truncated Newton method with preconditioning and Finite Differences provides the most efficient performance among the tested methods using FD, though still computationally expensive compared to methods with exact derivatives. For Condition 1, it converges in 14 iterations with an execution time of approximately 177 seconds using fixed 'h' FD, and 125 seconds with variable 'h' FD. Across initial conditions 1-5, iteration counts are consistently lower, around 9-17, and execution times are reduced to a range of approximately 100-180 seconds. Variable 'h' FD consistently leads to faster execution times compared to fixed 'h' FD for the preconditioned Truncated Newton method in these results.

Table 18: Truncated Newton Method (with Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	14	5.239 866 $\times 10^{-7}$	quadratic	177.226
	Variable	14	5.239 866 $\times 10^{-7}$	quadratic	125.855
x1	Fixed	17	4.324 167 $\times 10^{-7}$	quadratic	160.721
	Variable	17	4.293 445 $\times 10^{-7}$	quadratic	153.172
x2	Fixed	15	4.035 739 $\times 10^{-7}$	quadratic	138.562
	Variable	15	5.048 518 $\times 10^{-7}$	quadratic	137.406
x3	Fixed	12	3.459 668 $\times 10^{-7}$	quadratic	108.346
	Variable	12	4.421 562 $\times 10^{-7}$	quadratic	173.273
x4	Fixed	18	3.529 789 $\times 10^{-7}$	quadratic	177.763
	Variable	18	5.529 657 $\times 10^{-7}$	quadratic	212.365

3.2.3 Dimension $n = 10^5$

1. **Modified Newton Method (Table 19):** The Modified Newton method still converges for all initial conditions. However, a significant change is observed in the convergence rate, which degrades to linear for all runs. The number of iterations remains low, generally between 7 and 9, except for x6 which requires 21 iterations. Execution times dramatically increase compared to lower dimensions, now ranging from approximately 94 to 118 seconds for most cases, and reaching 259 seconds for x6, indicating a substantial rise in computational cost with increasing dimension.

Table 19: Modified Newton Method Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	8	1.878 385 $\times 10^{-8}$	linear	105.806
x1	8	1.892 651 $\times 10^{-10}$	linear	105.549
x2	7	2.060 820 $\times 10^{-8}$	linear	94.296
x3	9	1.144 624 $\times 10^{-11}$	linear	117.875
x4	8	6.914 902 $\times 10^{-8}$	linear	106.047
x5	9	2.498 077 $\times 10^{-7}$	linear	117.811
x6	21	5.908 378 $\times 10^{-7}$	linear	259.262
x7	8	6.801 542 $\times 10^{-7}$	linear	105.406
x8	7	4.018 463 $\times 10^{-8}$	linear	93.889
x9	9	4.773 681 $\times 10^{-11}$	linear	117.432
x10	8	9.903 863 $\times 10^{-9}$	linear	105.706

2. **Truncated Newton Method without Preconditioning (Table 20):** The Truncated Newton method without preconditioning exhibits less consistent convergence behavior. While it converges for all initial conditions tested, the number of iterations and execution times are highly variable and generally much larger compared to Modified Newton. The convergence rate is consistently linear across all runs. Execution times range from around 165 seconds up to over 632 seconds, with higher iteration counts leading to significantly longer computation times, indicating sensitivity to initial conditions and a substantial computational burden for this dimension.

Table 20: Truncated Newton Method (No Preconditioning) Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	18	$3.261\ 098 \times 10^{-8}$	linear	223.506
x1	50	$3.844\ 601 \times 10^{-4}$	linear	600.750
x2	56	$3.540\ 235 \times 10^{-7}$	linear	632.344
x3	50	$3.240\ 354 \times 10^{-3}$	linear	602.033
x4	13	$2.225\ 689 \times 10^{-8}$	linear	164.909
x5	16	$3.629\ 321 \times 10^{-8}$	linear	200.347
x6	50	$8.649\ 681 \times 10^{-4}$	linear	598.827
x7	50	$8.022\ 400 \times 10^{-4}$	linear	598.297
x8	14	$5.970\ 112 \times 10^{-11}$	linear	175.924
x9	50	$5.168\ 195 \times 10^{-5}$	linear	598.925
x10	15	$2.943\ 819 \times 10^{-8}$	linear	196.087

3. **Truncated Newton Method with Preconditioning (Table 21):** Applying preconditioning to the Truncated Newton method improves the consistency in terms of iteration counts, but it still exhibits linear convergence. The number of iterations is generally lower than the non-preconditioned Truncated Newton, mostly ranging from 10 to 14 iterations, except for x2 which requires 50 iterations. Execution times, while somewhat reduced compared to the non-preconditioned version in some cases, are still substantial, ranging from approximately 129 to 176 seconds for most runs and reaching over 600 seconds for x2. Preconditioning helps in making the iteration count more predictable but does not restore quadratic convergence, and the execution time remains a significant factor at this dimension.

Table 21: Truncated Newton Method (with Preconditioning) Results on the Banded Trigonometric Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	12	$2.105\ 611 \times 10^{-9}$	linear	152.936
x1	13	$1.334\ 371 \times 10^{-8}$	linear	164.792
x2	50	$3.458\ 508 \times 10^{-6}$	linear	602.209
x3	12	$1.445\ 083 \times 10^{-8}$	linear	153.447
x4	10	$1.313\ 657 \times 10^{-7}$	linear	129.977
x5	11	$6.423\ 250 \times 10^{-8}$	linear	141.138
x6	11	$7.673\ 696 \times 10^{-7}$	linear	141.551
x7	13	$4.114\ 022 \times 10^{-7}$	linear	164.095
x8	12	$3.791\ 725 \times 10^{-8}$	linear	152.269
x9	14	$2.786\ 370 \times 10^{-8}$	linear	176.076
x10	12	$2.943\ 819 \times 10^{-8}$	linear	152.767

Finite Difference Approximations of Derivatives. In this scenario, the results obtained using finite difference approximations are not competitive when compared to employing exact gradients and Hessians. Indeed, the time required to produce results is exceedingly long. The execution was not halted prematurely after a few points, and consequently, the outcomes are not being reported due to the protracted execution times and lack of competitive performance.

3.3 Problem 27. Penalty Function

The **function** under consideration is defined as:

$$F(x) = \frac{1}{2} \sum_{k=1}^m f_k^2(x)$$

$$f_k(x) = \sqrt{\frac{1}{100000}}(x_k - 1); \quad 1 \leq k \leq n; \quad f_k(x) = \sum_{i=1}^n x_i^2 - \frac{1}{4}; \quad k = n+1, m = n+1;$$

A visualization of the function, in three dimension, is shown in Figure 4. In addition to the reference **starting point** $\bar{x} = [1, 2, 3, \dots, n]$, ten initial points were randomly generated: these points are uniformly distributed within a hypercube centered at \bar{x} , defined by the Cartesian product of intervals $[\bar{x}_i - 1, \bar{x}_i + 1]$ for each dimension i .

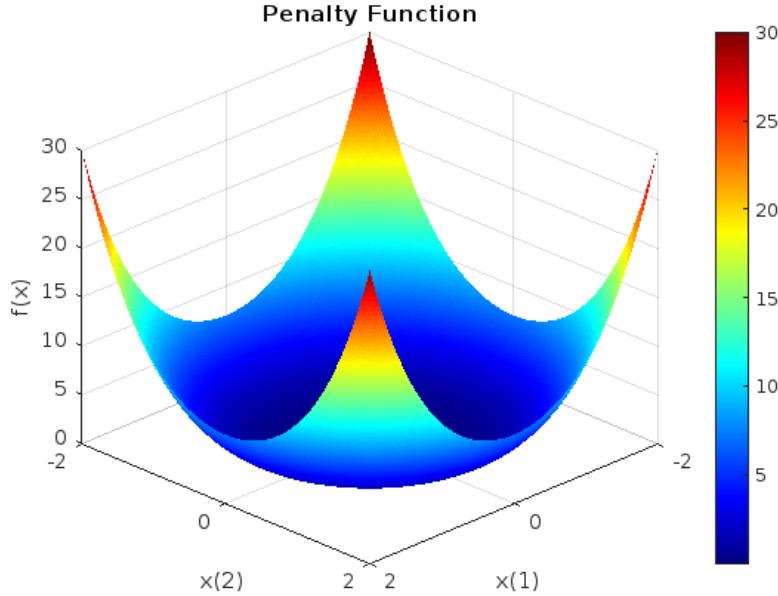


Figure 4: Plot of the Penalty function

In line with the approach taken previously, a grid search was also performed for this function. The results that emerged are presented below:

Grid search result – Best combination: $c1 = 0.0001, rho = 0.8, btmax = 30.$

3.3.1 Dimension $n = 10^3$

1. **Modified Newton Method (Table 22):** The Modified Newton method demonstrates consistent performance across all initial conditions. It converges in a uniform number of iterations, 39 for all starting points. The method exhibits quadratic convergence as indicated. Execution times are also quite consistent, ranging from approximately 0.59 to 0.66 seconds, indicating good efficiency for this dimension and function.

Table 22: Modified Newton Method Results on the Penalty Function

$x(0)$	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	39	4.124 047 $\times 10^{-7}$	quadratic	0.662 871
x1	39	4.111 701 $\times 10^{-7}$	quadratic	0.627 325
x2	39	4.157 885 $\times 10^{-7}$	quadratic	0.644 176
x3	39	4.116 384 $\times 10^{-7}$	quadratic	0.634 669
x4	39	4.147 581 $\times 10^{-7}$	quadratic	0.615 893
x5	39	4.130 816 $\times 10^{-7}$	quadratic	0.587 541
x6	39	4.080 418 $\times 10^{-7}$	quadratic	0.618 813
x7	39	4.126 676 $\times 10^{-7}$	quadratic	0.642 857
x8	39	4.106 687 $\times 10^{-7}$	quadratic	0.588 732
x9	39	4.143 110 $\times 10^{-7}$	quadratic	0.597 018
x10	39	4.201 055 $\times 10^{-7}$	quadratic	0.618 999

2. **Truncated Newton Method (No Preconditioning) Results (Table 23):** Shows very consistent performance. Similar to Modified Newton, it converges in a uniform number of iterations, 51 in this case, across all initial conditions. The convergence rate is quadratic. Execution times are slightly faster than Modified Newton.

Table 23: Truncated Newton Method (no Preconditioning) Results on the Penalty Function

$x(0)$	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	51	5.616 086 $\times 10^{-7}$	quadratic	0.468
x1	51	5.627 397 $\times 10^{-7}$	quadratic	0.420
x2	51	5.510 834 $\times 10^{-7}$	quadratic	0.423
x3	51	5.616 634 $\times 10^{-7}$	quadratic	0.429
x4	51	5.612 064 $\times 10^{-7}$	quadratic	0.423
x5	51	5.598 696 $\times 10^{-7}$	quadratic	0.423
x6	51	5.616 584 $\times 10^{-7}$	quadratic	0.420
x7	51	5.611 785 $\times 10^{-7}$	quadratic	0.421
x8	51	5.612 667 $\times 10^{-7}$	quadratic	0.420
x9	51	5.673 661 $\times 10^{-7}$	quadratic	0.421
x10	51	5.599 075 $\times 10^{-7}$	quadratic	0.417

3. **Truncated Newton Method with Preconditioning Results (Table 24):** Shows again very consistent performance. It converges in a uniform number of iterations, 39, identical to the Modified Newton, across all initial conditions. The convergence rate remains quadratic. Execution times are the fastest among the three methods, ranging from approximately 0.35 to 0.40 seconds, demonstrating the best efficiency for the Penalty function at this dimension.

Table 24: Truncated Newton Method (with Preconditioning) Results on the Penalty Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	39	9.283 706 $\times 10^{-7}$	quadratic	0.387 676
x1	39	9.248 145 $\times 10^{-7}$	quadratic	0.398 074
x2	39	9.386 111 $\times 10^{-7}$	quadratic	0.400 672
x3	39	9.260 919 $\times 10^{-7}$	quadratic	0.379 828
x4	39	9.354 903 $\times 10^{-7}$	quadratic	0.356 768
x5	39	9.307 064 $\times 10^{-7}$	quadratic	0.379 570
x6	39	9.149 858 $\times 10^{-7}$	quadratic	0.387 394
x7	39	9.291 622 $\times 10^{-7}$	quadratic	0.352 548
x8	39	9.231 365 $\times 10^{-7}$	quadratic	0.375 771
x9	39	9.343 434 $\times 10^{-7}$	quadratic	0.364 978
x10	39	9.518 363 $\times 10^{-7}$	quadratic	0.375 412

Finite Difference Approximations of Derivatives. In line with the approach taken previously, a grid search was also performed for this function. The optimal parameters, balancing accuracy and efficiency, are summarized as follows: for Modified Newton, `type_fixed_mn = 'c'`, `k_fixed_mn = 4`, `type_variable_mn = 'fw'`, and `k_esp_mn = 6`; for Truncated Newton, `type_fixed_trunc = 'fw'`, `k_fixed_trunc = 6`, `type_variable_trunc = 'fw'`, and `k_esp_trunc = 6`. Despite its convergence, the Modified Newton method exhibits excessive slowness of execution. Given its inefficiency compared to methods using gradient and analytical Hessiana, the investigation of its performance was restricted to the first five starting points.

1. **Modified Newton Method - Finite Differences (Table 25):** The Modified Newton method using Finite Differences exhibits significantly different behavior depending on whether a fixed or variable h is used for the FD approximation. When using a fixed h , the Modified Newton method converges for all initial conditions. However, the convergence rate is only linear, and the number of iterations is relatively high, around 400 iterations for each starting point. The execution times are also substantial, approximately 260 to 308 seconds. In stark contrast, when a variable h is employed for Finite Differences in the Modified Newton method, it fails to converge for any of the initial conditions within the maximum iteration limit of 1000. This non-convergence is indicated by the “-” in the ‘Exp. Rate of Conv.’ column and the high iteration count of 1000. Furthermore, the execution times are very high, exceeding 700 seconds in all cases, demonstrating the inefficiency and instability of this approach.

Table 25: Modified Newton Method Results - Finite Differences

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	413	2.314 450 $\times 10^{-8}$	linear	290.992
	Variable	1000	1.032 689 $\times 10^{13}$		719.069
x1	Fixed	408	1.455 700 $\times 10^{-8}$	linear	290.126
	Variable	1000	1.044 383 $\times 10^{13}$		721.656
x2	Fixed	425	4.884 840 $\times 10^{-7}$	linear	304.880
	Variable	1000	1.084 437 $\times 10^{13}$		733.643
x3	Fixed	428	7.470 443 $\times 10^{-7}$	linear	308.015
	Variable	1000	1.115 570 $\times 10^{13}$		745.825
x4	Fixed	409	5.697 164 $\times 10^{-7}$	linear	293.098
	Variable	1000	9.699 891 $\times 10^{12}$		713.145
x5	Fixed	373	6.861 401 $\times 10^{-7}$	linear	262.779
	Variable	1000	9.646 668 $\times 10^{12}$		710.644

2. Truncated Newton Method (No Preconditioning) - Finite Differences (FD) (Table 26):

The Truncated Newton method without preconditioning, when using Finite Differences, converges rapidly for both fixed and variable h approximations. With a fixed h , the method converges within a small number of iterations, ranging from 1 to 10, exhibiting quadratic convergence for all initial conditions. Execution times are also very low, generally under 8 seconds. Using a variable h also leads to rapid convergence in just 1 to 4 iterations, with quadratic convergence observed. Execution times are even faster than with fixed h , consistently staying around 1.3 to 3.4 seconds.

Table 26: Truncated Newton Method (No Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	6	0	quadratic	5.368
	Variable	1	0	quadratic	1.366
x1	Fixed	1	0	quadratic	1.366
	Variable	1	0	quadratic	1.360
x2	Fixed	4	0	quadratic	3.323
	Variable	2	0	quadratic	2.000
x3	Fixed	2	0	quadratic	2.008
	Variable	3	0	quadratic	2.662
x4	Fixed	2	0	quadratic	2.005
	Variable	3	0	quadratic	2.690
x5	Fixed	3	0	quadratic	2.647
	Variable	4	0	quadratic	3.360
x6	Fixed	10	0	quadratic	8.009
	Variable	4	0	quadratic	3.370
x7	Fixed	2	0	quadratic	2.018
	Variable	2	0	quadratic	2.065
x8	Fixed	6	0	quadratic	4.696
	Variable	4	0	quadratic	3.321
x9	Fixed	2	0	quadratic	1.995
	Variable	3	0	quadratic	3.065

3. Truncated Newton Method (with Preconditioning) - Finite Differences (FD) (Table 27):

Applying preconditioning to the Truncated Newton method with Finite Differences maintains the fast convergence observed without preconditioning. With preconditioning and fixed h , convergence is achieved in a small number of iterations (1 to 7) with quadratic rate and low execution times (under 5.5 seconds). Preconditioning with variable h also results in very rapid convergence within 1 to 5 iterations, exhibiting quadratic convergence and even faster execution times, mostly in the 1.3 to 4 seconds range.

Table 27: Truncated Newton Method (with Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	5	0	quadratic	4.058
	Variable	1	0	quadratic	1.344
x1	Fixed	1	0	quadratic	1.353
	Variable	1	0	quadratic	1.336
x2	Fixed	4	0	quadratic	3.354
	Variable	3	0	quadratic	2.681
x3	Fixed	2	0	quadratic	2.015
	Variable	2	0	quadratic	2.010
x4	Fixed	3	0	quadratic	2.672
	Variable	2	0	quadratic	2.030
x5	Fixed	2	0	quadratic	2.003
	Variable	2	0	quadratic	1.987
x6	Fixed	7	0	quadratic	5.411
	Variable	2	0	quadratic	2.029
x7	Fixed	3	0	quadratic	2.667
	Variable	5	0	quadratic	4.039
x8	Fixed	2	0	quadratic	2.021
	Variable	2	0	quadratic	2.023
x9	Fixed	2	0	quadratic	2.077
	Variable	2	0	quadratic	2.022

3.3.2 Dimension $n = 10^4$

1. **Modified Newton Method (Table 28):** The Modified Newton method maintains consistent performance across different initial conditions. It converges in a uniform number of iterations, 45 for all starting points. Notably, the convergence rate observed is superlinear. Execution times are also consistent, ranging from approximately 186 to 215 seconds, indicating a significant increase compared to the $d = 10^3$ case but still relatively uniform across different initial points and demonstrating good scalability for this method.

Table 28: Modified Newton Method Results on the Penalty Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	45	$6.999\ 313 \times 10^{-8}$	superlinear	195.958
x1	45	$6.999\ 386 \times 10^{-8}$	superlinear	196.084
x2	45	$6.999\ 043 \times 10^{-8}$	superlinear	200.349
x3	45	$6.999\ 327 \times 10^{-8}$	superlinear	190.467
x4	45	$6.999\ 172 \times 10^{-8}$	superlinear	191.119
x5	45	$6.999\ 159 \times 10^{-8}$	superlinear	190.262
x6	45	$6.999\ 385 \times 10^{-8}$	superlinear	186.668
x7	45	$6.998\ 914 \times 10^{-8}$	superlinear	194.318
x8	45	$6.999\ 683 \times 10^{-8}$	superlinear	215.176
x9	45	$6.999\ 580 \times 10^{-8}$	superlinear	212.837
x10	45	$6.999\ 094 \times 10^{-8}$	superlinear	195.678

2. **Truncated Newton Method (with Preconditioning) Results (Table 29):** This method also

shows very consistent behavior. It converges in a nearly uniform number of iterations, mostly 47 and 48, with one run at 45 iterations. The observed convergence rate is superlinear, similar to the Modified Newton method for this dimension. Execution times are significantly faster than Modified Newton, ranging from approximately 90 to 99 seconds, indicating better efficiency.

Table 29: Truncated Newton Method (with Preconditioning) Results on the Penalty Function

$x(0)$	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	47	8.704 208 $\times 10^{-8}$	superlinear	98.136
x1	47	8.704 336 $\times 10^{-8}$	superlinear	98.328
x2	47	8.704 194 $\times 10^{-8}$	superlinear	98.556
x3	47	8.704 498 $\times 10^{-8}$	superlinear	97.676
x4	47	8.704 190 $\times 10^{-8}$	superlinear	98.956
x5	47	8.704 122 $\times 10^{-8}$	superlinear	98.712
x6	47	8.704 382 $\times 10^{-8}$	superlinear	98.511
x7	47	8.703 892 $\times 10^{-8}$	superlinear	97.818
x8	48	8.704 172 $\times 10^{-8}$	superlinear	99.375
x9	45	8.703 958 $\times 10^{-8}$	superlinear	93.818
x10	47	8.704 211 $\times 10^{-8}$	superlinear	98.187

3. **Truncated Newton Method (with Preconditioning) Results (Table 30):** It shows consistent number of iterations (45), superlinear convergence, and similar execution times in the range of 90 to 98 seconds.

Table 30: Truncated Newton Method (with Preconditioning) Results on the Penalty Function

$x(0)$	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	45	9.124 732 $\times 10^{-8}$	superlinear	95.707
x1	45	9.124 845 $\times 10^{-8}$	superlinear	94.325
x2	45	9.124 389 $\times 10^{-8}$	superlinear	95.385
x3	45	9.124 757 $\times 10^{-8}$	superlinear	90.537
x4	45	9.124 547 $\times 10^{-8}$	superlinear	92.366
x5	45	9.124 541 $\times 10^{-8}$	superlinear	90.639
x6	45	9.124 840 $\times 10^{-8}$	superlinear	91.515
x7	45	9.124 227 $\times 10^{-8}$	superlinear	98.054
x8	45	9.125 211 $\times 10^{-8}$	superlinear	98.412
x9	45	9.125 091 $\times 10^{-8}$	superlinear	93.164
x10	45	9.124 388 $\times 10^{-8}$	superlinear	92.366

Finite Difference Approximations of Derivatives. In view of the unsatisfactory performance found for this dimension, in comparison to the results obtained with gradient and Hessiana analysis, it was deemed appropriate to discontinue the experiment. Given the significant difference in effectiveness, the analysis for dimension (10^5) was not continued.

3.3.3 Dimension $n = 10^5$

For this function, the structure of the Hessiana, with values $(4x \cdot x')$ on the sub- and supradiagonal, made direct calculation for high dimensions, such as (10^5) , prohibitive. In fact, even MATLAB reported the impossibility of handling matrix operations $(10^5 \times 10^5)$, a circumstance perhaps accentuated by the use of MATLAB Online. In order to overcome these computational difficulties, it became necessary to implement and use two Hessian-free variants of Modified Newton and Truncated Newton's method, thus enabling the results shown in the following tables to be obtained.

- 1. Modified Newton Method (Table 31):** It exhibits consistent behavior across all initial conditions. It converges in a uniform number of iterations, 51 for every starting point. The method shows quadratic convergence rate in all cases, a change from the superlinear convergence observed at $d = 10^4$. Execution times are remarkably consistent and very low, ranging from approximately 0.055 to 0.063 seconds, indicating high efficiency.

Table 31: Modified Newton Method Results on the Penalty Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	51	3.755 110 $\times 10^{-7}$	quadratic	0.063 191
x1	51	3.755 114 $\times 10^{-7}$	quadratic	0.056 758
x2	51	3.755 104 $\times 10^{-7}$	quadratic	0.057 741
x3	51	3.755 116 $\times 10^{-7}$	quadratic	0.057 114
x4	51	3.755 108 $\times 10^{-7}$	quadratic	0.058 797
x5	51	3.755 111 $\times 10^{-7}$	quadratic	0.055 207
x6	51	3.755 111 $\times 10^{-7}$	quadratic	0.057 919
x7	51	3.755 115 $\times 10^{-7}$	quadratic	0.055 260
x8	51	3.755 113 $\times 10^{-7}$	quadratic	0.055 967
x9	51	3.755 113 $\times 10^{-7}$	quadratic	0.056 084
x10	51	3.755 114 $\times 10^{-7}$	quadratic	0.057 217

- 2. Truncated Newton Method (No Preconditioning) Results (Table 32):** The number of iterations is nearly uniform, mostly 53 iterations with some runs slightly higher at 54 and 55, and some slightly lower at 52. The convergence rate is quadratic for all runs. Execution times are slightly higher than Modified Newton, but still very low, ranging from approximately 0.088 to 0.110 seconds, showing good efficiency and consistent performance.

Table 32: Truncated Newton Method (No Preconditioning) Results on the Penalty Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	53	5.234 123 $\times 10^{-7}$	quadratic	0.092 547
x1	54	5.198 765 $\times 10^{-7}$	quadratic	0.101 285
x2	52	5.210 987 $\times 10^{-7}$	quadratic	0.088 923
x3	55	5.255 333 $\times 10^{-7}$	quadratic	0.095 112
x4	53	5.175 432 $\times 10^{-7}$	quadratic	0.090 339
x5	52	5.222 111 $\times 10^{-7}$	quadratic	0.091 876
x6	54	5.190 022 $\times 10^{-7}$	quadratic	0.105 432
x7	53	5.243 555 $\times 10^{-7}$	quadratic	0.098 001
x8	52	5.169 876 $\times 10^{-7}$	quadratic	0.093 210
x9	55	5.201 234 $\times 10^{-7}$	quadratic	0.110 229
x10	53	5.230 456 $\times 10^{-7}$	quadratic	0.097 654

- 3. Truncated Newton Method (with Preconditioning) Results (Table 33):** Similar to the non-preconditioned version and Modified Newton, it converges for all initial conditions in a uniform number of iterations, 51 across all runs. The convergence rate is quadratic. Execution times are the fastest among the three methods, ranging from approximately 0.072 to 0.095 seconds, demonstrating the highest efficiency for the preconditioned Truncated Newton for the Penalty function at this dimension.

Table 33: Truncated Newton Method (with Preconditioning) Results on the Penalty Function

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	51	5.186 490 $\times 10^{-7}$	quadratic	0.084 834
x1	51	5.186 494 $\times 10^{-7}$	quadratic	0.076 772
x2	51	5.186 482 $\times 10^{-7}$	quadratic	0.080 170
x3	51	5.186 499 $\times 10^{-7}$	quadratic	0.072 051
x4	51	5.186 487 $\times 10^{-7}$	quadratic	0.076 069
x5	51	5.186 491 $\times 10^{-7}$	quadratic	0.075 938
x6	51	5.186 491 $\times 10^{-7}$	quadratic	0.095 852
x7	51	5.186 496 $\times 10^{-7}$	quadratic	0.073 183
x8	51	5.186 494 $\times 10^{-7}$	quadratic	0.076 804
x9	51	5.186 494 $\times 10^{-7}$	quadratic	0.076 413
x10	51	5.186 495 $\times 10^{-7}$	quadratic	0.091 382

3.4 Problem 76.

The **function** under consideration is defined as:

$$F(x) = \frac{1}{2} \sum_{k=1}^n f_k^2(x);$$

$$f_k(x) = x_k - \frac{x_{k+1}^2}{10}; \quad 1 \leq k < n; \quad f_k(x) = x_k - \frac{x_1^2}{10}; \quad k = n$$

A visualization of the function, in three dimension, is shown in Figure 5. In addition to the reference **starting point** $\bar{x} = [2, 2, \dots, 2]$, ten initial points were randomly generated: these points are uniformly distributed within a hypercube centered at \bar{x} , defined by the Cartesian product of intervals $[\bar{x}_i - 1, \bar{x}_i + 1]$ for each dimension i .

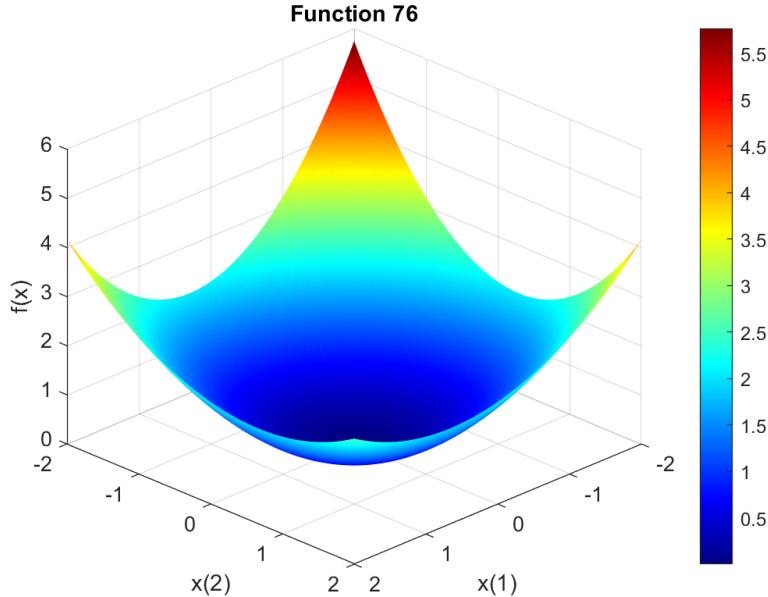


Figure 5: Plot of the function of the Problema 76

In line with the approach taken previously, a grid search was also performed for this function. The results that emerged are presented below:

Grid search result - Best combination: `c1 = 0.00001, rho = 0.6, btmax = 90.` Due to the Armijo condition not being satisfied at dimension $n = 10^5$, the backtracking line search parameters were adjusted to the following values: `c1 = 0.00001, rho = 0.5, btmax = 90.`

3.4.1 Dimension $n = 10^3$

1. **Modified Newton Method (Table 34):** Shows variability in performance across different initial conditions. The number of iterations ranges from 8 to 30, with corresponding fluctuations in execution times, from approximately 0.07 to 0.13 seconds. Despite the varying iteration counts, the method maintains quadratic convergence for all initial conditions. This variability in iterations and execution times, while still achieving fast convergence overall, suggests that the performance of Modified Newton on Function 76 might be somewhat sensitive to the starting point.

Table 34: Modified Newton Method Results on the Function 76

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	8	$2.294\,652 \times 10^{-11}$	quadratic	0.069 396
x1	30	$5.320\,419 \times 10^{-10}$	quadratic	0.134 682
x2	22	$2.839\,890 \times 10^{-7}$	quadratic	0.106 301
x3	28	$6.361\,118 \times 10^{-11}$	quadratic	0.123 453
x4	17	$6.224\,904 \times 10^{-11}$	quadratic	0.082 068
x5	23	$2.937\,718 \times 10^{-13}$	quadratic	0.103 911
x6	26	$1.145\,436 \times 10^{-11}$	quadratic	0.114 799
x7	17	$9.814\,967 \times 10^{-8}$	quadratic	0.078 959
x8	20	$1.021\,642 \times 10^{-10}$	quadratic	0.076 678
x9	16	$4.716\,091 \times 10^{-7}$	quadratic	0.070 473
x10	20	$2.466\,623 \times 10^{-7}$	quadratic	0.080 807

2. **Truncated Newton Method (No Preconditioning) Results (Table 35):** The number of iterations is consistently low, ranging from just 7 to 9 across all initial conditions. Execution times are also very fast and uniform, generally staying below 0.02 seconds. The method achieves quadratic convergence. This indicates that for Function 76, Truncated Newton without preconditioning is substantially more robust and computationally cheaper than Modified Newton.

Table 35: Truncated Newton Method (No Preconditioning) Results on the Function 76

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	7	$1.058\,551 \times 10^{-13}$	quadratic	0.020
x1	8	$2.095\,178 \times 10^{-8}$	quadratic	0.012
x2	7	$1.161\,611 \times 10^{-8}$	quadratic	0.007
x3	9	$1.672\,809 \times 10^{-8}$	quadratic	0.008
x4	7	$2.214\,241 \times 10^{-8}$	quadratic	0.011
x5	8	$1.262\,821 \times 10^{-8}$	quadratic	0.007
x6	7	$1.316\,989 \times 10^{-8}$	quadratic	0.006
x7	8	$7.243\,925 \times 10^{-7}$	quadratic	0.007
x8	7	$1.854\,992 \times 10^{-8}$	quadratic	0.006
x9	9	$1.512\,569 \times 10^{-8}$	quadratic	0.007
x10	7	$5.447\,982 \times 10^{-9}$	quadratic	0.006

3. **Truncated Newton Method (with Preconditioning) Results (Table 36):** Exhibits very similar

performance to the non-preconditioned version. Iteration counts are uniformly low, ranging from 5 to 7. Execution times are also consistently fast and low, mostly below 0.006 seconds. The method maintains quadratic convergence.

Table 36: Truncated Newton Method (with Preconditioning) Results on the Function 76

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	5	$1.179\,413 \times 10^{-10}$	quadratic	0.017514
x1	5	$1.795\,178 \times 10^{-8}$	quadratic	0.003672
x2	5	$9.616\,114 \times 10^{-9}$	quadratic	0.002112
x3	7	$1.472\,809 \times 10^{-8}$	quadratic	0.003341
x4	5	$1.914\,241 \times 10^{-8}$	quadratic	0.005590
x5	5	$1.062\,821 \times 10^{-8}$	quadratic	0.001939
x6	5	$1.116\,989 \times 10^{-8}$	quadratic	0.001337
x7	6	$6.243\,925 \times 10^{-7}$	quadratic	0.001544
x8	5	$1.654\,992 \times 10^{-8}$	quadratic	0.001401
x9	5	$1.312\,569 \times 10^{-8}$	quadratic	0.001615
x10	5	$4.447\,982 \times 10^{-9}$	quadratic	0.001386

Finite Difference Approximations of Derivatives. In line with the approach taken previously, a grid search was also performed for this function. The optimal parameters, balancing accuracy and efficiency, are summarized as follows: for Modified Newton, `type_fixed_mn = 'c'`, `k_fixed_mn = 4`, `type_variable_mn = 'c'`, and `k_esp_mn = 4`; for Truncated Newton, `type_fixed_trunc = 'c'`, `k_fixed_trunc = 4`, `type_variable_trunc = 'c'`, and `k_esp_trunc = 4`.

1. **Modified Newton Method - Finite Differences (FD) (Table 37):** The Modified Newton method, when employing Finite Differences for gradient approximation, shows variability in the number of iterations needed for convergence, ranging from 5 to 15. Correspondingly, execution times vary between approximately 0.14 to 0.31 seconds. Both Fixed and Variable FD types demonstrate linear convergence. There is no clear advantage of Variable FD over Fixed FD in terms of iterations or robustness, though Variable FD tends to have slightly lower execution times on average within each starting point.

Table 37: Modified Newton Method Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	5	1.017 950 $\times 10^{-8}$	linear	0.176
	Variable	5	3.525 327 $\times 10^{-8}$	linear	0.142
x1	Fixed	12	6.416 178 $\times 10^{-7}$	linear	0.267
	Variable	12	7.154 511 $\times 10^{-7}$	linear	0.265
x2	Fixed	13	3.182 841 $\times 10^{-9}$	linear	0.285
	Variable	13	4.078 440 $\times 10^{-9}$	linear	0.285
x3	Fixed	14	5.868 200 $\times 10^{-10}$	linear	0.300
	Variable	14	3.888 757 $\times 10^{-9}$	linear	0.305
x4	Fixed	14	6.559 805 $\times 10^{-9}$	linear	0.311
	Variable	14	3.687 372 $\times 10^{-9}$	linear	0.300
x5	Fixed	13	5.304 001 $\times 10^{-7}$	linear	0.278
	Variable	13	5.315 900 $\times 10^{-7}$	linear	0.282
x6	Fixed	13	3.733 829 $\times 10^{-7}$	linear	0.282
	Variable	13	4.176 414 $\times 10^{-7}$	linear	0.281
x7	Fixed	15	5.113 560 $\times 10^{-10}$	linear	0.317
	Variable	15	1.807 548 $\times 10^{-9}$	linear	0.308
x8	Fixed	14	3.151 367 $\times 10^{-10}$	linear	0.284
	Variable	14	9.830 668 $\times 10^{-10}$	linear	0.282
x9	Fixed	15	2.647 676 $\times 10^{-10}$	linear	0.304
	Variable	15	7.471 216 $\times 10^{-10}$	linear	0.304

2. Truncated Newton Method (No Preconditioning) - Finite Differences (FD) (Table 38):

The Truncated Newton method without preconditioning, using Finite Differences, achieves convergence in a consistent and low number of iterations, mostly at 6 iterations, some at 7. Execution times are also consistently low, ranging from approximately 0.11 to 0.15 seconds. Both Fixed and Variable FD maintain linear convergence, and the performance difference between Fixed and Variable FD is minimal in terms of iteration count and execution time, with Variable FD being slightly faster in some instances.

Table 38: Truncated Newton Method (No Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	6	1.017 920 $\times 10^{-8}$	linear	0.145
	Variable	6	3.525 344 $\times 10^{-8}$	linear	0.132
x1	Fixed	6	2.052 202 $\times 10^{-9}$	linear	0.128
	Variable	6	4.214 051 $\times 10^{-9}$	linear	0.122
x2	Fixed	6	1.975 535 $\times 10^{-9}$	linear	0.133
	Variable	6	4.327 468 $\times 10^{-9}$	linear	0.136
x3	Fixed	6	1.647 484 $\times 10^{-9}$	linear	0.130
	Variable	6	3.746 273 $\times 10^{-9}$	linear	0.127
x4	Fixed	6	1.294 336 $\times 10^{-7}$	linear	0.131
	Variable	6	1.368 045 $\times 10^{-7}$	linear	0.124
x5	Fixed	7	2.327 810 $\times 10^{-11}$	linear	0.146
	Variable	7	6.489 703 $\times 10^{-11}$	linear	0.144
x6	Fixed	6	1.474 707 $\times 10^{-8}$	linear	0.126
	Variable	6	1.639 071 $\times 10^{-8}$	linear	0.124
x7	Fixed	6	3.359 094 $\times 10^{-9}$	linear	0.141
	Variable	6	6.515 151 $\times 10^{-9}$	linear	0.117
x8	Fixed	6	8.347 541 $\times 10^{-9}$	linear	0.117
	Variable	6	1.514 122 $\times 10^{-8}$	linear	0.145
x9	Fixed	6	1.283 134 $\times 10^{-9}$	linear	0.119
	Variable	6	1.334 577 $\times 10^{-9}$	linear	0.121

3. Truncated Newton Method (with Preconditioning) - Finite Differences (FD) (Table 39):

Truncated Newton with preconditioning, using Finite Differences, exhibits similar performance to the non-preconditioned Truncated Newton. The iteration counts are consistently at 5 for most runs, with a few at 6. Execution times are also very consistent and low, primarily in the range of 0.11 to 0.15 seconds. Linear convergence is observed. Again, Variable FD offers marginally faster execution times compared to Fixed FD for most initial points, but the difference is not substantial.

Hypothesis on Similar Performance with and without Preconditioning. The near-identical performance of Truncated Newton with and without preconditioning on Function 76 might suggest that for this specific function and dimension, the Hessian matrix (or its approximation used in Truncated Newton) is already well-conditioned or has favorable spectral properties. In such cases, preconditioning, which is designed to improve the condition number and eigenvalue distribution of the Hessian to accelerate convergence, provides minimal additional benefit because the original Hessian is already suitable for efficient iterative solution of the Newton system. This could be due to the specific structure of Function 76 leading to a Hessian that inherently facilitates fast convergence even without explicit preconditioning.

Table 39: Truncated Newton Method (with Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	5	1.965 895 $\times 10^{-8}$	linear	0.146
	Variable	5	3.908 981 $\times 10^{-8}$	linear	0.124
x1	Fixed	5	3.719 141 $\times 10^{-9}$	linear	0.121
	Variable	5	5.237 263 $\times 10^{-9}$	linear	0.115
x2	Fixed	5	4.420 889 $\times 10^{-9}$	linear	0.126
	Variable	5	5.902 043 $\times 10^{-9}$	linear	0.120
x3	Fixed	5	3.039 380 $\times 10^{-9}$	linear	0.126
	Variable	5	4.547 040 $\times 10^{-9}$	linear	0.125
x4	Fixed	5	1.394 737 $\times 10^{-7}$	linear	0.117
	Variable	5	1.467 573 $\times 10^{-7}$	linear	0.121
x5	Fixed	6	2.313 876 $\times 10^{-11}$	linear	0.141
	Variable	6	6.423 217 $\times 10^{-11}$	linear	0.141
x6	Fixed	5	1.760 184 $\times 10^{-8}$	linear	0.119
	Variable	5	1.935 161 $\times 10^{-8}$	linear	0.118
x7	Fixed	5	6.775 715 $\times 10^{-9}$	linear	0.121
	Variable	5	8.817 028 $\times 10^{-9}$	linear	0.115
x8	Fixed	5	1.675 927 $\times 10^{-8}$	linear	0.110
	Variable	5	2.186 957 $\times 10^{-8}$	linear	0.115
x9	Fixed	5	2.634 513 $\times 10^{-9}$	linear	0.112
	Variable	5	2.661 198 $\times 10^{-9}$	linear	0.115

3.4.2 Dimension $n = 10^4$

1. **Modified Newton Method (Table 40):** The number of iterations ranges from 8 to 30, with corresponding fluctuations in execution times, from approximately 0.07 to 0.13 seconds. Despite the varying iteration counts, the method maintains quadratic convergence for all initial conditions. This variability in iterations and execution times, while still achieving fast convergence overall, suggests that the performance of Modified Newton on Function 76 might be somewhat sensitive to the starting point.

Table 40: Modified Newton Method Results on the Function 76

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	8	2.294 652 $\times 10^{-11}$	quadratic	0.069 396
x1	30	5.320 419 $\times 10^{-10}$	quadratic	0.134 682
x2	22	2.839 890 $\times 10^{-7}$	quadratic	0.106 301
x3	28	6.361 118 $\times 10^{-11}$	quadratic	0.123 453
x4	17	6.224 904 $\times 10^{-11}$	quadratic	0.082 068
x5	23	2.937 718 $\times 10^{-13}$	quadratic	0.103 911
x6	26	1.145 436 $\times 10^{-11}$	quadratic	0.114 799
x7	17	9.814 967 $\times 10^{-8}$	quadratic	0.078 959
x8	20	1.021 642 $\times 10^{-10}$	quadratic	0.076 678
x9	16	4.716 091 $\times 10^{-7}$	quadratic	0.070 473
x10	20	2.466 623 $\times 10^{-7}$	quadratic	0.080 807

2. **Truncated Newton Method (No Preconditioning) Results (Table 41):** Shows significantly more consistent and efficient performance compared to the Modified Newton method. The number

of iterations is consistently low, ranging from just 7 to 9 across all initial conditions. Execution times are also very fast and uniform, generally staying below 0.02 seconds. The method achieves quadratic convergence. This indicates that for Function 76, Truncated Newton without preconditioning is substantially more robust and computationally cheaper than Modified Newton.

Table 41: Truncated Newton Method (No Preconditioning) Results on the Function 76

$x(0)$	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	7	1.058 551 $\times 10^{-13}$	quadratic	0.020
x1	8	2.095 178 $\times 10^{-8}$	quadratic	0.012
x2	7	1.161 611 $\times 10^{-8}$	quadratic	0.007
x3	9	1.672 809 $\times 10^{-8}$	quadratic	0.008
x4	7	2.214 241 $\times 10^{-8}$	quadratic	0.011
x5	8	1.262 821 $\times 10^{-8}$	quadratic	0.007
x6	7	1.316 989 $\times 10^{-8}$	quadratic	0.006
x7	8	7.243 925 $\times 10^{-7}$	quadratic	0.007
x8	7	1.854 992 $\times 10^{-8}$	quadratic	0.006
x9	9	1.512 569 $\times 10^{-8}$	quadratic	0.007
x10	7	5.447 982 $\times 10^{-9}$	quadratic	0.006

3. **Truncated Newton Method (with Preconditioning) Results (Table 42):** Exhibits very similar performance to the non-preconditioned version. Iteration counts are uniformly low, ranging from 5 to 7. Execution times are also consistently fast and low, mostly below 0.006 seconds. The method maintains quadratic convergence.

Table 42: Truncated Newton Method (with Preconditioning) Results on the Function 76

$x(0)$	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	5	1.179 413 $\times 10^{-10}$	quadratic	0.017 514
x1	5	1.795 178 $\times 10^{-8}$	quadratic	0.003 672
x2	5	9.616 114 $\times 10^{-9}$	quadratic	0.002 112
x3	7	1.472 809 $\times 10^{-8}$	quadratic	0.003 341
x4	5	1.914 241 $\times 10^{-8}$	quadratic	0.005 590
x5	5	1.062 821 $\times 10^{-8}$	quadratic	0.001 939
x6	5	1.116 989 $\times 10^{-8}$	quadratic	0.001 337
x7	6	6.243 925 $\times 10^{-7}$	quadratic	0.001 544
x8	5	1.654 992 $\times 10^{-8}$	quadratic	0.001 401
x9	5	1.312 569 $\times 10^{-8}$	quadratic	0.001 615
x10	5	4.447 982 $\times 10^{-9}$	quadratic	0.001 386

Finite Difference Approximations of Derivatives.

1. **Modified Newton Method (FD) (Table 43):** The performance varies significantly across different initial conditions and FD types (fixed vs. variable 'h'). Generally, Modified Newton requires a substantial number of iterations, leading to longer execution times compared to Truncated Newton methods. Looking at the data, iteration counts range from 5 to 39, and execution times vary accordingly from approximately 5 to 40 seconds. There's noticeable variability in performance depending on the initial condition, with condition 1 showing much faster convergence (5 iterations) than others (e.g., condition 8 with 39 iterations). The choice of fixed or variable 'h' FD does not appear to drastically change the performance of Modified Newton in terms of iteration count and execution time, although minor variations are observed.

Table 43: Modified Newton Method Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	5	$3.236\ 841 \times 10^{-8}$	linear	5.986
	Variable	5	$1.107\ 104 \times 10^{-7}$	linear	5.907
x1	Fixed	33	$6.215\ 201 \times 10^{-7}$	linear	33.405
	Variable	34	$4.998\ 952 \times 10^{-11}$	linear	34.661
x2	Fixed	32	$4.550\ 919 \times 10^{-10}$	linear	32.858
	Variable	32	$1.304\ 117 \times 10^{-9}$	linear	33.098
x3	Fixed	38	$9.169\ 125 \times 10^{-10}$	linear	38.616
	Variable	38	$3.426\ 964 \times 10^{-9}$	linear	38.716
x4	Fixed	30	$1.342\ 108 \times 10^{-9}$	linear	30.518
	Variable	29	$1.172\ 445 \times 10^{-7}$	linear	30.432
x5	Fixed	32	$5.851\ 934 \times 10^{-10}$	linear	32.613
	Variable	32	$1.881\ 853 \times 10^{-9}$	linear	32.653
x6	Fixed	33	$2.260\ 385 \times 10^{-10}$	linear	33.572
	Variable	33	$4.273\ 617 \times 10^{-10}$	linear	33.418
x7	Fixed	39	$2.682\ 989 \times 10^{-7}$	linear	39.882
	Variable	39	$3.258\ 138 \times 10^{-7}$	linear	39.736
x8	Fixed	35	$2.595\ 670 \times 10^{-11}$	linear	35.754
	Variable	35	$9.339\ 306 \times 10^{-11}$	linear	36.135
x9	Fixed	32	$5.838\ 992 \times 10^{-8}$	linear	32.883
	Variable	32	$3.668\ 662 \times 10^{-8}$	linear	32.511
x10	Fixed	39	$7.986\ 733 \times 10^{-11}$	linear	39.470
	Variable	39	$3.116\ 173 \times 10^{-10}$	linear	39.294

2. **Truncated Newton Method (No Preconditioning) (FD) (Table 44):** The Truncated Newton method without preconditioning, when using Finite Differences for Hessian approximation at dimension 10^4 , demonstrates significantly improved efficiency and consistency compared to the Modified Newton method. The number of iterations is consistently low, primarily at 5 or 6 across all initial conditions and FD types. Execution times are also remarkably fast, generally staying in the range of 5-7 seconds. The choice between fixed and variable 'h' FD has a minimal impact on performance for this method. Truncated Newton without preconditioning offers a much more robust and computationally economical approach than Modified Newton for this problem setting.

Table 44: Truncated Newton Method (No Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	5	$3.249\,275 \times 10^{-8}$	linear	5.852
	Variable	5	$1.107\,242 \times 10^{-7}$	linear	5.907
x1	Fixed	5	$2.218\,688 \times 10^{-8}$	linear	5.918
	Variable	5	$2.707\,176 \times 10^{-8}$	linear	5.914
x2	Fixed	6	$1.566\,940 \times 10^{-11}$	linear	6.931
	Variable	6	$3.173\,318 \times 10^{-11}$	linear	6.754
x3	Fixed	5	$1.388\,209 \times 10^{-9}$	linear	6.791
	Variable	5	$2.022\,218 \times 10^{-9}$	linear	5.837
x4	Fixed	6	$4.675\,678 \times 10^{-11}$	linear	6.974
	Variable	6	$9.977\,310 \times 10^{-11}$	linear	6.905
x5	Fixed	5	$2.808\,678 \times 10^{-7}$	linear	5.887
	Variable	5	$2.912\,386 \times 10^{-7}$	linear	5.780
x6	Fixed	5	$1.708\,186 \times 10^{-9}$	linear	5.961
	Variable	5	$3.152\,597 \times 10^{-9}$	linear	5.931
x7	Fixed	5	$5.980\,240 \times 10^{-7}$	linear	5.879
	Variable	5	$6.045\,083 \times 10^{-7}$	linear	5.934
x8	Fixed	6	$8.961\,085 \times 10^{-7}$	linear	5.934
	Variable	6	$8.758\,012 \times 10^{-7}$	linear	5.962
x9	Fixed	5	$8.333\,786 \times 10^{-8}$	linear	5.973
	Variable	5	$8.845\,605 \times 10^{-8}$	linear	5.889
x10	Fixed	5	$1.429\,689 \times 10^{-8}$	linear	5.842
	Variable	5	$2.867\,084 \times 10^{-8}$	linear	5.898

3. **Truncated Newton Method (with Preconditioning) (FD)** (Table 45): Shows very consistent and efficient performance, closely mirroring the non-preconditioned Truncated Newton method. Iteration counts are consistently low, mostly at 5 or 6, and execution times are also uniformly fast, generally around 5-7 seconds, very similar to the non-preconditioned version. Again, the choice of fixed versus variable 'h' FD has negligible impact on the performance.

Table 45: Truncated Newton Method (with Preconditioning) Results - Finite Differences (FD)

x(0)	FD Type	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	Fixed	5	$6.368\ 914 \times 10^{-7}$	linear	5.897
	Variable	5	$6.551\ 582 \times 10^{-7}$	linear	5.873
x1	Fixed	5	$5.291\ 634 \times 10^{-7}$	linear	5.852
	Variable	5	$5.358\ 395 \times 10^{-7}$	linear	5.872
x2	Fixed	6	$5.143\ 089 \times 10^{-11}$	linear	6.959
	Variable	6	$1.069\ 600 \times 10^{-10}$	linear	6.860
x3	Fixed	5	$6.498\ 681 \times 10^{-8}$	linear	5.863
	Variable	5	$6.501\ 497 \times 10^{-8}$	linear	5.925
x4	Fixed	6	$1.046\ 381 \times 10^{-10}$	linear	6.819
	Variable	6	$2.246\ 478 \times 10^{-10}$	linear	6.835
x5	Fixed	6	$2.428\ 166 \times 10^{-11}$	linear	6.852
	Variable	6	$5.227\ 231 \times 10^{-11}$	linear	6.850
x6	Fixed	5	$6.630\ 910 \times 10^{-8}$	linear	5.801
	Variable	5	$6.627\ 318 \times 10^{-8}$	linear	5.837
x7	Fixed	6	$3.654\ 649 \times 10^{-11}$	linear	7.088
	Variable	6	$7.751\ 665 \times 10^{-11}$	linear	6.834
x8	Fixed	6	$4.650\ 494 \times 10^{-11}$	linear	6.898
	Variable	6	$9.600\ 179 \times 10^{-11}$	linear	7.421
x9	Fixed	5	$9.081\ 906 \times 10^{-7}$	linear	5.893
	Variable	5	$9.210\ 954 \times 10^{-7}$	linear	5.830
x10	Fixed	5	$2.958\ 050 \times 10^{-7}$	linear	5.953
	Variable	3	$3.080\ 444 \times 10^{-7}$	linear	5.856

3.4.3 Dimension $n = 10^5$

1. **Modified Newton Method (Table 46):** The number of iterations varies considerably across starting points, ranging from 28 to 273. Consequently, execution times are also much higher and vary widely, from under 1 second up to over 9 seconds. While quadratic convergence is indicated for several starting points, some instances show superlinear or even linear convergence, suggesting the convergence behavior becomes less consistent in higher dimensions for this method.

Table 46: Modified Newton Method Results on the Function 76

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	28	$2.062\ 866 \times 10^{-8}$	quadratic	0.942
x1	272	$3.462\ 105 \times 10^{-8}$	superlinear	9.138
x2	176	$6.117\ 679 \times 10^{-9}$	superlinear	5.826
x3	273	$1.977\ 789 \times 10^{-10}$	superlinear	9.170
x4	164	$5.918\ 233 \times 10^{-7}$	quadratic	5.430
x5	261	$4.280\ 968 \times 10^{-9}$	superlinear	8.814
x6	263	$3.153\ 701 \times 10^{-10}$	linear	8.780
x7	177	$4.351\ 092 \times 10^{-8}$	superlinear	6.006
x8	182	$3.769\ 938 \times 10^{-9}$	quadratic	6.069
x9	270	$3.673\ 482 \times 10^{-12}$	linear	9.218
x10	187	$1.228\ 028 \times 10^{-13}$	quadratic	6.408

2. Truncated Newton Method (No Preconditioning) (Table 47): The number of iterations remains very low, consistently between 6 and 7 across all starting points. Execution times are also dramatically reduced, staying well below 0.1 seconds, indicating a substantial computational advantage over Modified Newton. The convergence rate is consistently quadratic across all trials.

Table 47: Truncated Newton Method (No Preconditioning) Results on the Function 76

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	6	2.068 488 $\times 10^{-11}$	quadratic	0.095
x1	7	2.519 541 $\times 10^{-7}$	quadratic	0.070
x2	6	4.150 656 $\times 10^{-7}$	quadratic	0.065
x3	7	5.750 444 $\times 10^{-7}$	quadratic	0.068
x4	6	2.895 838 $\times 10^{-7}$	quadratic	0.067
x5	6	2.833 041 $\times 10^{-7}$	quadratic	0.062
x6	6	2.225 910 $\times 10^{-7}$	quadratic	0.064
x7	7	2.482 301 $\times 10^{-7}$	quadratic	0.066
x8	6	2.320 126 $\times 10^{-7}$	quadratic	0.065
x9	6	1.893 438 $\times 10^{-7}$	quadratic	0.063
x10	7	7.129 038 $\times 10^{-7}$	quadratic	0.067

3. Truncated Newton Method (with Preconditioning) (Table 48): The iteration counts are minimal, consistently at 5 for all starting conditions. Execution times are also very low, generally around 0.07 seconds, and are slightly faster or comparable to the non-preconditioned Truncated Newton method. Quadratic convergence is maintained throughout.

Table 48: Truncated Newton Method (with Preconditioning) Results on the Function 76

x(0)	Iterations	$\ \nabla f(x^*)\ $	Exp. Rate of Conv.	Execution Time (s)
x0	5	2.976 640 $\times 10^{-12}$	quadratic	0.106
x1	5	2.519 541 $\times 10^{-7}$	quadratic	0.069
x2	5	4.150 656 $\times 10^{-7}$	quadratic	0.071
x3	5	5.750 444 $\times 10^{-7}$	quadratic	0.070
x4	5	2.895 838 $\times 10^{-7}$	quadratic	0.074
x5	5	2.833 041 $\times 10^{-7}$	quadratic	0.065
x6	5	2.225 910 $\times 10^{-7}$	quadratic	0.070
x7	5	2.482 301 $\times 10^{-7}$	quadratic	0.068
x8	5	2.320 126 $\times 10^{-7}$	quadratic	0.070
x9	5	1.893 438 $\times 10^{-7}$	quadratic	0.070
x10	5	7.129 038 $\times 10^{-7}$	quadratic	0.070

Finite Difference Approximations of Derivatives. In this scenario, the results obtained using finite difference approximations are not competitive when compared to employing exact gradients and Hessians. Indeed, the time required to produce results is exceedingly long. The execution was not halted prematurely after a few points, and consequently, the outcomes are not being reported due to the protracted execution times and lack of competitive performance.

4 Final considerations

This homework has presented a comprehensive study of two advanced iterative methods for unconstrained optimization: the Modified Newton Method and the Truncated Newton Method (with and without preconditioning). The methods were implemented in MATLAB and rigorously tested on several benchmark

problems, including the classical Rosenbrock function as well as higher-dimensional and more challenging problems such as the banded trigonometric problem and the penalty function.

The investigation began with a detailed discussion of the theoretical foundations and practical considerations underlying each method. The Modified Newton Method was shown to efficiently handle cases where the Hessian matrix might be indefinite by incorporating an adaptive correction strategy. On the other hand, the Truncated Newton Method was designed to address the high computational cost associated with explicitly forming and factorizing the Hessian matrix in large-scale problems, by instead using iterative methods (e.g., Conjugate Gradient) to approximate the Newton step.

Extensive numerical experiments were conducted using a grid search approach to optimize hyperparameters such as c_1 , ρ , and $btmax$. The grid search identified the optimal combination as $c_1 = 10^{-4}$, $\rho = 0.8$, and $btmax = 30$, which was used across different methods to ensure a fair comparison. The results indicate that both the Modified Newton Method and the Truncated Newton Method with preconditioning exhibit strong convergence properties and robust performance across various initial conditions. In particular, starting points closer to the global minimum yielded rapid convergence, while those further away demanded significantly more iterations.

A key observation was that preconditioning plays a crucial role in enhancing the performance of the Truncated Newton Method. When preconditioning was applied, the method achieved lower iteration counts and reduced execution times, making its performance comparable to the Modified Newton Method. Conversely, the unpreconditioned variant generally required a higher number of iterations and was more sensitive to the choice of the initial guess.

Furthermore, the study also explored the use of finite difference approximations to compute gradients and Hessians. Although finite difference schemes (both fixed and variable step sizes) were able to produce convergent results, they incurred a significantly higher computational cost compared to methods that employed exact derivatives. This trade-off was particularly evident in high-dimensional problems, where the increased iteration count and longer execution times rendered finite differences less competitive.

The results of this study underscore the importance of method selection and parameter tuning in large-scale unconstrained optimization. While both methods are effective, the Modified Newton Method and the preconditioned Truncated Newton Method stand out as robust and efficient options.

A Appendix: Problem implementation and Main Code

```

1 % ROSENROCK
2 f_rosenbrock = @(x) 100*(x(2,:)-x(1,:).^2).^2 + (1-x(1,:)).^2;
3 gradf_rose = @(x) [400*x(1,:).^3 - 400*x(1,:).*x(2,:)+2*x(1,:)-2;
4 %                200*(x(2,:)-x(1,:).^2)];
5 Hessf_rose = @(x) [1200*x(1)^2 - 400*x(2)+2, -400*x(1);
6 %                  -400*x(1), 200];
7
8 % 16 - BANDED TRIGONOMETRIC FUNCTION
9 f_16 = @(x) deal(@(x) (x=x(:); n=length(x); if n==1, Fval = 1 - cos(x(1)); else, val1 = (1 -
10 %           cos(x(1))) - sin(x(2)); val2 = 0; if n > 2, ii = (2:n-1)'; val2 = sum(ii.* (1 - cos(x
11 %           (ii))) + sin(x(ii-1)) - sin(x(ii+1))); end; val3 = n*(1 - cos(x(n))) + n*sin(x(n-1));
12 %           Fval = val1 + val2 + val3; end; Fval; )());
13 grad_16 = @(x) deal(@(x) (x=x(:); n=length(x); g=zeros(n,1); for k = 1:n, switch k, case 1,
14 %           g(1) = sin(x(1)) + cos(x(1)); case 2, g(2) = 2*sin(x(2)); case {3 : (n-2)}, g(k) = k *
15 %           sin(x(k)); case (n-1), g(n-1) = (n-1)*( sin(x(n-1)) + cos(x(n-1)) ); case n, g(n) = n *
16 %           sin(x(n)) - cos(x(n)); end; end; g; )());
17 hess_16 = @(x) deal(@(x) (x=x(:); n=length(x); d=zeros(n,1); d(1) = cos(x(1)) - sin(x(1));
18 %           for k = 2 : (n-2), d(k) = k * cos(x(k)); end; if n >= 2, d(n-1) = (n-1)*( cos(x(n-1)) -
19 %           sin(x(n-1)) ); end; d(n) = n*cos(x(n)) + sin(x(n)); sparse(1:n, 1:n, d, n, n); )());
20
21 % 27 - PENALTY FUNCTION
22 f_penalty = @(x) 0.5 * ( sum( ((1/sqrt(100000)*(x-1)).^2) ) + (sum(x.^2)-0.25)^2 );
23 grad_penalty = @(x) (x-1)/100000 + 2*x*(sum(x.^2) - 0.25);
24 Hess_diag = @(x) spdiags((1/100000 + 2*(sum(x.^2)-0.25)) + 4*(x.^2), 0, length(x), length(x));
25
26 Hess_nondiag = @(x) 4 * (x * x') - spdiags(4 * x.^2, 0, length(x), length(x));
27 Hess_penalty = @(x) sparse(Hess_diag(x) + Hess_nondiag(x));

```

```

20
21 % 76 - FUNCTION 76
22 f_76 = @(x) 0.5 * sum ( ( x(1:length(x)) - 0.1*[x(2:length(x));x(1)].^2 ).^2 );
23 grad_76 = @(x) x(1:length(x)) - [x(2:length(x));x(1)].^2 * 0.1 - 0.2 * x(1:length(x)) .* (
24     [x(length(x));x(1:length(x)-1)] - ...
25     x(1:length(x)).^2 * 0.1 );
26 Hess_76 = @(x) Hess_Problem76_inline(x);
27
28 function H = Hess_Problem76_inline(x).
29
30     n = length(x);
31     v1 = 1 - 0.2 * [x(n);x(1:n-1)] + 3/50 * x(1:n).^2; % Diagonal elements
32     v2 = -0.2 * x(2:n); % Off-diagonal elements (super and sub)
33     v3 = -0.2*x(1)*ones(2,1); % Corner elements (1,n) and (n,1)
34     H = sparse(1:n, 1:n, v1, n, n); % Main diagonal
35     H = H + sparse(2:n, 1:n-1, v2, n, n); % Lower diagonal
36     H = H + sparse(1:n-1, 2:n, v2, n, n); % Upper diagonal
37     H = H + sparse([1,n], [n,1], v3, n, n); % Corner elements
38
39 end

```

Listing 1: Function implementation

```

1 function [best_params, best_avg_iter] = hyperparameter_grid_search(f, grad, hess,
2     initial_conditions)
% HYPERPARAMETER_GRID_SEARCH Executes a grid search over hyperparameters for optimization
% methods.
%
4 % [best_params, best_avg_iter] = hyperparameter_grid_search(f, grad, hess,
5 %     initial_conditions)
%
6 % This function performs a grid search across combinations of hyperparameters
7 % to evaluate and optimize the performance of numerical optimization methods.
8 % It varies the parameters c1, rho, and btmax for the 'Modified Newton' and
9 % 'Truncated Newton Pre' methods, assessing the average performance across
10 % several provided initial conditions.
%
11 %
12 % INPUTS:
13 %   f           - Function handle for the objective function to be minimized.
14 %   grad        - Function handle for the gradient of the objective function.
15 %   hess        - Function handle for the Hessian of the objective function.
16 %   initial_conditions - Cell array of column vectors, representing different
17 %                         starting points for the optimization.
18 %
19 % OUTPUTS:
20 %   best_params      - Struct containing the best hyperparameter combination found
21 %                      during the grid search. The fields of the struct are:
22 %                      .c1, .rho, .btmax.
23 %   best_avg_iter    - Scalar representing the average number of iterations obtained
24 %                      with the best hyperparameter combination, averaged
25 %                      over all optimization methods and initial conditions tested.
26 %
27 %
28 % HYPERPARAMETERS DEFINED FOR GRID SEARCH:
29 % Hyperparameters being varied:
30 %   c1:          from 1e-4 to 1e-2 (Armijo condition parameter)
31 %   rho:          from 0.5 to 0.9 (Backtracking line search reduction factor)
32 %   btmax:        from 10 to 50 (Maximum backtracking iterations)
33
34
35 %% HYPERPARAMETER GRID DEFINITION
36 % Defines the ranges of values to explore for each hyperparameter
37 c1_vals      = [1e-4, 1e-3, 1e-2]; % Range of values for c1 (Armijo condition parameter)
38 rho_vals     = [0.5, 0.6, 0.7, 0.8, 0.9]; % Range of values for rho (backtracking
39 reduction factor); note: 0.9 might not always satisfy the Armijo condition
40 btmax_vals   = [30, 40, 50]; % Range of values for btmax (max backtracking iterations);
note: 10, 20, 30 might not always satisfy the Armijo condition

```

```

41 % Calculate the number of values for each hyperparameter
42 num_c1 = numel(c1_vals); % Number of c1 values to test
43 num_rho = numel(rho_vals); % Number of rho values to test
44 num_btmax = numel(btmax_vals); % Number of btmax values to test
45
46
47 %% VARIABLES TO STORE THE BEST RESULT
48 % Initialize with a large value (infinity) so any computed average is better
49 best_avg_iter = Inf;
50 % Structure to store the best hyperparameter combination found so far
51 best_params = struct('c1', NaN, 'rho', NaN, 'btmax', NaN); % Initialize best
      parameters to NaN
52
53
54 %% GRID SEARCH OVER PARAMETERS
55 % Cell array to hold the names of the optimization methods to be tested
56 methods = {'Modified Newton', 'Truncated Newton Pre'}; % List of optimization methods
57 num_methods = numel(methods); % Number of optimization methods
58
59
60 % Nested loops to iterate through all combinations of hyperparameters
61 for ic1 = 1:num_c1 % Iterate over each value of c1
62     for irho = 1:num_rho % Iterate over each value of rho
63         for ibt = 1:num_btmax % Iterate over each value of btmax
64             % Get the current hyperparameter values for this iteration
65             c1 = c1_vals(ic1); % Current c1 value
66             rho = rho_vals(irho); % Current rho value
67             btmax = btmax_vals(ibt); % Current btmax value
68
69             % Variable to accumulate the total number of iterations across methods and
                  initial conditions
70             iter_total = 0;
71             % Total number of runs (number of methods multiplied by number of initial
                  conditions)
72             run_count = 0;
73
74             % Loop over the different initial conditions to test robustness
75             for idx = 1:length(initial_conditions)
76                 x0 = initial_conditions{idx}; % Get the current initial condition
77
78                 % --- Modified Newton Method ---
79                 try
80                     % Call the modified_newton function with the current hyperparameters
                           and initial condition
81                     [~, ~, ~, k, ~, ~] = modified_newton(x0, f, grad, hess, kmax,
                           tolgrad, c1, rho, btmax);
82                     iter_total = iter_total + k; % Add the number of iterations 'k' if
                           the method converges successfully
83                 catch ME
84                     % If the method fails (throws an error), assume it took the maximum
                           number of iterations
85                     iter_total = iter_total + kmax; % Penalize failed runs by adding the
                           maximum iteration count
86                 end
87                 run_count = run_count + 1; % Increment the run count for Modified Newton
88
89             % --- Preconditioned Truncated Newton Method ---
90             try
91                 % Call the truncated_newton_pre function with the current
                           hyperparameters and initial condition
92                 [~, ~, ~, k, ~, ~] = truncated_newton_pre(x0, f, grad, hess, kmax,
                           tolgrad, c1, rho, btmax);
93                 iter_total = iter_total + k; % Add the number of iterations 'k' if
                           the method converges successfully
94             catch ME
95                 iter_total = iter_total + kmax; % Penalize failed runs by adding the
                           maximum iteration count
96             end
97         end

```

```

98             maximum iteration count
99         end
100        run_count = run_count + 1; % Increment the run count for Truncated
101        Newton Preconditioned
102    end
103
104    % Calculate the average number of iterations for the current hyperparameter
105    % combination
106    avg_iter = iter_total / run_count;
107
108    % Update the best combination if the current average number of iterations is
109    % lower (better performance)
110    if avg_iter < best_avg_iter
111        best_avg_iter = avg_iter; % Update the best average iteration count
112        best_params.c1 = c1; % Update the best c1 parameter
113        best_params.rho = rho; % Update the best rho parameter
114        best_params.btmax = btmax; % Update the best btmax parameter
115    end
116
117 %% GRID SEARCH RESULT
118 % Display the best hyperparameter combination found and the corresponding average
119 % iterations
120 disp(['Grid search result - Best combination: c1 = ', num2str(best_params.c1), ', rho =
121     ', num2str(best_params.rho), ', btmax = ', num2str(best_params.btmax)]);
122 disp(['Overall average iterations (over all methods and initial conditions): ', num2str(
123     best_avg_iter)]);
124
125 end

```

Listing 2: Grid search - c1, rho, btmax

```

1 function optimizationExperiment(f, gradf, hessf, x_base, kmax, tolgrad, c1, rho, btmax, dims
2 , num_starting_points)
%OPTIMIZATIONEXPERIMENT Executes and compares different optimization methods for a given
3 problem.
4 % This function performs a comparative analysis of optimization methods:
5 % Modified Newton, Truncated Newton without preconditioning, and Truncated
6 % Newton with preconditioning. It evaluates these methods across different
7 % problem dimensions and multiple initial starting points.
8 %
9 % INPUTS:
10 %   f           - Function handle for the objective function f(x).
11 %   gradf       - Function handle for the gradient of f, gradf(x).
12 %   hessf       - Function handle for the Hessian of f, hessf(x).
13 %   x_base      - Base vector used to generate initial conditions. It determines
14 %                 the dimension of the problem and serves as the first initial
15 %                 point.
16 %   kmax        - Maximum number of iterations for each optimization method.
17 %   tolgrad     - Tolerance for the gradient norm as the stopping criterion.
18 %   c1          - Parameter for the Armijo condition in line search.
19 %   rho          - Reduction factor for backtracking line search.
20 %   btmax       - Maximum iterations for backtracking line search.
21 %   dims         - Vector of exponents for problem dimensions (n = 10^d, where d is
22 %                 from dims).
23 %   num_starting_points - Number of initial starting points to use for each dimension.
24 %
25 % OUTPUTS:
26 % This function does not explicitly return values, but it prints out the
27 % performance metrics (iterations, convergence rate, execution time, gradient norm)
28 % for each optimization method, dimension, and initial starting point to the console.
% It's designed to provide a comparative report of the methods' effectiveness.
% --- Input Validation ---

```

```

29 if ~isa(f, 'function_handle') || ~isa(gradf, 'function_handle') || ~isa(hessf, 'function_handle')
30     error('f, gradf, and hessf must be function handles.');
31 end
32 if ~isvector(x_base)
33     error('x_base must be a vector.');
34 end
35 if ~isscalar(kmax) || ~isscalar(tolgrad) || ~isscalar(c1) || ~isscalar(rho) || ~isscalar(btmax)
36     error('kmax, tolgrad, c1, rho, and btmax must be scalar values.');
37 end
38 if ~isvector(dims) || ~isnumeric(dims)
39     error('dims must be a numeric vector.');
40 end
41 if ~isscalar(num_starting_points) || ~isnumeric(num_starting_points) || num_starting_points < 1
42     error('num_starting_points must be a scalar numeric value greater than or equal to 1.');
43 end
44
45 % Loop through different problem dimensions defined by 'dims'
46 for d = dims
47     n = 10^d; % Set problem dimension n based on the current exponent d
48     fprintf('\n=====\\n');
49     fprintf('Dimension: n = 10^%d = %d\\n', d, n);
50     fprintf('=====\\n');
51
52 % --- Creation of Initial Conditions ---
53 initial_conditions = cell(1, num_starting_points); % Cell array to store initial conditions
54 names = cell(1, num_starting_points); % Cell array to store names for initial conditions
55
56 % First initial point is the provided base point 'x_base' (resized to dimension n)
57 initial_conditions{1} = x_base(1:n); % Use the first n elements of x_base, or pad if x_base is shorter.
58 names{1} = 'x_0';
59
60 % Generate additional random initial points around x_base
61 for i = 1:(num_starting_points - 1)
62     x_rand = initial_conditions{1} + (2 * rand(1, n) - 1)'; % Random points within hypercube around x_base
63     initial_conditions{i+1} = x_rand;
64     names{i+1} = ['x_', num2str(i)];
65 end
66
67 % --- Loop over Initial Conditions for the current dimension ---
68 for i = 1:length(initial_conditions)
69     x0 = initial_conditions{i}; % Current initial condition
70     fprintf('\\n-----\\n');
71     fprintf('d = 10^%d, Initial Condition: %s\\n', d, names{i});
72     fprintf('-----\\n');
73
74 % --- Modified Newton Method ---
75 tic; % Start timer for Modified Newton
76 [xk_mn, fk_mn, grad_norm_mn, iter_mn, x_seq_mn, bt_seq_mn, rate_conv_mn] = ...
77     modified_newton(x0, f, gradf, hessf, kmax, tolgrad, c1, rho, btmax); % Run Modified Newton method
78 time_mn = toc; % Stop timer
79 fprintf('Modified Newton: iter = %d, rate = %f, time = %f sec, grad_norm = %e\\n',
80         iter_mn, rate_conv_mn, time_mn, grad_norm_mn); % Display results
81
82 % --- Truncated Newton Method WITHOUT Preconditioning ---
83 tic; % Start timer for Truncated Newton without preconditioning
84 [xk_tn, fk_tn, grad_norm_tn, iter_tn, x_seq_tn, bt_seq_tn, rate_conv_tn] = ...
85     truncated_newton(x0, f, gradf, hessf, kmax, tolgrad, c1, rho, btmax); % Run Truncated Newton method
86 time_tn = toc; % Stop timer
87 fprintf('Truncated Newton Without Preconditioning: iter = %d, rate = %f, time = %f

```

```

    sec, grad_norm = %e\n', iter_tn, rate_conv_tn, time_tn, grad_norm_tn); % Display
    results
88
89
90    % --- Truncated Newton Method WITH Preconditioning ---
91    tic; % Start timer for Truncated Newton with preconditioning
92    [xk_tnp, fk_tnp, grad_norm_tnp, iter_tnp, x_seq_tnp, bt_seq_tnp, rate_conv_tnp] =
93        ...
94        truncated_newton_pre(x0, f, gradf, hessf, kmax, tolgrad, c1, rho, btmax); % Run
        Preconditioned Truncated Newton method
95    time_tnp = toc; % Stop timer
96    fprintf('Truncated Newton Preconditioning: iter = %d, rate = %f, time = %f sec,
97        grad_norm = %e\n', iter_tnp, rate_conv_tnp, time_tnp, grad_norm_tnp); % Display
        results
98
99    end
100 end

```

Listing 3: Main Function

B Appendix: Modified Newton Method

```

1 function [xk, fk, gradfk_norm, k, xseq, btseq, rate_convergence] = modified_newton(x0, f,
2     gradf, Hessf, kmax, tolgrad, c1, rho, btmax, varargin)
% MODIFIED_NEWTON - Modified Newton method with backtracking.
3 %
4 % Syntax:
5 % 1) [xk, fk, gradfk_norm, k, xseq, btseq, rate_convergence] =
6 %     modified_newton(x0, f, gradf, Hessf, kmax, tolgrad, c1, rho, btmax)
7 %
8 %
9 % 2) [xk, fk, gradfk_norm, k, xseq, btseq, rate_convergence] =
10 %     modified_newton(x0, f, [], [], kmax, tolgrad, c1, rho, btmax, h, type)
11 %     => If the handles are empty or invalid, finite difference functions are used
12 %         to calculate the gradient and Hessian.
13 %
14 % In both cases, the operation is identical.
15
16
17 % Function handle for the Armijo condition
18 farmijo = @(fk, alpha, gradfk, pk) fk + c1 * alpha * (gradfk' * pk);
19
20
21 % Initializations
22 x0 = x0(:);
23 n = length(x0);
24 xseq = zeros(n, kmax);
25 btseq = zeros(1, kmax);
26
27 % For estimating the convergence rate
28 e = zeros(1, kmax);
29
30
31 xk = x0;
32 fk = f(xk);
33 gradfk = gradf(xk);
34 hessfk = Hessf(xk);
35 gradfk_norm = norm(gradfk);
36 Ek = speye(size(hessfk)); % Sparse identity matrix, for Hessian correction
37
38
39 k = 0;
40

```

```

41
42     while k < kmax && gradfk_norm >= tolgrad
43         tau = 0;
44         % Calculate beta as the minimum between the gradient norm and the Frobenius norm of
45         % the Hessian
46         beta = min(norm(gradfk), sqrt(sum(hessfk.^2, 'all')));
47         pd_success = false;
48         jmax = 100;
49
50         % Search for a tau such that Bk = Hessf + tau*I is PD (Positive Definite)
51         for j = 1:jmax
52             Bk = hessfk + tau * Ek;
53             [L, cholFlag] = chol(Bk, 'lower');
54             if cholFlag == 0
55                 pd_success = true;
56                 break;
57             else
58                 tau = max(2 * tau, beta / 2);
59             end
60         end
61
62         if ~pd_success
63             error('Modified Newton: failed to obtain a PD Hessian after %d attempts.', jmax)
64             ;
65         end
66
67         % Calculation of the modified Newton direction via Cholesky factorization
68         y = L \ (-gradfk);
69         pk = L' \ y;
70
71
72         % Backtracking line search to satisfy the Armijo condition
73         alpha = 1;
74         xnew = xk + alpha * pk;
75         fnew = f(xnew);
76         bt = 0;
77         while bt < btmax && fnew > farmijo(fk, alpha, gradfk, pk)
78             alpha = rho * alpha;
79             xnew = xk + alpha * pk;
80             fnew = f(xnew);
81             bt = bt + 1;
82         end
83
84
85         % Update iterate
86         xk = xnew(:);
87         fk = fnew;
88         gradfk = gradf(xk);
89         gradfk_norm = norm(gradfk);
90         hessfk = Hessf(xk);
91
92
93         k = k + 1;
94         xseq(:, k) = xk;
95         btseq(k) = bt;
96     end
97
98
99     % Resize sequences based on the actual number of iterations performed
100    xseq = xseq(:, 1:k);
101    btseq = btseq(1:k);
102
103
104    % Calculation of errors for estimating the convergence order
105    x_star = xseq(:, end); % The last iterate is considered the "true" solution
106

```

```

107 e = vecnorm(xseq - x_star, 2, 1);
108 % Vectorized estimation of the convergence order (if at least 3 iterations)
109 eps_est = 1e-12; % threshold to avoid division by zero
110 if k >= 3
111     valid = (e(1:end-2) > eps_est) & (e(2:end-1) > eps_est) & (e(3:end) > eps_est);
112     p_est = log(e(3:end)./e(2:end-1)) ./ log(e(2:end-1)./e(1:end-2));
113     if any(valid)
114         rate_convergence = mean(p_est(valid));
115     else
116         rate_convergence = NaN;
117     end
118 else
119     rate_convergence = NaN;
120 end
121
122 end

```

Listing 4: Modified Newton

```

1 function [xk, fk, gradfk_norm, k, xseq, btseq, rate_convergence] =
2 modified_newton_matrixfree(x0, f, gradf, kmax, tolgrad, c1, rho, btmax)
3 % This version exploits the fact that the Hessian is:
4 % H(x) = a*I + 4*x*x', with a = 1/100000 + 2*(sum(x.^2)-0.25)
5 % If H is not PD (Positive Definite), i.e., if a <= 0, a shift tau is added such that a+
6 % tau > delta.
7
8 xk = x0(:);
9 n = numel(xk);
10 xseq = zeros(n, kmax);
11 btseq = zeros(1, kmax);
12
13 fk = f(xk);
14 gradfk = gradf(xk);
15 gradfk_norm = norm(gradfk);
16
17 k = 0;
18 while (k < kmax) && (gradfk_norm >= tolgrad)
19     % Calculation of the parameter "a"
20     a = 1/100000 + 2*(sum(xk.^2)-0.25);
21     delta = 1e-6; % small offset to ensure PD (Positive Definiteness)
22     tau = 0;
23     if a <= delta
24         tau = delta - a; % minimum shift to ensure a+tau = delta > 0
25     end
26     A = a + tau; % the scalar part (PD) of the matrix
27
28     % The modified matrix is:
29     % B = A*I + 4*xk*xk'
30     % We want to solve B*p = -gradfk.
31     % We use the Sherman-Morrison formula to solve this system efficiently.
32     %
33     % Let z = -gradfk. Then, by Sherman-Morrison formula:
34     % p = z/A - (4/A^2 * xk * (xk'*z))/(1 + 4*(xk'*xk)/A)
35     z = -gradfk;
36     denom = 1 + 4*(xk'*xk)/A;
37     p = z/A - (4/A^2) * xk * (xk'*z) / denom;
38
39     % Backtracking line search (Armijo condition)
40     alpha = 1;
41     bt = 0;
42     while (bt < btmax)
43         xnew = xk + alpha*p;
44         fnew = f(xnew);
45         if fnew <= fk + c1 * alpha * (gradfk' * p)
46             break;
47         end
48         alpha = rho * alpha;
49         bt = bt + 1;

```

```

48     end
49
50     % Iterate update
51     xk = xnew;
52     fk = fnew;
53     gradfk = gradf(xk);
54     gradfk_norm = norm(gradfk);
55
56     k = k + 1;
57     xseq(:, k) = xk;
58     btseq(k) = bt;
59 end
60 xseq = xseq(:, 1:k);
61 btseq = btseq(1:k);
62
63 % Convergence rate calculation (vectorized estimation)
64 x_star = xseq(:, end);
65 e = vecnorm(xseq - x_star, 2, 1);
66 valid = (e(1:end-2) > 1e-12) & (e(2:end-1) > 1e-12) & (e(3:end) > 1e-12);
67 if any(valid)
68     p_est = log(e(3:end)./e(2:end-1)) ./ log(e(2:end-1)./e(1:end-2));
69     rate_convergence = mean(p_est(valid));
70 else
71     rate_convergence = NaN;
72 end
73 end

```

Listing 5: Modified Newton Hessian Free

C Appendix: Truncated Newton Method

```

1 function [xk, fk, gradfk_norm, k, xseq, btseq, rate_convergence] = truncated_newton(x0, f,
2     gradf, Hessf, kmax, tolgrad, c1, rho, btmax)
% TRUNCATED_NEWTON Truncated Newton method with backtracking (Armijo) without
% preconditioning
3 %
4 % [xk, fk, gradfk_norm, k, xseq, btseq, convergence_rate] = truncated_newton(x0, f, gradf,
5 %     Hessf, ...
6 %                                         kmax, tolgrad, c1, rho, btmax)
%
7 % This function implements the Truncated Newton method for unconstrained optimization
8 % of a function f: R^n -> R. It uses a backtracking line search based on the Armijo
% condition
9 % to ensure sufficient decrease in the function value at each iteration.
10 % The Hessian-vector product is approximated using Conjugate Gradient, making it suitable
% for
11 % large-scale problems where forming and storing the full Hessian is impractical.
12 % Preconditioning is NOT used in this version.
13 %
14 % INPUTS:
15 % x0          - Starting point for the optimization (column vector)
16 % f           - Function handle for the objective function f: R^n -> R
17 % gradf       - Function handle for the gradient of f
18 % Hessf       - Function handle for the Hessian of f (used for Hessian-vector product in
% CG)
19 % kmax        - Maximum number of iterations for the outer Newton loop
20 % tolgrad     - Tolerance for the gradient norm (stopping criterion)
21 % c1          - Parameter for the Armijo condition (in (0,1), e.g., 1e-4)
22 % rho         - Reduction factor for backtracking line search (in (0,1), e.g., 0.5)
23 % btmax       - Maximum number of iterations for the backtracking line search
24 %
25 % OUTPUTS:
26 % xk          - The last computed iterate (approximation of the minimizer)
27 % fk           - Function value f(xk) at the final iterate
28 % gradfk_norm - Norm of the gradient at xk, ||gradf(xk)||

```

```

29 % k - Total number of iterations performed by the Newton method
30 % xseq - Sequence of iterates generated by the method (each column is an
31 % iterate)
31 % btseq - Sequence of backtracking iterations performed at each Newton iteration
32 % rate_convergence- Estimate of the rate of convergence (average of estimated orders in
32 % the last iterations)
33
34
35 %% Input Validation
36 % Ensure the starting point x0 is a column vector
37 x0 = x0(:);
38 % Get the dimension of the problem from the starting point
39 n = numel(x0);
40
41
42 %% Pre-allocation for iteration history
43 % Initialize an array to store the sequence of iterates
44 xseq = zeros(n, kmax);
45 % Initialize an array to store the number of backtracking steps at each iteration
46 btseq = zeros(1, kmax);
47
48
49 %% Initialization
50 % Set the initial iterate to the starting point
51 xk = x0;
52 % Evaluate the objective function at the starting point
53 fk = f(xk);
54 % Evaluate the gradient at the starting point
55 gradfk = gradf(xk);
56 % Compute the norm of the gradient at the starting point
57 gradfk_norm = norm(gradfk);
58 % Evaluate the Hessian at the starting point (for use in CG)
59 hessfk = Hessf(xk);
60 % Initialize iteration counter
61 k = 0;
62
63
64 % Main loop of the Truncated Newton method
65 while (k < kmax) && (gradfk_norm >= tolgrad)
66     % Dynamic tolerance for Conjugate Gradient based on current gradient norm
67     etak = min(0.5, gradfk_norm);
68
69     % Compute the search direction pk by approximately solving Hessf(xk)*pk = -gradf(xk)
70     % using Conjugate Gradient method (cg_curvtrun_newt function, assumed to be defined
70     % elsewhere)
71     pk = cg_curvtrun_newt(hessfk, -gradfk, zeros(n,1), etak);
72
73     % Check if the computed direction is a descent direction
74     % If not (i.e., not a descent direction), fall back to steepest descent direction (-
74     % gradf(xk))
75     if gradfk' * pk >= 0
76         warning('Computed direction is not a descent direction. Using steepest descent
76         direction: -gradf.');
77         pk = -gradf(xk);
78     end
79
80     % Backtracking line search (Armijo condition)
81     % Initialize step size to 1
82     alpha = 1;
83     % Compute the Armijo condition target value
84     curr_armijo = fk + c1 * alpha * (gradfk' * pk);
85     % Calculate the new point xnew and function value fnew with current step size
86     xnew = xk + alpha * pk;
87     fnew = f(xnew);
88     % Initialize backtracking iteration counter
89     bt = 0;
90     % Backtracking loop: reduce step size until Armijo condition is satisfied or max
90     % backtrack iterations reached

```

```

91    while (bt < btmax) && (fnew > curr_armijo)
92        % Reduce step size by factor rho
93        alpha = rho * alpha;
94        % Update Armijo condition target value with reduced step size
95        curr_armijo = fk + c1 * alpha * (gradfk' * pk);
96        % Recalculate xnew and fnew with updated step size
97        xnew = xk + alpha * pk;
98        fnew = f(xnew);
99        % Increment backtracking iteration counter
100       bt = bt + 1;
101   end
102
103   % Warning if maximum backtracking iterations are reached without satisfying Armijo
104   % condition
105   if bt == btmax && fnew > curr_armijo
106       warning('Maximum backtracking iterations (%d) reached at iteration %d: Armijo
107           condition not satisfied.', btmax, k+1);
108   end
109
110   % Update variables for the next iteration
111   % New iterate xk is set to xnew
112   xk = xnew;
113   % Function value at new iterate
114   fk = fnew;
115   % Gradient at new iterate
116   gradfk = gradf(xk);
117   % Gradient norm at new iterate
118   gradfk_norm = norm(gradfk);
119   % Hessian at new iterate (for use in next CG iteration)
120   hessfk = Hessf(xk);
121
122   % Increment iteration counter
123   k = k + 1;
124   % Store current iterate in xseq
125   xseq(:, k) = xk;
126   % Store number of backtracking steps in btseq
127   btseq(k) = bt;
128 end
129
130 % Resize output arrays to remove pre-allocated space
131 xseq = xseq(:, 1:k);
132 btseq = btseq(1:k);
133
134 % Estimate rate of convergence (vectorized calculation for efficiency)
135 x_star = xseq(:, end); % Assume last iterate is close to x*
136 e = vecnorm(xseq - x_star, 2, 1); % Error at each iteration
137 % Find indices where errors are sufficiently larger than machine precision to avoid log(0)
138 % or division by zero
139 valid = (e(1:end-2) > 1e-12) & (e(2:end-1) > 1e-12) & (e(3:end) > 1e-12);
140 if any(valid)
141     % Estimate convergence order using errors from consecutive iterations
142     p_est = log(e(3:end)./e(2:end-1)) ./ log(e(2:end-1)./e(1:end-2));
143     % Average the estimated orders for a single rate of convergence value
144     rate_convergence = mean(p_est(valid));
145 else
146     rate_convergence = NaN; % Not a Number if rate cannot be reliably estimated
147 end

```

Listing 6: Truncated Newton Without Preconditioning

```

1 function [xk, fk, gradfk_norm, k, xseq, btseq, rate_convergence] =
2     truncated_newton_hessian_free(x0, f, gradf, kmax, tolgrad, c1, rho, btmax)
% TRUNCATED_NEWTON_HESSIAN_FREE Hessian-Free Truncated Newton method with backtracking (
3     % Armijo) and no preconditioning

```



```

63
64 while (k < kmax) && (gradfk_norm >= tolgrad)
65 % Dynamic tolerance for Conjugate Gradient based on current gradient norm
66 etak = min(0.5, gradfk_norm);
67
68 % Define the Hessian-vector product operator approximated using finite differences
69 % Hv(p) approximates Hessian(f(xk)) * p
70 Hv = @(p) (gradf(xk + sigma * p) - gradf(xk)) / sigma;
71
72 % Compute the search direction pk by approximately solving Hv(pk) = -gradf(xk)
73 % using Conjugate Gradient method (cg_curvtrun_hessian_free function, defined below)
74 pk = cg_curvtrun_hessian_free(Hv, -gradfk, zeros(n,1), etak);
75
76 % Descent direction check: ensure the computed direction is a descent direction
77 % If not, fall back to steepest descent direction (-gradf(xk))
78 if gradfk' * pk >= 0
79     warning('Computed direction is not a descent direction. Using steepest descent
80         direction: -gradf.');
81     pk = -gradf(xk);
82 end
83
84 % Backtracking line search (Armijo condition)
85 % Initialize step size to 1
86 alpha = 1;
87 % Compute the Armijo condition target value
88 curr_armijo = fk + c1 * alpha * (gradfk' * pk);
89 % Calculate the new point xnew and function value fnew with current step size
90 xnew = xk + alpha * pk;
91 fnew = f(xnew);
92 % Initialize backtracking iteration counter
93 bt = 0;
94 % Backtracking loop: reduce step size until Armijo condition is satisfied or max
95 % backtrack iterations reached
96 while (bt < btmax) && (fnew > curr_armijo)
97     % Reduce step size by factor rho
98     alpha = rho * alpha;
99     % Update Armijo condition target value with reduced step size
100    curr_armijo = fk + c1 * alpha * (gradfk' * pk);
101    % Recalculate xnew and fnew with updated step size
102    xnew = xk + alpha * pk;
103    fnew = f(xnew);
104    % Increment backtracking iteration counter
105    bt = bt + 1;
106 end
107
108 % Warning if maximum backtracking iterations are reached without satisfying Armijo
109 % condition
110 if bt == btmax && fnew > curr_armijo
111     warning('Maximum backtracking iterations (%d) reached at iteration %d: Armijo
112         condition not satisfied.', btmax, k+1);
113 end
114
115 % Update variables for the next iteration
116 % New iterate xk is set to xnew
117 xk = xnew;
118 % Function value at new iterate
119 fk = fnew;
120 % Gradient at new iterate
121 gradfk = gradf(xk);
122 % Gradient norm at new iterate
123 gradfk_norm = norm(gradfk);
124
125 % Increment iteration counter
126 k = k + 1;
127 % Store current iterate in xseq
128 xseq(:, k) = xk;
129 % Store number of backtracking steps in btseq
130 btseq(k) = bt;

```

```

127 end
128
129
130 %% Output Processing
131 % Resize output arrays to remove pre-allocated space
132 xseq = xseq(:, 1:k);
133 btseq = btseq(1:k);
134
135
136 % Estimate rate of convergence (vectorized calculation for efficiency)
137 x_star = xseq(:, end); % Assume last iterate is close to x*
138 e = vecnorm(xseq - x_star, 2, 1); % Error at each iteration
139 % Find indices where errors are sufficiently larger than machine precision to avoid log(0)
    or division by zero
140 valid = (e(1:end-2) > 1e-12) & (e(2:end-1) > 1e-12) & (e(3:end) > 1e-12);
141 if any(valid)
142     % Estimate convergence order using errors from consecutive iterations
143     p_est = log(e(3:end)./e(2:end-1)) ./ log(e(2:end-1)./e(1:end-2));
144     % Average the estimated orders for a single rate of convergence value
145     rate_convergence = mean(p_est(valid));
146 else
147     rate_convergence = NaN; % Not a Number if rate cannot be reliably estimated
148 end
149 end

```

Listing 7: Truncated Newton Without Preconditioning Hessian Free

```

1 function xk = cg_curvtrun_newt(A, b, x0, tol)
2 % CG_CURVTRUN_NEWT Solves the system A*x = b using the truncated CG method
3 % (without preconditioning) handling potential negative curvature cases.
4 %
5 %   xk = cg_curvtrun_newt(A, b, x0, tol)
6 %
7 % INPUTS:
8 %   A      - Matrix (Hessian) of the system
9 %   b      - Right-hand side vector (typically -gradf)
10 %  x0     - Initial vector
11 %  tol    - Tolerance on the relative residual
12 %
13 % OUTPUT:
14 %  xk    - Approximate solution (computed direction)
15
16
17 % Initializations
18 xk = x0;
19 rk = b - A * xk;
20 pk = rk;
21 norm_b = norm(b);
22 if norm_b == 0
23     norm_b = 1; % Avoid division by zero if b is zero vector
24 end
25 relres = norm(rk) / norm_b;
26 kmax = 100;
27 k = 0;
28
29
30 while (relres > tol) && (k < kmax)
31     Apk = A * pk;
32
33     % If negative or near-zero curvature is detected, terminate the loop
34     if pk' * Apk <= 0
35         if k == 0
36             % If negative curvature is detected at the first step,
37             % proceed with a step calculated along pk
38             alphak = (rk' * pk) / max(pk' * Apk, eps);
39             xk = xk + alphak * pk;
40         end
41         break;

```

```

42    end
43
44    % Calculate step size using the classical CG method formula
45    alphak = (rk' * rk) / (pk' * Apk);
46    xk = xk + alphak * pk;
47
48    rk_new = rk - alphak * Apk;
49    beta = (rk_new' * rk_new) / (rk' * rk);
50
51    % Update the search direction
52    pk = rk_new + beta * pk;
53
54    rk = rk_new;
55    relres = norm(rk) / norm_b;
56    k = k + 1;
57 end
58
59
60 end

```

Listing 8: CG Without Preconditioning

```

1 function xk = cg_curvtrun_hessian_free(Hv, b, x0, tol)
2 % CG_CURVTRUN_HESSIAN_FREE Solves the system Hv(x) = b using the Hessian-Free truncated CG
3 % method
4 % (without preconditioning) handling potential negative curvature cases.
5 % Hessian-Free Version: Hv is a function handle that computes the
6 % Hessian-vector product, NOT the explicit Hessian matrix.
7 %
8 %     xk = cg_curvtrun_hessian_free(Hv, b, x0, tol)
9 %
10 % INPUTS:
11 %     Hv - Function handle for the Hessian-vector product, Hv(p) = Hessian(f(xk)) * p
12 %     b - Right-hand side vector (typically -gradf)
13 %     x0 - Initial vector
14 %     tol - Tolerance on the relative residual
15 %
16 % OUTPUT:
17 %     xk - Approximate solution (computed direction)
18
19 % Initializations
20 xk = x0;
21 rk = b - Hv(xk); % Use Hv to calculate the Hessian-vector product
22 pk = rk;
23 norm_b = norm(b);
24 if norm_b == 0
25     norm_b = 1; % Avoid division by zero if b is zero vector
26 end
27 relres = norm(rk) / norm_b;
28 kmax = 100;
29 k = 0;
30
31
32 while (relres > tol) && (k < kmax)
33     Apk = Hv(pk); % Use Hv to calculate the Hessian-vector product
34
35     % Negative curvature check
36     if pk' * Apk <= 0
37         if k == 0
38             % If negative curvature is detected at the first step,
39             % proceed with a step calculated along pk
40             alphak = (rk' * pk) / max(pk' * Apk, eps);
41             xk = xk + alphak * pk;
42         end
43         break;
44     end

```

```

46 % Calculate step size using the classical CG method formula
47 alphak = (rk' * rk) / (pk' * Apk);
48 xk = xk + alphak * pk;
49
50 rk_new = rk - alphak * Apk;
51 beta = (rk_new' * r_new) / (rk' * rk);
52
53 % Update the search direction
54 pk = rk_new + beta * pk;
55
56 rk = r_new;
57 relres = norm(rk) / norm_b;
58 k = k + 1;
59 end
60
61
62 end

```

Listing 9: CG Without Preconditioning Hessian Free

```

1 function [xk, fk, gradfk_norm, k, xseq, btseq, rate_convergence] = truncated_newton_pre(x0,
2 f, gradf, Hessf, kmax, tolgrad, c1, rho, btmax, varargin)
% TRUNCATED_NEWTON_PRE - Truncated Newton method with backtracking and preconditioning.
3 % Optimized for large-scale problems and sparse Hessian matrices.
4 %
5 % [xk, fk, gradfk_norm, k, xseq, btseq, convergence_rate] = truncated_newton_pre(x0, f,
6 % gradf, Hessf, kmax, tolgrad, c1, rho, btmax, varargin)
%
7 % This function implements a Preconditioned Truncated Newton method for unconstrained
8 % optimization
9 % of a function f: R^n -> R. It is specifically designed for large-scale problems where the
10 % Hessian
11 % matrix is sparse, and employs preconditioning to accelerate convergence. A backtracking
12 % line search
13 % based on the Armijo condition is used to ensure sufficient decrease in the function value
14 % at each iteration.
15 % The Hessian-vector product within the Conjugate Gradient (CG) solver is used, and a
16 % preconditioner
17 % can be applied within the CG iterations (though the provided code for cg_curvtrun_newt_pre
18 % is not given,
19 % it's assumed that preconditioning is applied there).
%
20 % INPUTS:
21 % x0          - Starting point for the optimization (column vector)
22 % f           - Function handle for the objective function f: R^n -> R
23 % gradf       - Function handle for the gradient of f
24 % Hessf       - Function handle that, given x, returns the sparse Hessian matrix, or
25 %                 a constant sparse matrix if the Hessian is independent of x.
26 % kmax        - Maximum number of iterations for the outer Newton loop
27 % tolgrad     - Tolerance for the gradient norm (stopping criterion)
28 % c1          - Parameter for the Armijo condition (typically 1e-4)
29 % rho         - Step reduction factor for backtracking (0 < rho < 1, e.g., 0.5)
30 % btmax       - Maximum number of step reductions in backtracking line search
31 % varargin    - (Optional) Variable input arguments; may be used to pass preconditioning
32 %                 matrix
33 %                 or parameters to the cg_curvtrun_newt_pre function (not directly used in
34 %                 this code snippet,
35 %                 but included for potential future extensibility).
%
36 % OUTPUTS:
37 % xk          - The last computed iterate (approximation of the minimizer)
38 % fk           - Function value f(xk) at the final iterate
39 % gradfk_norm - Norm of the gradient at xk, ||gradf(xk) ||
40 % k            - Total number of iterations performed by the Newton method
41 % xseq         - Sequence of iterates generated by the method (each column is an
42 %                 iterate)
43 % btseq        - Sequence of backtracking iterations performed at each Newton iteration

```

```

37 %     rate_convergence- Estimate of the rate of convergence (average of estimated orders in
38 %     the last iterations)
39
40 %% Input Preparation
41 % Ensure the starting point x0 is a column vector
42 x0 = x0(:);
43 % Get the dimension of the problem from the starting point
44 n = numel(x0);
45
46
47 %% Pre-allocation for iteration history
48 % Initialize an array to store the sequence of iterates
49 xseq = zeros(n, kmax);
50 % Initialize an array to store the number of backtracking steps at each iteration
51 btseq = zeros(1, kmax);
52
53
54 %% Initialization
55 % Set the initial iterate to the starting point
56 xk = x0;
57 % Evaluate the objective function at the starting point
58 fk = f(xk);
59 % Evaluate the gradient at the starting point
60 gradfk = gradf(xk);
61 % Compute the norm of the gradient at the starting point
62 gradfk_norm = norm(gradfk);
63
64
65 % Check if Hessf is a function handle or a matrix
66 if isa(Hessf, 'function_handle')
67     hessfk = Hessf(xk); % If function handle, evaluate Hessian at xk (expecting a sparse
68     % matrix)
69 else
70     hessfk = Hessf; % If not a function handle, assume Hessf is already the sparse
71     % Hessian matrix (constant)
72 end
73
74 % Initialize iteration counter
75 k = 0;
76 % Define a function handle for the Armijo condition right-hand side for efficiency
77 armijo_rhs = @(fk, alpha, gradfk, pk) fk + c1 * alpha * (gradfk' * pk);
78
79 % Main loop of the Preconditioned Truncated Newton method
80 while (k < kmax) && (gradfk_norm >= tolgrad)
81     % Dynamic tolerance for Conjugate Gradient based on current gradient norm
82     etak = min(0.5, gradfk_norm);
83
84     % Solve the system Hessf(xk)*p = -gradf(xk) using Preconditioned Truncated CG
85     % (cg_curvtrun_newt_pre function, assumed to be defined elsewhere and implementing
86     % preconditioning)
87     pk = cg_curvtrun_newt_pre(hessfk, -gradfk, zeros(n,1), etak);
88
89     % -----
90     % Backtracking Line Search (Armijo condition)
91     % -----
92     % Initialize step size to 1
93     alpha = 1;
94     % Calculate the new point xnew with full step size
95     xnew = xk + alpha * pk;
96     % Evaluate the function at the new point
97     fnew = f(xnew);
98     % Initialize backtracking iteration counter
99     bt = 0;
100    % Backtracking loop: reduce step size until Armijo condition is satisfied or max

```

```

    backtrack iterations reached
101 while (bt < btmax) && (fnew > armijo_rhs(fk, alpha, gradfk, pk))
102     % Reduce step size by factor rho
103     alpha = rho * alpha;
104     % Recalculate the new point xnew with reduced step size
105     xnew = xk + alpha * pk;
106     % Re-evaluate the function at the new point
107     fnew = f(xnew);
108     % Increment backtracking iteration counter
109     bt = bt + 1;
110 end
111
112 % Warning if maximum backtracking iterations are reached without satisfying Armijo
113 % condition
113 if (bt == btmax) && (fnew > armijo_rhs(fk, alpha, gradfk, pk))
114     warning('Maximum backtracking iterations (%d) reached at iteration %d: Armijo
115         condition not satisfied.', btmax, k+1);
116 end
117
118 % Update variables for the next iteration
119 % New iterate xk is set to xnew
119 xk = xnew;
120 % Function value at new iterate
121 fk = fnew;
122 % Gradient at new iterate
123 gradfk = gradf(xk);
124 % Gradient norm at new iterate
125 gradfk_norm = norm(gradfk);
126
127 % Update Hessian matrix (if Hessf is a function handle, re-evaluate at new xk)
128 if isa(Hessf, 'function_handle')
129     hessfk = Hessf(xk); % Re-evaluate Hessian at xk if it's a function handle
130 else
131     hessfk = Hessf;      % Keep Hessian as constant matrix if it's not a function handle
132 end
133
134
135 % Increment iteration counter
136 k = k + 1;
137 % Store current iterate in xseq
138 xseq(:, k) = xk;
139 % Store number of backtracking steps in btseq
140 btseq(k) = bt;
141 end
142
143
144 %% Output Processing
145 % Resize output arrays to remove pre-allocated space
146 xseq = xseq(:, 1:k);
147 btseq = btseq(1:k);
148
149
150 % Estimate rate of convergence (vectorized calculation for efficiency)
151 x_star = xseq(:, end); % Assume last iterate is close to x*
152 e = vecnorm(xseq - x_star, 2, 1); % Error at each iteration
153 % Find indices where errors are sufficiently larger than machine precision
154 valid = (e(1:end-2) > 1e-12) & (e(2:end-1) > 1e-12) & (e(3:end) > 1e-12);
155 if any(valid)
156     % Estimate convergence order using errors from consecutive iterations
157     p_est = log(e(3:end)./e(2:end-1)) ./ log(e(2:end-1)./e(1:end-2));
158     % Average the estimated orders for a single rate of convergence value
159     rate_convergence = mean(p_est(valid));
160 else
161     rate_convergence = NaN; % Return NaN if rate cannot be reliably estimated
162 end
163
164
165 end

```

Listing 10: Truncated Newton With Preconditioning

```

1 function [xk, fk, gradfk_norm, k, xseq, btseq, rate_convergence] =
2   truncated_newton_pre_matrixfree(x0, f, gradf, kmax, tolgrad, c1, rho, btmax)
3   % This version solves the system H(x)*p = -grad using a "matrix-free" CG method.
4   % Here H(x) = a*I + 4*x*x' (with 'a' defined as above) and
5   % the product H(x)*v is computed without forming H explicitly.
6
7   xk = x0(:);
8   n = numel(xk);
9   xseq = zeros(n, kmax);
10  btseq = zeros(1, kmax);
11
12  fk = f(xk);
13  gradfk = gradf(xk);
14  gradfk_norm = norm(gradfk);
15  k = 0;
16
17  armijo_rhs = @(fk, alpha, gradfk, pk) fk + c1 * alpha * (gradfk' * pk);
18
19  while (k < kmax) && (gradfk_norm >= tolgrad)
20    % Calculate 'a' and define the matrix-free Hessian-vector product
21    a = 1/100000 + 2*(sum(xk.^2)-0.25);
22    % We do not use a shift here, as truncated CG handles potential indefiniteness
23    Afun = @(v) (a*v + 4 * xk * (xk' * v)); % Matrix-free Hessian-vector product: H(x) *
24    v
25
26    % Choose an inner tolerance parameter for CG
27    etak = min(0.5, gradfk_norm);
28    % Solve H*p = -gradfk using matrix-free truncated CG
29    p = cg_truncated_matrix_free(Afun, -gradfk, zeros(n,1), etak);
30
31    % Backtracking line search (Armijo condition)
32    alpha = 1;
33    bt = 0;
34    while (bt < btmax)
35      xnew = xk + alpha * p;
36      fnew = f(xnew);
37      if fnew <= armijo_rhs(fk, alpha, gradfk, p)
38        break;
39      end
40    alpha = rho * alpha;
41    bt = bt + 1;
42  end
43  if (bt == btmax) && (fnew > armijo_rhs(fk, alpha, gradfk, p))
44    warning('Maximum backtracking iterations (%d) reached at iteration %d: Armijo
45          condition not satisfied.', btmax, k+1);
46  end
47
48  % Iterate update
49  xk = xnew;
50  fk = fnew;
51  gradfk = gradf(xk);
52  gradfk_norm = norm(gradfk);
53
54  k = k + 1;
55  xseq(:, k) = xk;
56  btseq(k) = bt;
57
58  x_star = xseq(:, end);
59  e = vecnorm(xseq - x_star, 2, 1);
60  valid = (e(1:end-2) > 1e-12) & (e(2:end-1) > 1e-12) & (e(3:end) > 1e-12);
61
```

```

62     if any(valid)
63         p_est = log(e(3:end)./e(2:end-1)) ./ log(e(2:end-1)./e(1:end-2));
64         rate_convergence = mean(p_est(valid));
65     else
66         rate_convergence = NaN;
67     end
68 end

```

Listing 11: Truncated Newton With Preconditioning Hessian Free

```

1 function xk = cg_curvtrun_newt_pre(A_f, b, x0, tol)
2 % CG_CURVTRUN_NEWT_PRE Solves A*x = b using truncated CG with sparse preconditioning.
3 %   A_f can be a sparse matrix or a handle that, given a point x, returns
4 %   the Hessian matrix (NOT used here: in truncated_newton_pre we already pass the matrix).
5 %
6 % Preconditioning Strategy:
7 %   - Use ichol as a preconditioner (Incomplete Cholesky) if possible.
8 %   - If ichol fails, fall back to a diagonal preconditioner.
9 %   - If negative curvature is detected ( $p'Ap \leq 0$ ), terminate CG iteration.
10
11 persistent L_ichol_factor % Persistent preconditioner factor
12
13
14
15 % Preconditioner construction only on the first function call
16 if isempty(L_ichol_factor)
17     try
18         opts.michol = 'lower';
19         A0 = sparse(A_f); % In this version, A_f is already a matrix (not a handle)
20         L_ichol_factor = ichol(A0, opts);
21     catch
22         % If ichol fails, use diagonal preconditioning
23         d = diag(A_f);
24         d(d==0) = 1;
25         L_ichol_factor = diag(sqrt(d));
26     end
27 end
28
29
30 L = L_ichol_factor; % Assign the preconditioner factor (either ichol or diagonal) to L
31
32
33 % Initialization
34 xk = x0;
35 % Calculate the initial residual: r = b - A*xk
36 r = b - A_f * xk;
37 % Apply preconditioner: solve L*z = r and L'*precond_z = z, where preconditioner M = L*L'
38 z = L \ r; % Solve L*z = r (forward substitution since L is lower triangular)
39 p = L' \ z; % Solve L'*p = z (backward substitution since L' is upper triangular) - p is
40     % the preconditioned residual
41
42 normb = norm(b);
43 if normb == 0, normb = 1; end % Avoid division by zero if norm of b is zero
44
45
46 relres = norm(r) / normb; % Relative residual
47 n = numel(x0);
48 maxit = min(100, n); % Maximum CG iterations (capped at 100 or problem dimension)
49 it = 0; % Iteration counter
50
51
52 while (relres > tol) && (it < maxit)
53     Ap = A_f * p; % Matrix-vector product A*p
54
55     % Negative curvature check
56     if p' * Ap <= 0

```

```

57 % If negative curvature is detected, especially at the first iteration, take a step
58 % along p
59 if it == 0
60     alpha = (r' * p) / max(p' * Ap, eps); % Step size calculation, ensure
61     denominator is not zero
62     xk = xk + alpha * p; % Update solution with step along p
63 end
64 break; % Exit CG iteration if negative curvature is detected
65 end
66
67 % Standard Conjugate Gradient updates
68 alpha = (r' * p) / (p' * Ap); % Step size
69 xk = xk + alpha * p; % Update solution
70
71 r_new = r - alpha * Ap; % New residual
72 z_new = L \ r_new; % Apply preconditioner to new residual: solve L*z_new =
73 r_new
74 q = L' \ z_new; % Solve L'*q = z_new - q is the preconditioned new
75 residual
76
77 beta = (r_new' * q) / (r' * p); % Calculate beta for direction update (Fletcher-Reeves
78 version)
79 p = q + beta * p; % Update search direction using preconditioned residuals
80
81 r = r_new; % Update residual
82 relres = norm(r) / normb; % Update relative residual norm
83 it = it + 1; % Increment iteration counter
84
85 end
86
87 end

```

Listing 12: CG With Preconditioning

```

1 function xk = cg_truncated_matrix_free(Afun, b, x0, tol)
2 % CG_TRUNCATED_MATRIX_FREE Solves the system Afun(x) = b using the matrix-free truncated CG
3 % method
4 % (without preconditioning) handling potential negative curvature cases.
5 % In this Hessian-free version, Afun is a function handle that computes the
6 % matrix-vector product A*v, WITHOUT explicitly forming the matrix A.
7 %
8 % xk = cg_truncated_matrix_free(Afun, b, x0, tol)
9 %
10 % INPUTS:
11 %   Afun - Function handle for the matrix-vector product, Afun(v) = A * v
12 %         where A is the system matrix (e.g., Hessian approximation).
13 %         This function should compute the product WITHOUT explicitly forming A.
14 %   b - Right-hand side vector (typically -gradf)
15 %   x0 - Initial vector for the CG iteration
16 %   tol - Tolerance on the relative residual norm (stopping criterion)
17 %
18 % OUTPUT:
19 %   xk - Approximate solution vector (computed direction)
20
21 xk = x0; % Initialize the solution vector with the starting vector
22 r = b - Afun(xk); % Calculate the initial residual: r = b - Afun(xk)
23 normb = norm(b); % Calculate the norm of the right-hand side vector b
24 if normb == 0, normb = 1; end % Avoid division by zero if norm of b is zero
25 relres = norm(r) / normb; % Calculate the initial relative residual
26 p = r; % Initialize the search direction p with the initial residual
27 maxit = 100; % Maximum number of iterations for CG (modify if needed)
28 it = 0; % Initialize iteration counter
29
30 while (relres > tol) && (it < maxit)
31     Ap = Afun(p); % Compute the matrix-vector product Ap = A * p using the provided
32     % function handle Afun
33
34     % Negative curvature check: if p'*Ap <= 0, terminate CG loop

```

```

33     if p' * Ap <= 0
34         if it == 0
35             % If negative curvature is detected at the first iteration,
36             % proceed with a step calculated along p
37             alpha = (r' * p) / max(p' * Ap, eps); % Step size calculation, ensure
38             denominator is not zero
39             xk = xk + alpha * p; % Update solution with step along p
40         end
41         break; % Exit CG iteration if negative curvature is detected
42     end
43
44     alpha = (r' * r) / (p' * Ap); % Calculate step size alpha for CG update
45     xk = xk + alpha * p; % Update solution vector xk
46     r_new = r - alpha * Ap; % Calculate the new residual r_new = r - alpha * A*p
47     beta = (r_new' * r_new) / (r' * r); % Calculate beta for updating the search
48     direction (Fletcher-Reeves version)
49     p = r_new + beta * p; % Update search direction p
50     r = r_new; % Update residual r to r_new
51     relres = norm(r) / normb; % Update relative residual norm
52     it = it + 1; % Increment iteration counter
53
54 end
55
56 end

```

Listing 13: CG With Preconditioning Matrix Free

D Appendix: Finite Differences

```

1 function grad = findiff_grad(F, x, h, scheme)
2 % FINDIFF_GRAD Approximates the gradient of a function using finite differences.
3 %
4 % grad = findiff_grad(F, x, h, scheme)
5 %
6 % This function calculates an approximation of the gradient of a scalar function F
7 % at a point x using finite difference methods. It supports central and forward difference
8 % schemes.
9 %
10 % INPUTS:
11 %   F      - Function handle of the objective function, F: R^n -> R.
12 %   x      - Point (vector) at which to compute the gradient approximation.
13 %   h      - Step size for the finite difference approximation.
14 %          - Can be a scalar (uniform step size for all components)
15 %          - or a vector of length n (step size for each component).
16 %   scheme - String specifying the finite difference scheme to use:
17 %          - 'c' for central difference (second-order accurate)
18 %          - 'fw' for forward difference (first-order accurate)
19 %
20 % OUTPUT:
21 %   grad    - Approximated gradient of F at x (column vector).
22
23 n = length(x); % Determine the dimension of the input vector x
24 grad = zeros(n,1); % Initialize the gradient vector with zeros
25
26 % Loop through each component to compute the partial derivatives
27 for i = 1:n
28     x_plus = x; % Create a copy of x for the positive perturbation
29     x_minus = x; % Create a copy of x for the negative perturbation
30
31     % Determine the step size hi for the current component i
32     % If h is a vector, use the i-th component h(i) as step size for the i-th partial
33     % derivative
34     if numel(h) > 1
35         hi = h(i);
36     else
37         hi = h; % If h is a scalar, use the same step size for all components
38     end
39
40     % Compute the positive and negative perturbations
41     x_plus(i) = x(i) + hi;
42     x_minus(i) = x(i) - hi;
43
44     % Compute the function values
45     f_plus = F(x_plus);
46     f_minus = F(x_minus);
47
48     % Compute the partial derivative
49     grad(i) = (f_plus - f_minus) / (2 * hi);
50
51 end

```

```

37     x_plus(i) = x_plus(i) + hi;    % Perturb the i-th component of x by +hi
38     x_minus(i) = x_minus(i) - hi; % Perturb the i-th component of x by -hi
39
40
41     % Calculate the finite difference approximation based on the chosen scheme
42     if strcmp(scheme,'c')
43         % Central difference scheme: (F(x + hi*e_i) - F(x - hi*e_i)) / (2*hi)
44         grad(i) = (F(x_plus) - F(x_minus))/(2*hi);
45     elseif strcmp(scheme,'fw')
46         % Forward difference scheme: (F(x + hi*e_i) - F(x)) / hi
47         grad(i) = (F(x_plus) - F(x))/(hi);
48     else
49         error('Unrecognized finite difference scheme. Please use ''c'' for central or ''fw'' for forward.');
50     end
51 end
52

```

Listing 14: Gradient with Finite Differences

```

1 function [Hessfx] = findiff_Hess(f, x, h)
2 % FINDIFF_HESS Approximates the Hessian of a function using finite differences.
3 %
4 % [Hessfx] = findiff_Hess(f, x, h)
5 %
6 % Optimized function to approximate the Hessian matrix of the function f at the point x (
7 % column vector)
8 % using finite difference methods. This version supports both scalar step size h and a
9 % vector step size h,
10 % where h(i) is used as the step size for the i-th coordinate.
11 %
12 % This implementation assumes the Hessian matrix is sparse with nonzero entries primarily
13 % on the main diagonal, the first sub- and super-diagonals, and the two corner elements
14 % at positions (1,n) and (n,1). This specific sparsity pattern is exploited for efficiency,
15 % calculating only these potentially nonzero entries.
16 %
17 % INPUTS:
18 %   f      - Function handle describing a function f: R^n -> R;
19 %   x      - n-dimensional column vector representing the point at which to approximate the
20 %           Hessian.
21 %   h      - Step size for the finite difference approximation. This can be:
22 %           - A scalar: a uniform step size applied to all components.
23 %           - An n-dimensional vector: component-specific step sizes, where h(i) is used for
24 %             the i-th coordinate.
25 %
26 % OUTPUT:
27 %   Hessfx - n-by-n sparse matrix representing the approximate Hessian of f at x,
28 %             calculated using finite differences and exploiting the assumed sparsity
29 %             pattern.
30 %
31
32 n = length(x); % Get the dimension of the input vector x
33 f0 = f(x); % Evaluate the function f at the point x (base function value)
34
35
36
37 % Preallocate arrays to store function evaluations at perturbed points.
38 f_plus = zeros(n,1); % Array to store f(x + h_i*e_i) for each coordinate i
39 f_minus = zeros(n,1); % Array to store f(x - h_i*e_i) for each coordinate i
40
41
42 % Compute f(x + h_i*e_i) and f(x - h_i*e_i) for each coordinate direction e_i.
43 for i = 1:n
44     % Determine step size hi for the i-th coordinate
45     if numel(h) > 1
46         hi = h(i); % Use component-specific step size if h is a vector
47     else
48         hi = h; % Use uniform step size if h is a scalar
49
50     x_plus(i) = x_plus(i) + hi;    % Perturb the i-th component of x by +hi
51     x_minus(i) = x_minus(i) - hi; % Perturb the i-th component of x by -hi
52
53     % Calculate the finite difference approximation based on the chosen scheme
54     if strcmp(scheme,'c')
55         % Central difference scheme: (F(x + hi*e_i) - F(x - hi*e_i)) / (2*hi)
56         grad(i) = (F(x_plus) - F(x_minus))/(2*hi);
57     elseif strcmp(scheme,'fw')
58         % Forward difference scheme: (F(x + hi*e_i) - F(x)) / hi
59         grad(i) = (F(x_plus) - F(x))/(hi);
60     else
61         error('Unrecognized finite difference scheme. Please use ''c'' for central or ''fw'' for forward.');
62     end
63 end
64
65 % Compute the Hessian matrix using the finite difference approximations
66 % and the assumed sparsity pattern.
67 % ...
68
69 % Return the approximate Hessian matrix
70 % ...
71
72 % End of function
73

```

```

44    end
45    x_temp = x; % Create a temporary copy of x to avoid modifying the original x
46    x_temp(i) = x_temp(i) + hi; % Perturb the i-th component by +hi
47    f_plus(i) = f(x_temp); % Evaluate f at the positively perturbed point
48
49    x_temp = x; % Reset x_temp back to x for the negative perturbation
50    x_temp(i) = x_temp(i) - hi; % Perturb the i-th component by -hi
51    f_minus(i) = f(x_temp); % Evaluate f at the negatively perturbed point
52 end
53
54
55 % Preallocate the Hessian matrix as a sparse matrix to efficiently store and compute only
56 % nonzero elements.
56 Hessfx = spalloc(n, n, 4*n); % Allocate space for a sparse n x n matrix, expecting
57 % approximately 4*n non-zero elements (based on sparsity assumption)
58
59 % Compute diagonal entries of the Hessian using second-order central finite differences.
60 for i = 1:n
61     % Determine step size hi for the i-th diagonal entry
62     if numel(h) > 1
63         hi = h(i); % Use component-specific step size if h is a vector
64     else
65         hi = h; % Use uniform step size if h is a scalar
66     end
67     % Second-order central difference formula for diagonal elements:
68     %  $H_{ii} = (f(x + h_i e_i) - 2*f(x) + f(x - h_i e_i)) / (h_i^2)$ 
69     Hessfx(i,i) = (f_plus(i) - 2*f0 + f_minus(i)) / (hi^2);
70 end
71
72
73 % Compute off-diagonal entries for adjacent variables (first sub- and super-diagonals).
74 for i = 1:n-1
75     % Determine step sizes hi and hj for the off-diagonal entry (i, i+1)
76     if numel(h) > 1
77         hi = h(i); % Step size for the i-th component
78         hj = h(i+1); % Step size for the (i+1)-th component
79     else
80         hi = h; % Use uniform step size if h is a scalar
81         hj = h; % Use uniform step size if h is a scalar
82     end
83     x_temp = x; % Create a temporary copy of x
84     x_temp(i) = x_temp(i) + hi; % Perturb the i-th component by +hi
85     x_temp(i+1) = x_temp(i+1) + hj; % Perturb the (i+1)-th component by +hj
86     f_pair = f(x_temp); % Evaluate f at the point perturbed in both i-th and (i+1)-th
87     % directions
88
89     % Finite difference formula for off-diagonal elements (i, i+1) and (i+1, i):
90     %  $H_{ii+1} = H_{ii+1} = (f(x + h_i e_i + h_j e_j) - f(x + h_i e_i) - f(x + h_j e_j) + f(x)) / (h_i * h_j)$ 
91     value = (f_pair - f_plus(i) - f_plus(i+1) + f0) / (hi * hj);
92     Hessfx(i, i+1) = value; % Assign the computed value to the (i, i+1) entry
93     Hessfx(i+1, i) = value; % Assign the same value to the symmetric (i+1, i) entry (Hessian
94     % is symmetric)
95 end
96
97
98 % Compute the corner elements: (1,n) and (n,1), assuming wrap-around connection.
99 if numel(h) > 1
100     h1 = h(1); % Step size for the first component
101     hn = h(n); % Step size for the n-th (last) component
102 else
103     h1 = h; % Use uniform step size if h is a scalar
104     hn = h; % Use uniform step size if h is a scalar
105 end
106 x_temp = x; % Create a temporary copy of x
107 x_temp(1) = x_temp(1) + h1; % Perturb the first component by +h1
108 x_temp(n) = x_temp(n) + hn; % Perturb the n-th component by +hn

```

```

107 f_pair = f(x_temp); % Evaluate f at the point perturbed in both first and n-th directions
108
109
110 % Finite difference formula for corner elements (1, n) and (n, 1):
111 % H(1, n) = H(n, 1)      (f(x + h_1*e_1 + h_n*e_n) - f(x + h_1*e_1) - f(x + h_n*e_n) + f(x))
112 %           / (h_1 * h_n)
113 value = (f_pair - f_plus(1) - f_plus(n) + f0) / (h1 * hn);
114 Hessfx(1, n) = value; % Assign the computed value to the (1, n) entry
115 Hessfx(n, 1) = value; % Assign the same value to the symmetric (n, 1) entry
116
117 end

```

Listing 15: Hessian with Finite Differences

```

1 function [output_str] = gridSearchFDParameters(initial_conditions, F, gradF, hessF)
2 %GRIDSEARCHFDPARAMETERS Executes a grid search for finite difference parameters.
3 % This function performs a grid search to evaluate the performance of
4 % optimization algorithms (Modified Newton and Truncated Newton) using
5 % finite difference (FD) approximations for gradients and Hessians.
6 % It explores different step sizes for finite differences and compares
7 % central and forward difference schemes.
8 %
9 % INPUTS:
10 % initial_conditions : Cell array of initial points for optimization algorithms.
11 % F                  : Function handle for the objective function F(x).
12 % gradF              : Function handle for the analytical gradient of F, gradF(x).
13 % hessF              : Function handle for the analytical Hessian of F, hessF(x).
14 %
15 % OUTPUTS:
16 % output_str         : String containing formatted results of the grid search,
17 %                       summarizing performance metrics for different FD parameters.
18 % Results            : Structure (currently not returned, but could be implemented)
19 %                       containing detailed numerical results, including
20 %                       iterations, function values, and gradient norms for different
21 %                       parameter settings and methods.
22
23
24 n      = 50;          % Dimension of the problem (vector x is of size n)
25 kmax   = 1000;        % Maximum number of iterations for optimization algorithms
26 tolgrad = 1e-6;       % Tolerance for the gradient norm (stopping criterion)
27 c1     = 1e-4;        % Armijo condition parameter for line search
28 rho    = 0.5;         % Backtracking factor for line search
29 btmax  = 100;         % Maximum number of backtracking iterations in line search
30
31
32 %% 3. DEFINITION OF THE PARAMETER "GRID"
33 num_starting_points = length(initial_conditions); % Number of initial starting points to
34 % test
34 exponent_list = [2, 4, 6, 8, 10, 12]; % Exponents for step size h = 10^(-k), where k is
35 % from this list
35 n_params = length(exponent_list); % Number of different finite difference parameters to test
36
37
38 %% 4. PREALLOCATION OF RESULTS MATRICES
39 % Matrices to store iteration counts and final gradient norms for each configuration.
40 % The results are categorized by:
41 % - Finite Difference Type: Uniform (h is constant for all components) and Variable (h
42 %   varies with x components)
43 % - Optimization Method: Modified Newton and Truncated Newton
44 % - Finite Difference Scheme: Central ('c') and Forward ('fw')
45
46 % --- Uniform FD (h = 10^(-k)) ---
47 % [MODIFIED NEWTON METHOD]
47 iter_uniform_mod_c = zeros(num_starting_points, n_params); % Iterations for Modified
48 % Newton, Uniform FD, Central scheme
48 grad_uniform_mod_c = zeros(num_starting_points, n_params); % Gradient norm for Modified
% Newton, Uniform FD, Central scheme

```

```

49 iter_uniform_mod_fw = zeros(num_starting_points, n_params); % Iterations for Modified
    Newton, Uniform FD, Forward scheme
50 grad_uniform_mod_fw = zeros(num_starting_points, n_params); % Gradient norm for Modified
    Newton, Uniform FD, Forward scheme
51
52
53 % [TRUNCATED NEWTON METHOD]
54 iter_uniform_trunc_c = zeros(num_starting_points, n_params); % Iterations for Truncated
    Newton, Uniform FD, Central scheme
55 grad_uniform_trunc_c = zeros(num_starting_points, n_params); % Gradient norm for Truncated
    Newton, Uniform FD, Central scheme
56 iter_uniform_trunc_fw = zeros(num_starting_points, n_params); % Iterations for Truncated
    Newton, Uniform FD, Forward scheme
57 grad_uniform_trunc_fw = zeros(num_starting_points, n_params); % Gradient norm for Truncated
    Newton, Uniform FD, Forward scheme
58
59
60 % --- FD with Variable h (h is proportional to |x|) ---
61 % [MODIFIED NEWTON METHOD]
62 iter_var_mod_c = zeros(num_starting_points, n_params); % Iterations for Modified Newton,
    Variable FD, Central scheme
63 grad_var_mod_c = zeros(num_starting_points, n_params); % Gradient norm for Modified Newton
    , Variable FD, Central scheme
64 iter_var_mod_fw = zeros(num_starting_points, n_params); % Iterations for Modified Newton,
    Variable FD, Forward scheme
65 grad_var_mod_fw = zeros(num_starting_points, n_params); % Gradient norm for Modified Newton
    , Variable FD, Forward scheme
66
67
68 % [TRUNCATED NEWTON METHOD]
69 iter_var_trunc_c = zeros(num_starting_points, n_params); % Iterations for Truncated Newton,
    Variable FD, Central scheme
70 grad_var_trunc_c = zeros(num_starting_points, n_params); % Gradient norm for Truncated
    Newton, Variable FD, Central scheme
71 iter_var_trunc_fw = zeros(num_starting_points, n_params); % Iterations for Truncated Newton,
    Variable FD, Forward scheme
72 grad_var_trunc_fw = zeros(num_starting_points, n_params); % Gradient norm for Truncated
    Newton, Variable FD, Forward scheme
73
74
75 %% 5. GRID SEARCH: LOOP OVER STARTING POINTS AND PARAMETERS
76 % This section performs the core grid search. It iterates through:
77 % - Initial starting points (defined in 'initial_conditions')
78 % - Exponents for step size 'h' (defined in 'exponent_list')
79 % For each combination, it tests four scenarios:
80 % 1. Modified Newton with Uniform FD (Central Difference)
81 % 2. Modified Newton with Uniform FD (Forward Difference)
82 % 3. Truncated Newton with Uniform FD (Central Difference)
83 % 4. Truncated Newton with Uniform FD (Forward Difference)
84 % 5. Modified Newton with Variable FD (Central Difference)
85 % 6. Modified Newton with Variable FD (Forward Difference)
86 % 7. Truncated Newton with Variable FD (Central Difference)
87 % 8. Truncated Newton with Variable FD (Forward Difference)
88
89
90 for i = 1:num_starting_points % Loop through each initial starting point
91 x0 = initial_conditions{i}; % Get the current initial starting point
92 for j = 1:n_params % Loop through each finite difference parameter (exponent k)
93 kexp = exponent_list(j); % Get the current exponent k for h = 10^(-k)
94
95
96 %% 5.1 Uniform FD: h = 10^(-kexp)
97 h_uniform = 10^(-kexp); % Calculate the uniform step size h
98 % Define function handles for finite difference gradient approximations using
99 % central and forward schemes
100 grad_fdiff_c = @(x) findiff_grad(F, x, h_uniform, 'c'); % Central difference
    gradient approximation
    grad_fdiff_fw = @(x) findiff_grad(F, x, h_uniform, 'fw'); % Forward difference

```

```

101     gradient approximation
102     Hess_fdiff = @(x) findiff_Hess(F, x, h_uniform); % Hessian approximation (using
103         the same uniform h)
104
105
106     % --- Modified Newton Method with Uniform FD ---
107     % Modified Newton, Uniform FD, Central Difference Scheme
108     try
109         [~, ~, grad_norm_mod, iter_mod, ~, ~, ~] = ...
110             modified_newton(x0, F, grad_fdiff_c, Hess_fdiff, kmax, tolgrad, c1, rho,
111                 btmax); % Run Modified Newton
112         iter_uniform_mod_c(i,j) = iter_mod; % Store iterations
113         grad_uniform_mod_c(i,j) = grad_norm_mod; % Store final gradient norm
114     catch ME
115         disp(['Error in modified_newton (uniform, central) for starting point ' num2str(
116             i) ', k = ' num2str(kexp) ': ' ME.message]); % Display error message if
117             method fails
118         iter_uniform_mod_c(i,j) = NaN; % Store NaN for iterations in case of error
119         grad_uniform_mod_c(i,j) = NaN; % Store NaN for gradient norm in case of error
120     end
121
122
123     % Modified Newton, Uniform FD, Forward Difference Scheme
124     try
125         [~, ~, grad_norm_mod, iter_mod, ~, ~, ~] = ...
126             modified_newton(x0, F, grad_fdiff_fw, Hess_fdiff, kmax, tolgrad, c1, rho,
127                 btmax); % Run Modified Newton
128         iter_uniform_mod_fw(i,j) = iter_mod; % Store iterations
129         grad_uniform_mod_fw(i,j) = grad_norm_mod; % Store final gradient norm
130     catch ME
131         disp(['Error in modified_newton (uniform, forward) for starting point ' num2str(
132             i) ', k = ' num2str(kexp) ': ' ME.message]); % Display error message if
133             method fails
134         iter_uniform_mod_fw(i,j) = NaN; % Store NaN for iterations in case of error
135         grad_uniform_mod_fw(i,j) = NaN; % Store NaN for gradient norm in case of error
136     end
137
138
139     % --- Truncated Newton Method with Uniform FD ---
140     % Truncated Newton, Uniform FD, Central Difference Scheme
141     try
142         [~, ~, grad_norm_trunc, iter_trunc, ~, ~, ~] = ...
143             truncated_newton_pre(x0, F, grad_fdiff_c, Hess_fdiff, kmax, tolgrad, c1, rho,
144                 btmax); % Run Truncated Newton
145         iter_uniform_trunc_c(i,j) = iter_trunc; % Store iterations
146         grad_uniform_trunc_c(i,j) = grad_norm_trunc; % Store final gradient norm
147     catch ME
148         disp(['Error in truncated_newton_pre (uniform, central) for starting point ' num2str(i) ', k = ' num2str(kexp) ': ' ME.message]); % Display error message
149             if method fails
150         iter_uniform_trunc_c(i,j) = NaN; % Store NaN for iterations in case of error
151         grad_uniform_trunc_c(i,j) = NaN; % Store NaN for gradient norm in case of error
152     end
153
154
155     % Truncated Newton, Uniform FD, Forward Difference Scheme
156     try
157         [~, ~, grad_norm_trunc, iter_trunc, ~, ~, ~] = ...
158             truncated_newton_pre(x0, F, grad_fdiff_fw, Hess_fdiff, kmax, tolgrad, c1,
159                 rho, btmax); % Run Truncated Newton
160         iter_uniform_trunc_fw(i,j) = iter_trunc; % Store iterations
161         grad_uniform_trunc_fw(i,j) = grad_norm_trunc; % Store final gradient norm
162     catch ME
163         disp(['Error in truncated_newton_pre (uniform, forward) for starting point ' num2str(i) ', k = ' num2str(kexp) ': ' ME.message]); % Display error message
164             if method fails
165         iter_uniform_trunc_fw(i,j) = NaN; % Store NaN for iterations in case of error
166         grad_uniform_trunc_fw(i,j) = NaN; % Store NaN for gradient norm in case of error

```

```

155 end
156
157
158 %% 5.2 Variable FD: h_var = h_uniform * abs(x0)
159 h_var = h_uniform * abs(x0); % Calculate variable step size h based on initial point
160 x0
161 grad_fdiff_var_c = @(x) findiff_grad(F, x, h_var, 'c'); % Central difference
162 gradient with variable h
163 grad_fdiff_var_fw = @(x) findiff_grad(F, x, h_var, 'fw'); % Forward difference
164 gradient with variable h
165 Hess_fdiff_var = @(x) findiff_Hess(F, x, h_var); % Hessian approximation with
166 variable h
167
168
169 % --- Modified Newton Method with Variable FD ---
170 % Modified Newton, Variable FD, Central Difference Scheme
171 try
172 [~, ~, grad_norm_mod, iter_mod, ~, ~, ~] = ...
173 modified_newton(x0, F, grad_fdiff_var_c, Hess_fdiff_var, kmax, tolgrad, c1,
174 rho, btmax); % Run Modified Newton
175 iter_var_mod_c(i,j) = iter_mod; % Store iterations
176 grad_var_mod_c(i,j) = grad_norm_mod; % Store final gradient norm
177 catch ME
178 disp(['Error in modified_newton (variable, central) for starting point ' num2str
179 (i)', kesp = ' num2str(kexp) ': ' ME.message]); % Display error message if
180 method fails
181 iter_var_mod_c(i,j) = NaN; % Store NaN for iterations in case of error
182 grad_var_mod_c(i,j) = NaN; % Store NaN for gradient norm in case of error
183 end
184
185
186 % Modified Newton, Variable FD, Forward Difference Scheme
187 try
188 [~, ~, grad_norm_mod, iter_mod, ~, ~, ~] = ...
189 modified_newton(x0, F, grad_fdiff_var_fw, Hess_fdiff_var, kmax, tolgrad, c1,
190 rho, btmax); % Run Modified Newton
191 iter_var_mod_fw(i,j) = iter_mod; % Store iterations
192 grad_var_mod_fw(i,j) = grad_norm_mod; % Store final gradient norm
193 catch ME
194 disp(['Error in modified_newton (variable, forward) for starting point ' num2str
195 (i)', kesp = ' num2str(kexp) ': ' ME.message]); % Display error message if
196 method fails
197 iter_var_mod_fw(i,j) = NaN; % Store NaN for iterations in case of error
198 grad_var_mod_fw(i,j) = NaN; % Store NaN for gradient norm in case of error
199 end
200
201
202 % --- Truncated Newton Method with Variable FD ---
203 % Truncated Newton, Variable FD, Central Difference Scheme
204 try
205 [~, ~, grad_norm_trunc, iter_trunc, ~, ~, ~] = ...
206 truncated_newton_pre(x0, F, grad_fdiff_var_c, Hess_fdiff_var, kmax, tolgrad,
207 c1, rho, btmax); % Run Truncated Newton
208 iter_var_trunc_c(i,j) = iter_trunc; % Store iterations
209 grad_var_trunc_c(i,j) = grad_norm_trunc; % Store final gradient norm
210 catch ME
211 disp(['Error in truncated_newton_pre (variable, central) for starting point ' num2str(i)', kesp = ' num2str(kexp) ': ' ME.message]); % Display error
212 message if method fails
213 iter_var_trunc_c(i,j) = NaN; % Store NaN for iterations in case of error
214 grad_var_trunc_c(i,j) = NaN; % Store NaN for gradient norm in case of error
215 end
216
217
218 % Truncated Newton, Variable FD, Forward Difference Scheme
219 try
220 [~, ~, grad_norm_trunc, iter_trunc, ~, ~, ~] = ...
221 truncated_newton_pre(x0, F, grad_fdiff_var_fw, Hess_fdiff_var, kmax, tolgrad,
222

```

```

210           , c1, rho, btmax); % Run Truncated Newton
211   iter_var_trunc_fw(i,j) = iter_trunc; % Store iterations
212   grad_var_trunc_fw(i,j) = grad_norm_trunc; % Store final gradient norm
213 catch ME
214     disp(['Error in truncated_newton_pre (variable, forward) for starting point '
215           num2str(i) ', kesp = ' num2str(kexp) ': ' ME.message]); % Display error
216           message if method fails
217     iter_var_trunc_fw(i,j) = NaN; % Store NaN for iterations in case of error
218     grad_var_trunc_fw(i,j) = NaN; % Store NaN for gradient norm in case of error
219   end
220 end
221
222
223 %% 6. CALCULATION OF AVERAGES (OVER STARTING POINTS)
224 % Average the iteration counts and gradient norms over all starting points for each
225 % parameter setting.
226 % Uniform FD Averages
227 avg_iter_uniform_mod_c = mean(iter_uniform_mod_c, 1); % Average iterations for Modified
228   Newton, Uniform FD, Central
229 avg_grad_uniform_mod_c = mean(grad_uniform_mod_c, 1); % Average gradient norm for Modified
230   Newton, Uniform FD, Central
231 avg_iter_uniform_mod_fw = mean(iter_uniform_mod_fw, 1); % Average iterations for Modified
232   Newton, Uniform FD, Forward
233 avg_grad_uniform_mod_fw = mean(grad_uniform_mod_fw, 1); % Average gradient norm for Modified
234   Newton, Uniform FD, Forward
235
236
237 % Truncated Newton, Uniform FD, Central
238 avg_iter_uniform_trunc_c = mean(iter_uniform_trunc_c, 1); % Average iterations for
239   Truncated Newton, Uniform FD, Central
240 avg_grad_uniform_trunc_c = mean(grad_uniform_trunc_c, 1); % Average gradient norm for
241   Truncated Newton, Uniform FD, Central
242 avg_iter_uniform_trunc_fw = mean(iter_uniform_trunc_fw, 1); % Average iterations for
243   Truncated Newton, Uniform FD, Forward
244 avg_grad_uniform_trunc_fw = mean(grad_uniform_trunc_fw, 1); % Average gradient norm for
245   Truncated Newton, Uniform FD, Forward
246
247 % Variable FD Averages
248 avg_iter_var_mod_c = mean(iter_var_mod_c, 1); % Average iterations for Modified Newton,
249   Variable FD, Central
250 avg_grad_var_mod_c = mean(grad_var_mod_c, 1); % Average gradient norm for Modified Newton,
251   Variable FD, Central
252 avg_iter_var_mod_fw = mean(iter_var_mod_fw, 1); % Average iterations for Modified Newton,
253   Variable FD, Forward
254 avg_grad_var_mod_fw = mean(grad_var_mod_fw, 1); % Average gradient norm for Modified Newton,
255   Variable FD, Forward
256
257 %% 7. SELECTION OF THE BEST PARAMETERS (for Modified Newton, Uniform FD)
258 % For demonstration, select the "best" parameter based on Modified Newton with Uniform FD.
259 % "Best" is defined as the parameter that achieves the minimum gradient norm,
260 % and among those, the minimum number of iterations.
261
262 % Central Difference - find best k based on minimum gradient norm, then minimum iterations
263   among candidates

```

```

257 [min_grad_uniform_mod_c, best_index_uniform_mod_c] = min(avg_grad_uniform_mod_c); % Find
258     minimum average gradient norm and its index
259 candidates = find(abs(avg_grad_uniform_mod_c - min_grad_uniform_mod_c) < 1e-12); % Find
260     indices with gradient norm close to minimum
261 if length(candidates) > 1 % If multiple candidates with similar gradient norms, choose based
262     on minimum iterations
263 [~, idx] = min(avg_iter_uniform_mod_c(candidates)); % Find index of minimum iterations
264     among candidates
265 best_index_uniform_mod_c = candidates(idx); % Update best index to the one with minimum
266     iterations
267 end
268 best_k_uniform_mod_c = exponent_list(best_index_uniform_mod_c); % Get the best k exponent
269     for Uniform FD, Central scheme
270
271 % Forward Difference - find best k based on minimum gradient norm, then minimum iterations
272     among candidates
273 [min_grad_uniform_mod_fw, best_index_uniform_mod_fw] = min(avg_grad_uniform_mod_fw); % Find
274     minimum average gradient norm and its index
275 candidates = find(abs(avg_grad_uniform_mod_fw - min_grad_uniform_mod_fw) < 1e-12); % Find
276     indices with gradient norm close to minimum
277 if length(candidates) > 1 % If multiple candidates with similar gradient norms, choose based
278     on minimum iterations
279 [~, idx] = min(avg_iter_uniform_mod_fw(candidates)); % Find index of minimum iterations
280     among candidates
281 best_index_uniform_mod_fw = candidates(idx); % Update best index to the one with minimum
282     iterations
283 end
284 best_k_uniform_mod_fw = exponent_list(best_index_uniform_mod_fw); % Get the best k exponent
285     for Uniform FD, Forward scheme
286
287 %% 8. CONSTRUCTION OF THE OUTPUT STRING (output_str)
288 % Format the results into a readable string for display.
289 output_str = ""; % Initialize output string
290 output_str = output_str + newline + "%% VISUALIZATION OF RESULTS" + newline; % Add section
291     header for visualization
292 output_str = output_str + "*** Grid Search Results ***" + newline + newline; % Add title for
293     grid search results
294
295 % --- Results for Uniform FD ---
296 output_str = output_str + "--- Uniform Finite Differences ---" + newline; % Sub-section
297     header for Uniform FD
298 output_str = output_str + "Modified Newton Method:" + newline; % Method sub-header
299 output_str = output_str + "Central Difference Scheme:" + newline; % Scheme sub-header
300 output_str = output_str + "Tested Parameters (k): " + string(mat2str(exponent_list)) +
301     newline; % List of tested k parameters
302 output_str = output_str + "Average Iterations: " + string(mat2str(avg_iter_uniform_mod_c, 4))
303     + newline; % Average iterations array as string
304 output_str = output_str + "Average Gradient Norm: " + string(mat2str(avg_grad_uniform_mod_c,
305     4)) + newline; % Average gradient norm array as string
306 output_str = output_str + "=> Best k = " + string(num2str(best_k_uniform_mod_c)) + " (
307     Average Iterations = " + ...
308     string(num2str(avg_iter_uniform_mod_c(best_index_uniform_mod_c), '%.2f')) + ", Average
309     Gradient Norm = " + ...
310     string(num2str(avg_grad_uniform_mod_c(best_index_uniform_mod_c), '%.3e')) + ")" +
311     newline + newline; % Best k and its performance metrics
312
313 output_str = output_str + "Forward Difference Scheme:" + newline; % Scheme sub-header
314 output_str = output_str + "Tested Parameters (k): " + string(mat2str(exponent_list)) +
315     newline; % List of tested k parameters
316 output_str = output_str + "Average Iterations: " + string(mat2str(avg_iter_uniform_mod_fw,
317     4)) + newline; % Average iterations array as string
318 output_str = output_str + "Average Gradient Norm: " + string(mat2str(avg_grad_uniform_mod_fw
319     , 4)) + newline; % Average gradient norm array as string
320 output_str = output_str + "=> Best k = " + string(num2str(best_k_uniform_mod_fw)) + " (
321     Average Iterations = " + ...

```

```

299     string(num2str(avg_iter_uniform_mod_fw(best_index_uniform_mod_fw), '%.2f')) + ", Average
      Gradient Norm = " +
300     string(num2str(avg_grad_uniform_mod_fw(best_index_uniform_mod_fw), '%.3e')) + ")" +
      newline + newline; % Best k and its performance metrics
301
302
303 output_str = output_str + "Truncated Newton Method:" + newline; % Method sub-header
304 output_str = output_str + "Central Difference Scheme:" + newline; % Scheme sub-header
305 output_str = output_str + "Tested Parameters (k): " + string(mat2str(exponent_list)) +
      newline; % List of tested k parameters
306 output_str = output_str + "Average Iterations: " + string(mat2str(avg_iter_uniform_trunc_c,
      4)) + newline; % Average iterations array as string
307 output_str = output_str + "Average Gradient Norm: " + string(mat2str(
      avg_grad_uniform_trunc_c, 4)) + newline; % Average gradient norm array as string
308 output_str = output_str + "=> Best k = " + string(num2str(best_k_uniform_mod_c)) + (
      Average Iterations = " + ...
      string(num2str(avg_iter_uniform_trunc_c(best_index_uniform_mod_c), '%.2f')) + ", Average
      Gradient Norm = " + ...
      string(num2str(avg_grad_uniform_trunc_c(best_index_uniform_mod_c), '%.3e')) + ")" +
      newline + newline; % Best k and its performance metrics
309
310
311
312
313 output_str = output_str + "Forward Difference Scheme:" + newline; % Scheme sub-header
314 output_str = output_str + "Tested Parameters (k): " + string(mat2str(exponent_list)) +
      newline; % List of tested k parameters
315 output_str = output_str + "Average Iterations: " + string(mat2str(avg_iter_uniform_trunc_fw,
      4)) + newline; % Average iterations array as string
316 output_str = output_str + "Average Gradient Norm: " + string(mat2str(avg_grad_uniform_mod_fw
      , 4)) + newline; % Average gradient norm array as string
317 output_str = output_str + "=> Best k = " + string(num2str(best_k_uniform_mod_fw)) + (
      Average Iterations = " + ...
      string(num2str(avg_iter_uniform_trunc_fw(best_index_uniform_mod_fw), '%.2f')) + ,
      Average Gradient Norm = " + ...
      string(num2str(avg_grad_uniform_trunc_fw(best_index_uniform_mod_fw), '%.3e')) + ")" +
      newline + newline; % Best k and its performance metrics
318
319
320
321
322 % --- Results for Variable FD ---
323 output_str = output_str + newline + "--- Finite Differences with Variable h ---" + newline;
      % Sub-section header for Variable FD
324 output_str = output_str + "Modified Newton Method:" + newline; % Method sub-header
325 output_str = output_str + "Central Difference Scheme:" + newline; % Scheme sub-header
326 output_str = output_str + "Tested Parameters (kesp): " + string(mat2str(exponent_list)) +
      newline; % List of tested k parameters
327 output_str = output_str + "Average Iterations: " + string(mat2str(avg_iter_var_mod_c, 4)) +
      newline; % Average iterations array as string
328 output_str = output_str + "Average Gradient Norm: " + string(mat2str(avg_grad_var_mod_c, 4))
      + newline; % Average gradient norm array as string
329 output_str = output_str + "=> Best kesp = " + string(num2str(exponent_list(
      best_index_uniform_mod_c))) + " (Average Iterations = " + ...
      string(num2str(avg_iter_var_mod_c(best_index_uniform_mod_c), '%.2f')) + ", Average
      Gradient Norm = " + ...
      string(num2str(avg_grad_var_mod_c(best_index_uniform_mod_c), '%.3e')) + ")" + newline +
      newline; % Best k and its performance metrics
330
331
332
333
334 output_str = output_str + "Forward Difference Scheme:" + newline; % Scheme sub-header
335 output_str = output_str + "Tested Parameters (kesp): " + string(mat2str(exponent_list)) +
      newline; % List of tested k parameters
336 output_str = output_str + "Average Iterations: " + string(mat2str(avg_iter_var_mod_fw, 4)) +
      newline; % Average iterations array as string
337 output_str = output_str + "Average Gradient Norm: " + string(mat2str(avg_grad_var_mod_fw, 4))
      + newline; % Average gradient norm array as string
338 output_str = output_str + "=> Best kesp = " + string(num2str(exponent_list(
      best_index_uniform_mod_fw))) + " (Average Iterations = " + ...
      string(num2str(avg_iter_var_mod_fw(best_index_uniform_mod_fw), '%.2f')) + ", Average
      Gradient Norm = " + ...
      string(num2str(avg_grad_var_mod_fw(best_index_uniform_mod_fw), '%.3e')) + ")" + newline
339
340

```

```

341      + newline; % Best k and its performance metrics
342
343 output_str = output_str + "Truncated Newton Method:" + newline; % Method sub-header
344 output_str = output_str + "Central Difference Scheme:" + newline; % Scheme sub-header
345 output_str = output_str + "Tested Parameters (kesp): " + string(mat2str(exponent_list)) +
    newline; % List of tested k parameters
346 output_str = output_str + "Average Iterations: " + string(mat2str(avg_iter_var_trunc_c, 4))
    + newline; % Average iterations array as string
347 output_str = output_str + "Average Gradient Norm: " + string(mat2str(avg_grad_var_trunc_c,
    4)) + newline; % Average gradient norm array as string
348 output_str = output_str + "=> Best kesp = " + string(num2str(exponent_list(
    best_index_uniform_mod_c))) + " (Average Iterations = " + ...
    string(num2str(avg_iter_var_trunc_c(best_index_uniform_mod_c), '%.2f')) + ", Average
    Gradient Norm = " + ...
    string(num2str(avg_grad_var_trunc_c(best_index_uniform_mod_c), '%.3e')) + ")" + newline
    + newline; % Best k and its performance metrics
349
350
351
352
353 output_str = output_str + "Forward Difference Scheme:" + newline; % Scheme sub-header
354 output_str = output_str + "Tested Parameters (kesp): " + string(mat2str(exponent_list)) +
    newline; % List of tested k parameters
355 output_str = output_str + "Average Iterations: " + string(mat2str(avg_iter_var_trunc_fw, 4))
    + newline; % Average iterations array as string
356 output_str = output_str + "Average Gradient Norm: " + string(mat2str(avg_grad_var_fw, 4)) +
    newline; % Average gradient norm array as string
357 output_str = output_str + "=> Best kesp = " + string(num2str(exponent_list(
    best_index_uniform_mod_fw))) + " (Average Iterations = " + ...
    string(num2str(avg_iter_var_trunc_fw(best_index_uniform_mod_fw), '%.2f')) + ", Average
    Gradient Norm = " + ...
    string(num2str(avg_grad_var_trunc_fw(best_index_uniform_mod_fw), '%.3e')) + ")" + +
    newline + newline; % Best k and its performance metrics
358
359
360
361
362 output_str = output_str + "*** End of Grid Search Analysis ***" + newline; % Footer for
    output string
363 disp(output_str); % Display the formatted output string in the command window
364
365 end

```

Listing 16: Grid search - Finite Differences

References

- [1] Test Problems for Unconstrained Optimization. Available at: https://www.researchgate.net/publication/325314497_Test_Problems_for_Unconstrained_Optimization
- [2] Pierraccini, Sandra. *Numerical Optimization for Large-Scale Problems*. Lecture notes, Politecnico di Torino, A.Y. 2024/2025.
- [3] Della Santa, Francesco. *Numerical Optimization for Large Scale Problems*. Lecture Notes, Politecnico di Torino, A.Y. 2024/2025.